

Complete source code of `dlmalloc`

Written by Doug Lea
Typeset by John Wickerson

Contents

1	Introduction	3	12	System alloc setup	61
2	Compile-time options	8	13	Hooks	63
3	Mallinfo declarations	18	14	Debugging setup	65
4	Declarations of public routines	20	15	Bins	66
5	Internal <code>#includes</code>	31	16	Runtime Check Support	70
6	<code>size_t</code> and alignment properties	35	17	Setting mparams	73
7	MMAP preliminaries	36	18	Debugging Support	76
8	Locks	40	19	Statistics	82
9	Chunks	47	20	Operations on smallbins and trees	84
10	Segments	55	21	Direct-mmapping chunks	90
11	State	57	22	mSPACE management	92
			23	System allocation and deallocation	95
			24	Support for public routines	102
			25	Public routines	111
			26	User mspaces	119
			27	Postscript	129

Chapter 1

Introduction

```
/*
  This is a version (aka dlmalloc) of malloc/free/realloc written by
  Doug Lea and released to the public domain, as explained at
  4 http://creativecommons.org/licenses/publicdomain. Send questions,
  comments, complaints, performance data, etc to dl@cs.oswego.edu

  * Version 2.8.4 Wed May 27 09:56:23 2009 Doug Lea (dl at gee)

  Note: There may be an updated version of this malloc obtainable at
        ftp://gee.cs.oswego.edu/pub/misc/malloc.c
  Check before installing!
```

1.1 Quickstart

```
16 This library is all in one file to simplify the most common usage:
  ftp it, compile it (-O3), and link it into another program. All of
  the compile-time options default to reasonable values for use on
  most platforms. You might later want to step through various
  compile-time and dynamic tuning options.
```

```
For convenience, an include file for code using this malloc is at:
  ftp://gee.cs.oswego.edu/pub/misc/malloc-2.8.4.h
  You don't really need this .h file unless you call functions not
  24 defined in your system include files. The .h file contains only the
  excerpts from this file needed for using this malloc on ANSI C/C++
  systems, so long as you haven't changed compile-time options about
  naming and tuning parameters. If you do, then you can create your
  28 own malloc.h that does include all settings by cutting at the point
  indicated below. Note that you may already by default be using a C
  library containing a malloc that is based on some version of this
  malloc (for example in linux). You might still want to use the one
  32 in this file to customize settings or to avoid overheads associated
  with library versions.
```

1.2 Vital statistics

Supported pointer/size_t representation: 4 or 8 bytes
 size_t MUST be an unsigned type of the same width as
 pointers. (If you are using an ancient system that declares
 size_t as a signed type, or need it to be a different width
 than pointers, you can use a previous release of this malloc
 (e.g. 2.7.2) supporting these.)

Alignment: 8 bytes (default)
 This suffices for nearly all current machines and C compilers.
 However, you can define MALLOC_ALIGNMENT to be wider than this
 if necessary (up to 128bytes), at the expense of using more space.

Minimum overhead per allocated chunk: 4 or 8 bytes (if 4byte sizes)
 8 or 16 bytes (if 8byte sizes)
 Each malloced chunk has a hidden word of overhead holding size
 and status information, and additional cross-check word
 if FOOTERS is defined.

Minimum allocated size: 4-byte ptrs: 16 bytes (including overhead)
 8-byte ptrs: 32 bytes (including overhead)

Even a request for zero bytes (i.e., malloc(0)) returns a
 pointer to something of the minimum allocatable size.
 The maximum overhead wastage (i.e., number of extra bytes
 allocated than were requested in malloc) is less than or equal
 to the minimum size, except for requests \geq mmap_threshold that
 are serviced via mmap(), where the worst case wastage is about
 32 bytes plus the remainder from a system page (the minimal
 mmap unit); typically 4096 or 8192 bytes.

Security: static-safe; optionally more or less
 The "security" of malloc refers to the ability of malicious
 code to accentuate the effects of errors (for example, freeing
 space that is not currently malloc'ed or overwriting past the
 ends of chunks) in code that calls malloc. This malloc
 guarantees not to modify any memory locations below the base of
 heap, i.e., static variables, even in the presence of usage
 errors. The routines additionally detect most improper frees
 and reallocs. All this holds as long as the static bookkeeping
 for malloc itself is not corrupted by some other means. This
 is only one aspect of security -- these checks do not, and
 cannot, detect all possible programming errors.

If FOOTERS is defined nonzero, then each allocated chunk
 carries an additional check word to verify that it was malloced
 from its space. These check words are the same within each
 execution of a program using malloc, but differ across
 executions, so externally crafted fake chunks cannot be

freed. This improves security by rejecting frees/reallocs that could corrupt heap memory, in addition to the checks preventing writes to statics that are always on. This may further improve security at the expense of time and space overhead. (Note that FOOTERS may also be worth using with MSPACES.)

By default detected errors cause the program to abort (calling "abort()"). You can override this to instead proceed past errors by defining PROCEED_ON_ERROR. In this case, a bad free has no effect, and a malloc that encounters a bad address caused by user overwrites will ignore the bad address by dropping pointers and indices to all known memory. This may be appropriate for programs that should continue if at all possible in the face of programming errors, although they may run out of memory because dropped memory is never reclaimed.

If you don't like either of these options, you can define CORRUPTION_ERROR_ACTION and USAGE_ERROR_ACTION to do anything else. And if you are sure that your program using malloc has no errors or vulnerabilities, you can define INSECURE to 1, which might (or might not) provide a small performance improvement.

Thread-safety: NOT thread-safe unless USE_LOCKS defined

When USE_LOCKS is defined, each public call to malloc, free, etc is surrounded with either a pthread mutex or a win32 spinlock (depending on WIN32). This is not especially fast, and can be a major bottleneck. It is designed only to provide minimal protection in concurrent environments, and to provide a basis for extensions. If you are using malloc in a concurrent program, consider instead using nedmalloc (<http://www.nedprod.com/programs/portable/nedmalloc/>) or ptmalloc (See <http://www.malloc.de>), which are derived from versions of this malloc.

System requirements: Any combination of MORECORE and/or MMAP/MUNMAP

This malloc can use unix sbrk or any emulation (invoked using the CALL_MORECORE macro) and/or mmap/munmap or any emulation (invoked using CALL_MMAP/CALL_MUNMAP) to get and release system memory. On most unix systems, it tends to work best if both MORECORE and MMAP are enabled. On Win32, it uses emulations based on VirtualAlloc. It also uses common C library functions like memset.

Compliance: I believe it is compliant with the Single Unix Specification (See <http://www.unix.org>). Also SVID/XPG, ANSI C, and probably others as well.

1.3 Overview of algorithms

This is not the fastest, most space-conserving, most portable, or

most tunable malloc ever written. However it is among the fastest while also being among the most space-conserving, portable and tunable. Consistent balance across these factors results in a good general-purpose allocator for malloc-intensive programs.

In most ways, this malloc is a best-fit allocator. Generally, it chooses the best-fitting existing chunk for a request, with ties broken in approximately least-recently-used order. (This strategy normally maintains low fragmentation.) However, for requests less than 256bytes, it deviates from best-fit when there is not an exactly fitting available chunk by preferring to use space adjacent to that used for the previous small request, as well as by breaking ties in approximately most-recently-used order. (These enhance locality of series of small allocations.) And for very large requests ($\geq 256\text{Kb}$ by default), it relies on system memory mapping facilities, if supported. (This helps avoid carrying around and possibly fragmenting memory used only for large chunks.)

All operations (except `malloc_stats` and `mallinfo`) have execution times that are bounded by a constant factor of the number of bits in a `size_t`, not counting any clearing in `calloc` or copying in `realloc`, or actions surrounding `MORECORE` and `MMAP` that have times proportional to the number of non-contiguous regions returned by system allocation routines, which is often just 1. In real-time applications, you can optionally suppress segment traversals using `NO_SEGMENT_TRAVERSAL`, which assures bounded execution even when system allocators return non-contiguous spaces, at the typical expense of carrying around more memory and increased fragmentation.

The implementation is not very modular and seriously overuses macros. Perhaps someday all C compilers will do as good a job inlining modular code as can now be done by brute-force expansion, but now, enough of them seem not to.

Some compilers issue a lot of warnings about code that is dead/unreachable only on some platforms, and also about intentional uses of negation on unsigned types. All known cases of each can be ignored.

For a longer but out of date high-level description, see <http://gee.cs.oswego.edu/dl/html/malloc.html>

1.4 MSPACES

If `MSPACES` is defined, then in addition to `malloc`, `free`, etc., this file also defines `mSPACE_malloc`, `mSPACE_free`, etc. These are versions of `malloc` routines that take an "mSPACE" argument obtained using `create_mSPACE`, to control all internal bookkeeping. If `ONLY_MSPACES` is defined, only these versions are compiled. So if you would like to use this allocator for only some allocations, and your system `malloc` for others, you can compile with

ONLY_MSPACES and then do something like...

```
static mspace mymspace = create_mspace(0,0); // for example
#define mymalloc(bytes) mspace_malloc(mymspace, bytes)
```

(Note: If you only need one instance of an mspace, you can instead use "USE_DL_PREFIX" to relabel the global malloc.)

192 You can similarly create thread-local allocators by storing
mspaces as thread-locals. For example:

```
static __thread mspace tlms = 0;
void* tlmalloc(size_t bytes) {
196     if (tlms == 0) tlms = create_mspace(0, 0);
    return mspace_malloc(tlms, bytes);
}
void tlfree(void* mem) { mspace_free(tlms, mem); }
```

Unless FOOTERS is defined, each mspace is completely independent.
You cannot allocate from one and free to another (although
conformance is only weakly checked, so usage errors are not always
204 caught). If FOOTERS is defined, then each chunk carries around a tag
indicating its originating mspace, and frees are directed to their
originating spaces.

*/

Chapter 2

Compile-time options

```
/*
Be careful in setting #define values for numerical constants of type
size_t. On some systems, literal values are not automatically extended
212 to size_t precision unless they are explicitly casted. You can also
use the symbolic values MAX_SIZE_T, SIZE_T_ONE, etc below.

WIN32                                default: defined if _WIN32 defined
216   Defining WIN32 sets up defaults for MS environment and compilers.
Otherwise defaults are for unix. Beware that there seem to be some
cases where this malloc might not be a pure drop-in replacement for
Win32 malloc: Random-looking failures from Win32 GDI API's (eg;
220   SetDIBits()) may be due to bugs in some video driver implementations
when pixel buffers are malloc()'ed, and the region spans more than
one VirtualAlloc()'ed region. Because dlmalloc uses a small (64Kb)
default granularity, pixel buffers may straddle virtual allocation
224   regions more often than when using the Microsoft allocator. You can
avoid this by using VirtualAlloc() and VirtualFree() for all pixel
buffers rather than using malloc(). If this is not possible,
recompile this malloc with a larger DEFAULT_GRANULARITY.

MALLOC_ALIGNMENT                    default: (size_t)8
Controls the minimum alignment for malloc'ed chunks. It must be a
power of two and at least 8, even on machines for which smaller
232   alignments would suffice. It may be defined as larger than this
though. Note however that code and data structures are optimized for
the case of 8-byte alignment.

236   MSPACES                                default: 0 (false)
If true, compile in support for independent allocation spaces.
This is only supported if HAVE_MMAP is true.

240   ONLY_MSPACES                            default: 0 (false)
If true, only compile in mspace versions, not regular versions.

USE_LOCKS                            default: 0 (false)
244   Causes each call to each public routine to be surrounded with
pthread or WIN32 mutex lock/unlock. (If set true, this can be
overridden on a per-mspace basis for mspace versions.) If set to a
non-zero value other than 1, locks are used, but their
```


248 implementation is left out, so lock functions must be supplied manually,
 as described below.

USE_SPIN_LOCKS default: 1 iff USE_LOCKS and on x86 using gcc or MSC
 252 If true, uses custom spin locks for locking. This is currently
 supported only for x86 platforms using gcc or recent MS compilers.
 Otherwise, posix locks or win32 critical sections are used.

256 FOOTERS default: 0
 If true, provide extra checking and dispatching by placing
 information in the footers of allocated chunks. This adds
 space and time overhead.

INSECURE default: 0
 If true, omit checks for usage errors and heap space overwrites.

264 USE_DL_PREFIX default: NOT defined
 Causes compiler to prefix all public routines with the string 'dl'.
 This can be useful when you only want to use this malloc in one part
 of a program, using your regular system malloc elsewhere.

ABORT default: defined as abort()
 Defines how to abort on failed checks. On most systems, a failed
 check cannot die with an "assert" or even print an informative
 272 message, because the underlying print routines in turn call malloc,
 which will fail again. Generally, the best policy is to simply call
 abort(). It's not very useful to do more than this because many
 errors due to overwriting will show up as address faults (null, odd
 276 addresses etc) rather than malloc-triggered checks, so will also
 abort. Also, most compilers know that abort() does not return, so
 can better optimize code conditionally calling it.

280 PROCEED_ON_ERROR default: defined as 0 (false)
 Controls whether detected bad addresses cause them to be bypassed
 rather than aborting. If set, detected bad arguments to free and
 realloc are ignored. And all bookkeeping information is zeroed out
 284 upon a detected overwrite of freed heap space, thus losing the
 ability to ever return it from malloc again, but enabling the
 application to proceed. If PROCEED_ON_ERROR is defined, the
 static variable malloc_corruption_error_count is compiled in
 288 and can be examined to see if errors have occurred. This option
 generates slower code than the default abort policy.

DEBUG default: NOT defined
 292 The DEBUG setting is mainly intended for people trying to modify
 this code or diagnose problems when porting to new platforms.
 However, it may also be able to better isolate user errors than just
 using runtime checks. The assertions in the check routines spell
 296 out in more detail the assumptions and invariants underlying the
 algorithms. The checking is fairly extensive, and will slow down
 execution noticeably. Calling malloc_stats or mallinfo with DEBUG
 set will attempt to check every non-mmapped allocated and free chunk

300 in the course of computing the summaries.

ABORT_ON_ASSERT_FAILURE default: defined as 1 (true)

304 Debugging assertion failures can be nearly impossible if your
version of the assert macro causes malloc to be called, which will
lead to a cascade of further failures, blowing the runtime stack.
ABORT_ON_ASSERT_FAILURE cause assertions failures to call abort(),
which will usually make debugging easier.

MALLOC_FAILURE_ACTION default: sets errno to ENOMEM, or no-op on win32
The action to take before "return 0" when malloc fails to be able to
return memory because there is none available.

HAVE_MORECORE default: 1 (true) unless win32 or ONLY_MSPACES
True if this system supports sbrk or an emulation of it.

316 MORECORE default: sbrk

The name of the sbrk-style system routine to call to obtain more
memory. See below for guidance on writing custom MORECORE
functions. The type of the argument to sbrk/MORECORE varies across
320 systems. It cannot be size_t, because it supports negative
arguments, so it is normally the signed type of the same width as
size_t (sometimes declared as "intptr_t"). It doesn't much matter
though. Internally, we only call it with arguments less than half
324 the max value of a size_t, which should work across all reasonable
possibilities, although sometimes generating compiler warnings.

MORECORE_CONTIGUOUS default: 1 (true) if HAVE_MORECORE

328 If true, take advantage of fact that consecutive calls to MORECORE
with positive arguments always return contiguous increasing
addresses. This is true of unix sbrk. It does not hurt too much to
set it true anyway, since malloc copes with non-contiguities.
332 Setting it false when definitely non-contiguous saves time
and possibly wasted space it would take to discover this though.

MORECORE_CANNOT_TRIM default: NOT defined

336 True if MORECORE cannot release space back to the system when given
negative arguments. This is generally necessary only if you are
using a hand-crafted MORECORE function that cannot handle negative
arguments.

NO_SEGMENT_TRAVERSAL default: 0

344 If non-zero, suppresses traversals of memory segments
returned by either MORECORE or CALL_MMAP. This disables
merging of segments that are contiguous, and selectively
releasing them to the OS if unused, but bounds execution times.

HAVE_MMAP default: 1 (true)

348 True if this system supports mmap or an emulation of it. If so, and
HAVE_MORECORE is not true, MMAP is used for all system
allocation. If set and HAVE_MORECORE is true as well, MMAP is
primarily used to directly allocate very large blocks. It is also

352 used as a backup strategy in cases where MORECORE fails to provide space from system. Note: A single call to MUNMAP is assumed to be able to unmap memory that may have be allocated using multiple calls to MMAP, so long as they are adjacent.

HAVE_MREMAP default: 1 on linux, else 0
If true realloc() uses mremap() to re-allocate large blocks and extend or shrink allocation spaces.

MMAP_CLEARS default: 1 except on WINCE.
True if mmap clears memory so calloc doesn't need to. This is true for standard unix mmap using /dev/zero and on WIN32 except for WINCE.

USE_BUILTIN_FFS default: 0 (i.e., not used)
Causes malloc to use the builtin ffs() function to compute indices. Some compilers may recognize and intrinsify ffs to be faster than the
368 supplied C version. Also, the case of x86 using gcc is special-cased to an asm instruction, so is already as fast as it can be, and so this setting has no effect. Similarly for Win32 under recent MS compilers. (On most x86s, the asm version is only slightly faster than the C version.)

malloc_getpagesize default: derive from system includes, or 4096.
The system page size. To the extent possible, this malloc manages memory from the system in page-size units. This may be (and
376 usually is) a function rather than a constant. This is ignored if WIN32, where page size is determined using getSystemInfo during initialization.

380 USE_DEV_RANDOM default: 0 (i.e., not used)
Causes malloc to use /dev/random to initialize secure magic seed for stamping footers. Otherwise, the current time is used.

384 NO_MALLINFO default: 0
If defined, don't compile "mallinfo". This can be a simple way of dealing with mismatches between system declarations and those in this file.

MALLINFO_FIELD_TYPE default: size_t
The type of the fields in the mallinfo struct. This was originally defined as "int" in SVID etc, but is more usefully defined as
392 size_t. The value is used only if HAVE_USR_INCLUDE_MALLOC_H is not set

REALLOC_ZERO_BYTES_FREES default: not defined
This should be set if a call to realloc with zero bytes should
396 be the same as a call to free. Some people think it should. Otherwise, since this malloc returns a unique pointer for malloc(0), so does realloc(p, 0).

400 LACKS_UNISTD_H, LACKS_FCNTL_H, LACKS_SYS_PARAM_H, LACKS_SYS_MMAN_H
LACKS_STRINGS_H, LACKS_STRING_H, LACKS_SYS_TYPES_H, LACKS_ERRNO_H
LACKS_STDLIB_H default: NOT defined unless on WIN32
Define these if your system does not have these header files.

404 You might need to manually insert some of the declarations they provide.

```

DEFAULT_GRANULARITY      default: page size if MORECORE_CONTIGUOUS,
                          system_info.dwAllocationGranularity in WIN32,
408                          otherwise 64K.

```

Also settable using `mallopt(M_GRANULARITY, x)`

The unit for allocating and deallocating memory from the system. On most systems with contiguous `MORECORE`, there is no reason to make this more than a page. However, systems with `MMAP` tend to either require or encourage larger granularities. You can increase this value to prevent system allocation functions to be called so often, especially if they are slow. The value must be at least one page and must be a power of two. Setting to 0 causes initialization to either page size or win32 region size. (Note: In previous versions of `malloc`, the equivalent of this option was called "TOP_PAD")

```

DEFAULT_TRIM_THRESHOLD    default: 2MB

```

Also settable using `mallopt(M_TRIM_THRESHOLD, x)`

The maximum amount of unused top-most memory to keep before releasing via `malloc_trim` in `free()`. Automatic trimming is mainly useful in long-lived programs using contiguous `MORECORE`. Because trimming via `sbrk` can be slow on some systems, and can sometimes be wasteful (in cases where programs immediately afterward allocate more large chunks) the value should be high enough so that your overall system performance would improve by releasing this much memory. As a rough guide, you might set to a value close to the average size of a process (program) running on your system. Releasing this much memory would allow such a process to run in memory. Generally, it is worth tuning trim thresholds when a program undergoes phases where several large chunks are allocated and released in ways that can reuse each other's storage, perhaps mixed with phases where there are no such chunks at all. The trim value must be greater than page size to have any useful effect. To disable trimming completely, you can set to `MAX_SIZE_T`. Note that the trick some people use of allocating a huge space and then freeing it at program startup, in an attempt to reserve system memory, doesn't have the intended effect under automatic trimming, since that memory will immediately be returned to the system.

```

444 DEFAULT_MMAP_THRESHOLD    default: 256K

```

Also settable using `mallopt(M_MMAP_THRESHOLD, x)`

The request size threshold for using `MMAP` to directly service a request. Requests of at least this size that cannot be allocated using already-existing space will be serviced via `mmap`. (If enough normal freed space already exists it is used instead.) Using `mmap` segregates relatively large chunks of memory so that they can be individually obtained and released from the host system. A request serviced through `mmap` is never reused by any other request (at least not directly; the system may just so happen to remap successive requests to the same locations). Segregating space in this way has the benefits that: Mmapped space can always be individually released

```

456  back to the system, which helps keep the system level memory demands
    of a long-lived program low. Also, mapped memory doesn't become
    'locked' between other chunks, as can happen with normally allocated
460  chunks, which means that even trimming via malloc_trim would not
    release them. However, it has the disadvantage that the space
    cannot be reclaimed, consolidated, and then used to service later
    requests, as happens with normal chunks. The advantages of mmap
    nearly always outweigh disadvantages for "large" chunks, but the
464  value of "large" may vary across systems. The default is an
    empirically derived value that works well in most systems. You can
    disable mmap by setting to MAX_SIZE_T.

468  MAX_RELEASE_CHECK_RATE    default: 4095 unless not HAVE_MMAP
    The number of consolidated frees between checks to release
    unused segments when freeing. When using non-contiguous segments,
    especially with multiple mspaces, checking only for topmost space
472  doesn't always suffice to trigger trimming. To compensate for this,
    free() will, with a period of MAX_RELEASE_CHECK_RATE (or the
    current number of segments, if greater) try to release unused
    segments to the OS when freeing chunks that result in
476  consolidation. The best value for this parameter is a compromise
    between slowing down frees with relatively costly checks that
    rarely trigger versus holding on to unused memory. To effectively
    disable, set to MAX_SIZE_T. This may lead to a very slight speed
480  improvement at the expense of carrying around more memory.
    */

    /* Version identifier to allow people to support multiple versions */
484  #ifndef DLMALLOC_VERSION
    #define DLMALLOC_VERSION 20804
    #endif /* DLMALLOC_VERSION */

488  #ifndef WIN32
    #ifdef _WIN32
    #define WIN32 1
    #endif /* _WIN32 */
492  #ifdef _WIN32_WCE
    #define LACKS_FCNTL_H
    #define WIN32 1
    #endif /* _WIN32_WCE */
496  #endif /* WIN32 */
    #ifdef WIN32
    #define WIN32_LEAN_AND_MEAN
    #include <windows.h>
500  #define HAVE_MMAP 1
    #define HAVE_MORECORE 0
    #define LACKS_UNISTD_H
    #define LACKS_SYS_PARAM_H
504  #define LACKS_SYS_MMAN_H
    #define LACKS_STRING_H
    #define LACKS_STRINGS_H
    #define LACKS_SYS_TYPES_H

```

```

508 #define LACKS_ERRNO_H
    #ifndef MALLOC_FAILURE_ACTION
    #define MALLOC_FAILURE_ACTION
    #endif /* MALLOC_FAILURE_ACTION */
512 #ifdef _WIN32_WCE /* WINCE reportedly does not clear */
    #define MMAP_CLEARS 0
    #else
    #define MMAP_CLEARS 1
516 #endif /* _WIN32_WCE */
    #endif /* WIN32 */

    #if defined(DARWIN) || defined(_DARWIN)
520 /* Mac OSX docs advise not to use sbrk; it seems better to use mmap */
    #ifndef HAVE_MORECORE
    #define HAVE_MORECORE 0
    #define HAVE_MMAP 1
524 /* OSX allocators provide 16 byte alignment */
    #ifndef MALLOC_ALIGNMENT
    #define MALLOC_ALIGNMENT ((size_t)16U)
    #endif
528 #endif /* HAVE_MORECORE */
    #endif /* DARWIN */

    #ifndef LACKS_SYS_TYPES_H
532 #include <sys/types.h> /* For size_t */
    #endif /* LACKS_SYS_TYPES_H */

    #if (defined(__GNUC__) && ((defined(__i386__) || defined(__x86_64__))) \
536         || (defined(_MSC_VER) && _MSC_VER>=1310))
    #define SPIN_LOCKS_AVAILABLE 1
    #else
    #define SPIN_LOCKS_AVAILABLE 0
540 #endif

    /* The maximum possible size_t value has all bits set */
    #define MAX_SIZE_T (~(size_t)0)

    #ifndef ONLY_MSPACES
    #define ONLY_MSPACES 0 /* define to a value */
    #else
548 #define ONLY_MSPACES 1
    #endif /* ONLY_MSPACES */
    #ifndef MSPACES
    #if ONLY_MSPACES
552 #define MSPACES 1
    #else /* ONLY_MSPACES */
    #define MSPACES 0
    #endif /* ONLY_MSPACES */
556 #endif /* MSPACES */
    #ifndef MALLOC_ALIGNMENT
    #define MALLOC_ALIGNMENT ((size_t)8U)
    #endif /* MALLOC_ALIGNMENT */

```



```

560 #ifndef FOOTERS
    #define FOOTERS 0
    #endif /* FOOTERS */
    #ifndef ABORT
564 #define ABORT abort()
    #endif /* ABORT */
    #ifndef ABORT_ON_ASSERT_FAILURE
    #define ABORT_ON_ASSERT_FAILURE 1
568 #endif /* ABORT_ON_ASSERT_FAILURE */
    #ifndef PROCEED_ON_ERROR
    #define PROCEED_ON_ERROR 0
    #endif /* PROCEED_ON_ERROR */
572 #ifndef USE_LOCKS
    #define USE_LOCKS 0
    #endif /* USE_LOCKS */
    #ifndef USE_SPIN_LOCKS
576 #if USE_LOCKS && SPIN_LOCKS_AVAILABLE
    #define USE_SPIN_LOCKS 1
    #else
    #define USE_SPIN_LOCKS 0
580 #endif /* USE_LOCKS && SPIN_LOCKS_AVAILABLE. */
    #endif /* USE_SPIN_LOCKS */
    #ifndef INSECURE
    #define INSECURE 0
584 #endif /* INSECURE */
    #ifndef HAVE_MMAP
    #define HAVE_MMAP 1
    #endif /* HAVE_MMAP */
588 #ifndef MMAP_CLEARS
    #define MMAP_CLEARS 1
    #endif /* MMAP_CLEARS */
    #ifndef HAVE_MREMAP
592 #ifdef linux
    #define HAVE_MREMAP 1
    #else /* linux */
    #define HAVE_MREMAP 0
596 #endif /* linux */
    #endif /* HAVE_MREMAP */
    #ifndef MALLOC_FAILURE_ACTION
    #define MALLOC_FAILURE_ACTION errno = ENOMEM;
600 #endif /* MALLOC_FAILURE_ACTION */
    #ifndef HAVE_MORECORE
    #if ONLY_MSPACES
    #define HAVE_MORECORE 0
604 #else /* ONLY_MSPACES */
    #define HAVE_MORECORE 1
    #endif /* ONLY_MSPACES */
    #endif /* HAVE_MORECORE */
608 #if !HAVE_MORECORE
    #define MORECORE_CONTIGUOUS 0
    #else /* !HAVE_MORECORE */
    #define MORECORE_DEFAULT sbrk

```

```

612 #ifndef MORECORE_CONTIGUOUS
    #define MORECORE_CONTIGUOUS 1
    #endif /* MORECORE_CONTIGUOUS */
    #endif /* HAVE_MORECORE */
616 #ifndef DEFAULT_GRANULARITY
    #if (MORECORE_CONTIGUOUS || defined(WIN32))
    #define DEFAULT_GRANULARITY (0) /* 0 means to compute in init_mparams */
    #else /* MORECORE_CONTIGUOUS */
620 #define DEFAULT_GRANULARITY ((size_t)64U * (size_t)1024U)
    #endif /* MORECORE_CONTIGUOUS */
    #endif /* DEFAULT_GRANULARITY */
    #ifndef DEFAULT_TRIM_THRESHOLD
624 #ifndef MORECORE_CANNOT_TRIM
    #define DEFAULT_TRIM_THRESHOLD ((size_t)2U * (size_t)1024U * (size_t)1024U)
    #else /* MORECORE_CANNOT_TRIM */
    #define DEFAULT_TRIM_THRESHOLD MAX_SIZE_T
628 #endif /* MORECORE_CANNOT_TRIM */
    #endif /* DEFAULT_TRIM_THRESHOLD */
    #ifndef DEFAULT_MMAP_THRESHOLD
    #if HAVE_MMAP
632 #define DEFAULT_MMAP_THRESHOLD ((size_t)256U * (size_t)1024U)
    #else /* HAVE_MMAP */
    #define DEFAULT_MMAP_THRESHOLD MAX_SIZE_T
    #endif /* HAVE_MMAP */
636 #endif /* DEFAULT_MMAP_THRESHOLD */
    #ifndef MAX_RELEASE_CHECK_RATE
    #if HAVE_MMAP
    #define MAX_RELEASE_CHECK_RATE 4095
640 #else
    #define MAX_RELEASE_CHECK_RATE MAX_SIZE_T
    #endif /* HAVE_MMAP */
    #endif /* MAX_RELEASE_CHECK_RATE */
644 #ifndef USE_BUILTIN_FFS
    #define USE_BUILTIN_FFS 0
    #endif /* USE_BUILTIN_FFS */
    #ifndef USE_DEV_RANDOM
648 #define USE_DEV_RANDOM 0
    #endif /* USE_DEV_RANDOM */
    #ifndef NO_MALLINFO
    #define NO_MALLINFO 0
652 #endif /* NO_MALLINFO */
    #ifndef MALLINFO_FIELD_TYPE
    #define MALLINFO_FIELD_TYPE size_t
    #endif /* MALLINFO_FIELD_TYPE */
656 #ifndef NO_SEGMENT_TRAVERSAL
    #define NO_SEGMENT_TRAVERSAL 0
    #endif /* NO_SEGMENT_TRAVERSAL */

```

```

660 /*
    mallopt tuning options.  SVID/XPG defines four standard parameter
    numbers for mallopt, normally defined in malloc.h.  None of these
    are used in this malloc, so setting them has no effect.  But this

```



```
664  malloc does support the following options.  
    */  
  
    #define M_TRIM_THRESHOLD      (-1)  
668  #define M_GRANULARITY          (-2)  
    #define M_MMAP_THRESHOLD      (-3)
```

Chapter 3

Mallinfo declarations

```
672 #if !NO_MALLINFO
    /*
        This version of malloc supports the standard SVID/XPG mallinfo
        routine that returns a struct containing usage properties and
676 statistics. It should work on any system that has a
        /usr/include/malloc.h defining struct mallinfo. The main
        declaration needed is the mallinfo struct that is returned (by-copy)
        by mallinfo(). The mallinfo struct contains a bunch of fields that
680 are not even meaningful in this version of malloc. These fields are
        are instead filled by mallinfo() with other numbers that might be of
        interest.

684 HAVE_USR_INCLUDE_MALLOC_H should be set if you have a
        /usr/include/malloc.h file that includes a declaration of struct
        mallinfo. If so, it is included; else a compliant version is
        declared below. These must be precisely the same for mallinfo() to
688 work. The original SVID version of this struct, defined on most
        systems with mallinfo, declares all fields as ints. But some others
        define as unsigned long. If your system defines the fields using a
        type of different width than listed here, you MUST #include your
692 system version and #define HAVE_USR_INCLUDE_MALLOC_H.
    */

    /* #define HAVE_USR_INCLUDE_MALLOC_H */

    #ifndef HAVE_USR_INCLUDE_MALLOC_H
    #include "/usr/include/malloc.h"
    #else /* HAVE_USR_INCLUDE_MALLOC_H */
700 #ifndef STRUCT_MALLINFO_DECLARED
    #define STRUCT_MALLINFO_DECLARED 1
    struct mallinfo {
        MALLINFO_FIELD_TYPE arena; /* non-mmapped space allocated from system */
704 MALLINFO_FIELD_TYPE ordblks; /* number of free chunks */
        MALLINFO_FIELD_TYPE smblks; /* always 0 */
        MALLINFO_FIELD_TYPE hblks; /* always 0 */
        MALLINFO_FIELD_TYPE hblkhd; /* space in mmapped regions */
708 MALLINFO_FIELD_TYPE usmblks; /* maximum total allocated space */
        MALLINFO_FIELD_TYPE fsmblks; /* always 0 */
```

```

    MALLINFO_FIELD_TYPE uordblks; /* total allocated space */
    MALLINFO_FIELD_TYPE ffordblks; /* total free space */
712 MALLINFO_FIELD_TYPE keepcost; /* releasable (via malloc_trim) space */
};
#endif /* STRUCT_MALLINFO_DECLARED */
#endif /* HAVE_USR_INCLUDE_MALLOC_H */
716 #endif /* NO_MALLINFO */

/*
    Try to persuade compilers to inline. The most critical functions for
720 inlining are defined as macros, so these aren't used for them.
*/

#ifndef FORCEINLINE
724 #if defined(__GNUC__)
#define FORCEINLINE __inline __attribute__((always_inline))
    #elif defined(_MSC_VER)
        #define FORCEINLINE __forceinline
728 #endif
#endif
#ifndef NOINLINE
    #if defined(__GNUC__)
732 #define NOINLINE __attribute__((noinline))
    #elif defined(_MSC_VER)
        #define NOINLINE __declspec(noinline)
    #else
736 #define NOINLINE
    #endif
#endif

740 #ifdef __cplusplus
extern "C" {
    #ifndef FORCEINLINE
        #define FORCEINLINE inline
744 #endif
    #endif /* __cplusplus */
    #ifndef FORCEINLINE
        #define FORCEINLINE
748 #endif

    #if !ONLY_MSPACES

```

Chapter 4

Declarations of public routines

```
#ifndef USE_DL_PREFIX
#define dlcalloc          calloc
#define dlfree            free
756 #define dlmalloc        malloc
#define dlmemalign        memalign
#define dlrealloc         realloc
#define dlvalloc          valloc
760 #define dlpvalloc        pvalloc
#define dlmallinfo        mallinfo
#define dlmallopt         mallopt
#define dlmalloc_trim     malloc_trim
764 #define dlmalloc_stats   malloc_stats
#define dlmalloc_usable_size malloc_usable_size
#define dlmalloc_footprint malloc_footprint
#define dlmalloc_max_footprint malloc_max_footprint
768 #define dlindependent_calloc independent_calloc
#define dlindependent_comalloc independent_comalloc
#endif /* USE_DL_PREFIX */

/*
  malloc(size_t n)
  Returns a pointer to a newly allocated chunk of at least n bytes, or
776 null if no space is available, in which case errno is set to ENOMEM
  on ANSI C systems.

  If n is zero, malloc returns a minimum-sized chunk. (The minimum
780 size is 16 bytes on most 32bit systems, and 32 bytes on 64bit
  systems.) Note that size_t is an unsigned type, so calls with
  arguments that would be negative if signed are interpreted as
  requests for huge amounts of space, which will often fail. The
784 maximum supported value of n differs across systems, but is in all
  cases less than the maximum representable value of a size_t.
  */
void* dlmalloc(size_t);

/*
  free(void* p)
```

```

    Releases the chunk of memory pointed to by p, that had been previously
792 allocated using malloc or a related routine such as realloc.
    It has no effect if p is null. If p was not malloced or already
    freed, free(p) will by default cause the current program to abort.
    */
796 void  dlfree(void*);

    /*
    calloc(size_t n_elements, size_t element_size);
800 Returns a pointer to n_elements * element_size bytes, with all locations
    set to zero.
    */
void* dlcalloc(size_t, size_t);

    /*
    realloc(void* p, size_t n)
    Returns a pointer to a chunk of size n that contains the same data
808 as does chunk p up to the minimum of (n, p's size) bytes, or null
    if no space is available.

    The returned pointer may or may not be the same as p. The algorithm
812 prefers extending p in most cases when possible, otherwise it
    employs the equivalent of a malloc-copy-free sequence.

    If p is null, realloc is equivalent to malloc.

    If space is not available, realloc returns null, errno is set (if on
    ANSI) and p is NOT freed.

820 if n is for fewer bytes than already held by p, the newly unused
    space is lopped off and freed if possible.  realloc with a size
    argument of zero (re)allocates a minimum-sized chunk.

824 The old unix realloc convention of allowing the last-free'd chunk
    to be used as an argument to realloc is not supported.
    */

828 void* dlrealloc(void*, size_t);

    /*
    memalign(size_t alignment, size_t n);
832 Returns a pointer to a newly allocated chunk of n bytes, aligned
    in accord with the alignment argument.

    The alignment argument should be a power of two. If the argument is
836 not a power of two, the nearest greater power is used.
    8-byte alignment is guaranteed by normal malloc calls, so don't
    bother calling memalign with an argument of 8 or less.

840 Overreliance on memalign is a sure way to fragment space.
    */
void* dlmemalign(size_t, size_t);

```

```

844  /*
      valloc(size_t n);
      Equivalent to memalign(pagesize, n), where pagesize is the page
      size of the system. If the pagesize is unknown, 4096 is used.
848  */
      void* dlvalloc(size_t);

      /*
852  mallopt(int parameter_number, int parameter_value)
      Sets tunable parameters The format is to provide a
      (parameter-number, parameter-value) pair. mallopt then sets the
      corresponding parameter to the argument value if it can (i.e., so
856  long as the value is meaningful), and returns 1 if successful else
      0. To workaround the fact that mallopt is specified to use int,
      not size_t parameters, the value -1 is specially treated as the
      maximum unsigned size_t value.

      SVID/XPG/ANSI defines four standard param numbers for mallopt,
      normally defined in malloc.h. None of these are use in this malloc,
      so setting them has no effect. But this malloc also supports other
864  options in mallopt. See below for details. Briefly, supported
      parameters are as follows (listed defaults are for "typical"
      configurations).

868  Symbol          param #  default    allowed param values
      M_TRIM_THRESHOLD -1    2*1024*1024  any    (-1 disables)
      M_GRANULARITY    -2      page size  any power of 2 >= page size
      M_MMAP_THRESHOLD -3      256*1024  any    (or 0 if no MMAP support)
872  */
      int dlmallopt(int, int);

      /*
876  malloc_footprint();
      Returns the number of bytes obtained from the system. The total
      number of bytes allocated by malloc, realloc etc., is less than this
      value. Unlike mallinfo, this function returns only a precomputed
880  result, so can be called frequently to monitor memory consumption.
      Even if locks are otherwise defined, this function does not use them,
      so results might not be up to date.
      */
884  size_t dlmalloc_footprint(void);

      /*
      malloc_max_footprint();
888  Returns the maximum number of bytes obtained from the system. This
      value will be greater than current footprint if deallocated space
      has been reclaimed by the system. The peak number of bytes allocated
      by malloc, realloc etc., is less than this value. Unlike mallinfo,
892  this function returns only a precomputed result, so can be called
      frequently to monitor memory consumption. Even if locks are
      otherwise defined, this function does not use them, so results might

```

```

    not be up to date.
896 */
    size_t dlmalloc_max_footprint(void);

    #if !NO_MALLINFO
900 /*
    mallinfo()
    Returns (by copy) a struct containing various summary statistics:

904 arena:      current total non-mmapped bytes allocated from system
    ordblks:   the number of free chunks
    smblks:    always zero.
    hblks:     current number of mmapped regions
908 hblkhd:     total bytes held in mmapped regions
    usmblks:   the maximum total allocated space. This will be greater
               than current total if trimming has occurred.
    fsmblks:   always zero
912 uordblks:   current total allocated space (normal or mmapped)
    fordblks:  total free space
    keepcost:  the maximum number of bytes that could ideally be released
               back to system via malloc_trim. ("ideally" means that
916             it ignores page restrictions etc.)

    Because these fields are ints, but internal bookkeeping may
    be kept as longs, the reported values may wrap around zero and
920 thus be inaccurate.
    */
    struct mallinfo dlmallinfo(void);
    #endif /* NO_MALLINFO */

    /*
    independent_calloc(size_t n_elements, size_t element_size, void* chunks[]);

928 independent_calloc is similar to calloc, but instead of returning a
    single cleared space, it returns an array of pointers to n_elements
    independent elements that can hold contents of size elem_size, each
    of which starts out cleared, and can be independently freed,
932 realloc'ed etc. The elements are guaranteed to be adjacently
    allocated (this is not guaranteed to occur with multiple callocs or
    mallocs), which may also improve cache locality in some
    applications.

    The "chunks" argument is optional (i.e., may be null, which is
    probably the most typical usage). If it is null, the returned array
    is itself dynamically allocated and should also be freed when it is
940 no longer needed. Otherwise, the chunks array must be of at least
    n_elements in length. It is filled in with the pointers to the
    chunks.

944 In either case, independent_calloc returns this pointer array, or
    null if the allocation failed. If n_elements is zero and "chunks"
    is null, it returns a chunk representing an array with zero elements

```

(which should be freed if not wanted).

Each element must be individually freed when it is no longer needed. If you'd like to instead be able to free all at once, you should instead use regular `calloc` and assign pointers into this space to represent elements. (In this case though, you cannot independently free elements.)

`independent_calloc` simplifies and speeds up implementations of many kinds of pools. It may also be useful when constructing large data structures that initially have a fixed number of fixed-sized nodes, but the number is not known at compile time, and some of the nodes may later need to be freed. For example:

```
struct Node { int item; struct Node* next; };
```

```
struct Node* build_list() {
```

```
    struct Node** pool;
```

```
    int n = read_number_of_nodes_needed();
```

```
    if (n <= 0) return 0;
```

```
    pool = (struct Node**)(independent_calloc(n, sizeof(struct Node), 0);
```

```
    if (pool == 0) die();
```

```
    // organize into a linked list...
```

```
    struct Node* first = pool[0];
```

```
    for (i = 0; i < n-1; ++i)
```

```
        pool[i]->next = pool[i+1];
```

```
    free(pool);    // Can now free the array (or not, if it is needed later)
```

```
    return first;
```

```
}
```

```
*/
```

```
void** dlindependent_calloc(size_t, size_t, void**);
```

```
/*
```

```
independent_comalloc(size_t n_elements, size_t sizes[], void* chunks[]);
```

`independent_comalloc` allocates, all at once, a set of `n_elements` chunks with sizes indicated in the "sizes" array. It returns an array of pointers to these elements, each of which can be independently freed, `realloc`'ed etc. The elements are guaranteed to be adjacently allocated (this is not guaranteed to occur with multiple `calloc`s or `malloc`s), which may also improve cache locality in some applications.

The "chunks" argument is optional (i.e., may be null). If it is null the returned array is itself dynamically allocated and should also be freed when it is no longer needed. Otherwise, the chunks array must be of at least `n_elements` in length. It is filled in with the pointers to the chunks.

In either case, `independent_comalloc` returns this pointer array, or null if the allocation failed. If `n_elements` is zero and `chunks` is null, it returns a chunk representing an array with zero elements

(which should be freed if not wanted).

Each element must be individually freed when it is no longer needed. If you'd like to instead be able to free all at once, you should instead use a single regular malloc, and assign pointers at particular offsets in the aggregate space. (In this case though, you cannot independently free elements.)

`independent_comallac` differs from `independent_calloc` in that each element may have a different size, and also that it does not automatically clear elements.

`independent_comalloc` can be used to speed up allocation in cases where several structs or objects must always be allocated at the same time. For example:

```
struct Head { ... }
struct Foot { ... }

void send_message(char* msg) {
    int msglen = strlen(msg);
    size_t sizes[3] = { sizeof(struct Head), msglen, sizeof(struct Foot) };
    void* chunks[3];
    if (independent_comalloc(3, sizes, chunks) == 0)
        die();
    struct Head* head = (struct Head*)(chunks[0]);
    char* body = (char*)(chunks[1]);
    struct Foot* foot = (struct Foot*)(chunks[2]);
    // ...
}
```

In general though, `independent_comalloc` is worth using only for larger values of `n_elements`. For small values, you probably won't detect enough difference from series of malloc calls to bother.

Overuse of `independent_comalloc` can increase overall memory usage, since it cannot reuse existing noncontiguous small chunks that might be available for some of the elements.

```
*/
void** dlindependent_comalloc(size_t, size_t*, void**);

/*
pvalloc(size_t n);
Equivalent to valloc(minimum-page-that-holds(n)), that is,
round up n to nearest pagesize.
*/
void* dlpvalloc(size_t);

/*
malloc_trim(size_t pad);
```

If possible, gives memory back to the system (via negative arguments to `sbrk`) if there is unused memory at the ‘high’ end of the `malloc` pool or in unused `MMAP` segments. You can call this after freeing large blocks of memory to potentially reduce the system-level memory requirements of a program. However, it cannot guarantee to reduce memory. Under some allocation patterns, some large free blocks of memory will be locked between two used chunks, so they cannot be given back to the system.

The ‘`pad`’ argument to `malloc_trim` represents the amount of free trailing space to leave untrimmed. If this argument is zero, only the minimum amount of memory to maintain internal data structures will be left. Non-zero arguments can be supplied to maintain enough trailing space to service future expected allocations without having to re-obtain memory from the system.

`Malloc_trim` returns 1 if it actually released any memory, else 0.

```
1068 */
int  dlmalloc_trim(size_t);
```

```
1072 /*
malloc_stats();
Prints on stderr the amount of space obtained from the system (both
via sbrk and mmap), the maximum amount (which may be more than
current if malloc_trim and/or munmap got called), and the current
1076 number of bytes allocated via malloc (or realloc, etc) but not yet
freed. Note that this is the number of bytes allocated, not the
number requested. It will be larger than the number requested
because of alignment and bookkeeping overhead. Because it includes
1080 alignment wastage as being in use, this figure may be greater than
zero even when no user-level chunks are allocated.
```

The reported current and maximum system memory can be inaccurate if a program makes other calls to system memory allocation functions (normally `sbrk`) outside of `malloc`.

`malloc_stats` prints only the most commonly interesting statistics. More information can be obtained by calling `mallinfo`.

```
1088 */
void  dlmalloc_stats(void);
```

```
1092 #endif /* ONLY_MSPACES */
```

```
/*
malloc_usable_size(void* p);
```

Returns the number of bytes you can actually use in an allocated chunk, which may be more than you requested (although often not) due to alignment and minimum size constraints. You can use this many bytes without worrying about overwriting other allocated objects. This is not a particularly great programming practice. `malloc_usable_size` can be more useful in

```

    debugging and assertions, for example:

    p = malloc(n);
    assert(malloc_usable_size(p) >= 256);
    */
1108 size_t dlmalloc_usable_size(void*);

#ifdef MSPACES

    /*
     mspace is an opaque type representing an independent
     region of space that supports mspace_malloc, etc.
1116 */
    typedef void* mspace;

    /*
1120 create_mspace creates and returns a new independent space with the
     given initial capacity, or, if 0, the default granularity size. It
     returns null if there is no system memory available to create the
     space. If argument locked is non-zero, the space uses a separate
1124 lock to control access. The capacity of the space will grow
     dynamically as needed to service mspace_malloc requests. You can
     control the sizes of incremental increases of this space by
     compiling with a different DEFAULT_GRANULARITY or dynamically
1128 setting with mallopt(M_GRANULARITY, value).
    */
    mspace create_mspace(size_t capacity, int locked);

1132 /*
     destroy_mspace destroys the given space, and attempts to return all
     of its memory back to the system, returning the total number of
     bytes freed. After destruction, the results of access to all memory
1136 used by the space become undefined.
    */
    size_t destroy_mspace(mspace msp);

1140 /*
     create_mspace_with_base uses the memory supplied as the initial base
     of a new mspace. Part (less than 128*sizeof(size_t) bytes) of this
     space is used for bookkeeping, so the capacity must be at least this
1144 large. (Otherwise 0 is returned.) When this initial space is
     exhausted, additional memory will be obtained from the system.
     Destroying this space will deallocate all additionally allocated
     space (if possible) but not the initial base.
1148 */
    mspace create_mspace_with_base(void* base, size_t capacity, int locked);

    /*
1152 mspace_track_large_chunks controls whether requests for large chunks
     are allocated in their own untracked mmaped regions, separate from
     others in this mspace. By default large chunks are not tracked,

```

```
which reduces fragmentation. However, such chunks are not
1156 necessarily released to the system upon destroy_mspace. Enabling
tracking by setting to true may increase fragmentation, but avoids
leakage when relying on destroy_mspace to release all memory
allocated using this space. The function returns the previous
1160 setting.
*/
int mspace_track_large_chunks(mspace msp, int enable);

/*
    mspace_malloc behaves as malloc, but operates within
    the given space.
1168 */
void* mspace_malloc(mspace msp, size_t bytes);

/*
1172 mspace_free behaves as free, but operates within
the given space.

    If compiled with FOOTERS==1, mspace_free is not actually needed.
1176 free may be called instead of mspace_free because freed chunks from
any space are handled by their originating spaces.
*/
void mspace_free(mspace msp, void* mem);

/*
    mspace_realloc behaves as realloc, but operates within
    the given space.

    If compiled with FOOTERS==1, mspace_realloc is not actually
    needed. realloc may be called instead of mspace_realloc because
    reallocated chunks from any space are handled by their originating
1188 spaces.
*/
void* mspace_realloc(mspace msp, void* mem, size_t newsize);

1192 /*
    mspace_calloc behaves as calloc, but operates within
    the given space.
*/
1196 void* mspace_calloc(mspace msp, size_t n_elements, size_t elem_size);

/*
    mspace_memalign behaves as memalign, but operates within
1200 the given space.
*/
void* mspace_memalign(mspace msp, size_t alignment, size_t bytes);

1204 /*
    mspace_independent_calloc behaves as independent_calloc, but
    operates within the given space.
```

```

    */
1208 void** mspace_independent_calloc(mspace msp, size_t n_elements,
                                   size_t elem_size, void* chunks[]);

    /*
1212 mspace_independent_comalloc behaves as independent_comalloc, but
    operates within the given space.
    */
void** mspace_independent_comalloc(mspace msp, size_t n_elements,
1216 size_t sizes[], void* chunks[]);

    /*
    mspace_footprint() returns the number of bytes obtained from the
1220 system for this space.
    */
size_t mspace_footprint(mspace msp);

1224 /*
    mspace_max_footprint() returns the peak number of bytes obtained from the
    system for this space.
    */
1228 size_t mspace_max_footprint(mspace msp);

    #if !NO_MALLINFO
1232 /*
    mspace_mallinfo behaves as mallinfo, but reports properties of
    the given space.
    */
1236 struct mallinfo mspace_mallinfo(mspace msp);
    #endif /* NO_MALLINFO */

    /*
1240 malloc_usable_size(void* p) behaves the same as malloc_usable_size;
    */
size_t mspace_usable_size(void* mem);

1244 /*
    mspace_malloc_stats behaves as malloc_stats, but reports
    properties of the given space.
    */
1248 void mspace_malloc_stats(mspace msp);

    /*
    mspace_trim behaves as malloc_trim, but
1252 operates within the given space.
    */
int mspace_trim(mspace msp, size_t pad);

1256 /*
    An alias for mallopt.
    */

```

```
int mspace_mallopt(int, int);

#endif /* MSPACES */

#ifdef __cplusplus
1264 }; /* end of extern "C" */
#endif /* __cplusplus */

/*
1268 =====
    To make a fully customizable malloc.h header file, cut everything
    above this line, put into file malloc.h, edit to suit, and #include it
    on the next line, as well as in programs that use this malloc.
1272 =====
*/

/* #include "malloc.h" */
```

Chapter 5

Internal #includes

```
1276 #ifdef WIN32
    #pragma warning( disable : 4146 ) /* no "unsigned" warnings */
    #endif /* WIN32 */

1280 #include <stdio.h>          /* for printing in malloc_stats */

    #ifndef LACKS_ERRNO_H
    #include <errno.h>         /* for MALLOC_FAILURE_ACTION */
1284 #endif /* LACKS_ERRNO_H */
    #if FOOTERS || DEBUG
    #include <time.h>          /* for magic initialization */
    #endif /* FOOTERS */
1288 #ifndef LACKS_STDLIB_H
    #include <stdlib.h>        /* for abort() */
    #endif /* LACKS_STDLIB_H */
    #ifdef DEBUG
1292 #if ABORT_ON_ASSERT_FAILURE
    #undef assert
    #define assert(x) if(!(x)) ABORT
    #else /* ABORT_ON_ASSERT_FAILURE */
1296 #include <assert.h>
    #endif /* ABORT_ON_ASSERT_FAILURE */
    #else /* DEBUG */
    #ifndef assert
1300 #define assert(x)
    #endif
    #define DEBUG 0
    #endif /* DEBUG */
1304 #ifndef LACKS_STRING_H
    #include <string.h>        /* for memset etc */
    #endif /* LACKS_STRING_H */
    #if USE_BUILTIN_FFS
1308 #ifndef LACKS_STRINGS_H
    #include <strings.h>       /* for ffs */
    #endif /* LACKS_STRINGS_H */
    #endif /* USE_BUILTIN_FFS */
1312 #if HAVE_MMAP
    #ifndef LACKS_SYS_MMAN_H
        /* On some versions of linux, mremap decl in mman.h needs __USE_GNU set */
```

```

    #if (defined(linux) && !defined(__USE_GNU))
1316 #define __USE_GNU 1
    #include <sys/mman.h>      /* for mmap */
    #undef __USE_GNU
    #else
1320 #include <sys/mman.h>      /* for mmap */
    #endif /* linux */
    #endif /* LACKS_SYS_MMAN_H */
    #ifndef LACKS_FCNTL_H
1324 #include <fcntl.h>
    #endif /* LACKS_FCNTL_H */
    #endif /* HAVE_MMAP */
    #ifndef LACKS_UNISTD_H
1328 #include <unistd.h>      /* for sbrk, sysconf */
    #else /* LACKS_UNISTD_H */
    #if !defined(__FreeBSD__) && !defined(__OpenBSD__) && !defined(__NetBSD__)
    extern void*      sbrk(ptrdiff_t);
1332 #endif /* FreeBSD etc */
    #endif /* LACKS_UNISTD_H */

    /* Declarations for locking */
1336 #if USE_LOCKS
    #ifndef WIN32
    #include <pthread.h>
    #if defined (__SVR4) && defined (__sun) /* solaris */
1340 #include <thread.h>
    #endif /* solaris */
    #else
    #ifndef _M_AMD64
1344 /* These are already defined on AMD64 builds */
    #ifdef __cplusplus
    extern "C" {
    #endif /* __cplusplus */
1348 LONG __cdecl _InterlockedCompareExchange(LONG volatile *Dest, \
                                           LONG Exchange, LONG Comp);

    LONG __cdecl _InterlockedExchange(LONG volatile *Target, LONG Value);
    #ifdef __cplusplus
1352 }
    #endif /* __cplusplus */
    #endif /* _M_AMD64 */
    #pragma intrinsic (_InterlockedCompareExchange)
1356 #pragma intrinsic (_InterlockedExchange)
    #define interlockedcompareexchange _InterlockedCompareExchange
    #define interlockedexchange _InterlockedExchange
    #endif /* Win32 */
1360 #endif /* USE_LOCKS */

    /* Declarations for bit scanning on win32 */
    #if defined(_MSC_VER) && _MSC_VER>=1300
1364 #ifndef BitScanForward /* Try to avoid pulling in WinNT.h */
    #ifdef __cplusplus
    extern "C" {

```



```
1420 #           endif
      #           endif
      #           endif
      #           endif
      #           endif
1424 # endif
      #endif
      #endif
```

Chapter 6

size_t and alignment properties

```
/* The byte and bit size of a size_t */
1428 #define SIZE_T_SIZE      (sizeof(size_t))
    #define SIZE_T_BITSIZE  (sizeof(size_t) << 3)

/* Some constants coerced to size_t */
1432 /* Annoying but necessary to avoid errors on some platforms */
    #define SIZE_T_ZERO      ((size_t)0)
    #define SIZE_T_ONE       ((size_t)1)
    #define SIZE_T_TWO       ((size_t)2)
1436 #define SIZE_T_FOUR       ((size_t)4)
    #define TWO_SIZE_T_SIZES (SIZE_T_SIZE<<1)
    #define FOUR_SIZE_T_SIZES (SIZE_T_SIZE<<2)
    #define SIX_SIZE_T_SIZES (FOUR_SIZE_T_SIZES+TWO_SIZE_T_SIZES)
1440 #define HALF_MAX_SIZE_T  (MAX_SIZE_T / 2U)

/* The bit mask value corresponding to MALLOC_ALIGNMENT */
#define CHUNK_ALIGN_MASK  (MALLOC_ALIGNMENT - SIZE_T_ONE)

/* True if address a has acceptable alignment */
#define is_aligned(A)      (((size_t)((A)) & (CHUNK_ALIGN_MASK)) == 0)

1448 /* the number of bytes to offset an address to align it */
#define align_offset(A)\
    (((size_t)(A) & CHUNK_ALIGN_MASK) == 0)? 0 :\
    ((MALLOC_ALIGNMENT - ((size_t)(A) & CHUNK_ALIGN_MASK)) & CHUNK_ALIGN_MASK)
```

Chapter 7

MMAP preliminaries

```
1452  /*
        If HAVE_MORECORE or HAVE_MMAP are false, we just define calls and
        checks to fail so compiler optimizer can delete code rather than
        using so many "#if"s.
1456  */

        /* MORECORE and MMAP must return MFAIL on failure */
1460  #define MFAIL                ((void*)(MAX_SIZE_T))
        #define CMFAIL          ((char*)(MFAIL)) /* defined for convenience */

        #if HAVE_MMAP

        #ifndef WIN32
        #define MUNMAP_DEFAULT(a, s)  munmap((a), (s))
        #define MMAP_PROT             (PROT_READ|PROT_WRITE)
1468  #if !defined(MAP_ANONYMOUS) && defined(MAP_ANON)
        #define MAP_ANONYMOUS        MAP_ANON
        #endif /* MAP_ANON */
        #ifdef MAP_ANONYMOUS
1472  #define MMAP_FLAGS              (MAP_PRIVATE|MAP_ANONYMOUS)
        #define MMAP_DEFAULT(s)      mmap(0, (s), MMAP_PROT, MMAP_FLAGS, -1, 0)
        #else /* MAP_ANONYMOUS */
        /*
1476  Nearly all versions of mmap support MAP_ANONYMOUS, so the following
        is unlikely to be needed, but is supplied just in case.
        */
        #define MMAP_FLAGS          (MAP_PRIVATE)
1480  static int dev_zero_fd = -1; /* Cached file descriptor for /dev/zero. */
        #define MMAP_DEFAULT(s) ((dev_zero_fd < 0) ? \
            (dev_zero_fd = open("/dev/zero", O_RDWR), \
             mmap(0, (s), MMAP_PROT, MMAP_FLAGS, dev_zero_fd, 0)) : \
1484             mmap(0, (s), MMAP_PROT, MMAP_FLAGS, dev_zero_fd, 0))
        #endif /* MAP_ANONYMOUS */

        #define DIRECT_MMAP_DEFAULT(s) MMAP_DEFAULT(s)

        #else /* WIN32 */
```

```

/* Win32 MMAP via VirtualAlloc */
1492 static FORCEINLINE void* win32mmap(size_t size) {
    void* ptr = VirtualAlloc(0, size, MEM_RESERVE|MEM_COMMIT, PAGE_READWRITE);
    return (ptr != 0)? ptr: MFAIL;
}

/* For direct MMAP, use MEM_TOP_DOWN to minimize interference */
static FORCEINLINE void* win32direct_mmap(size_t size) {
    void* ptr = VirtualAlloc(0, size, MEM_RESERVE|MEM_COMMIT|MEM_TOP_DOWN,
1500     PAGE_READWRITE);
    return (ptr != 0)? ptr: MFAIL;
}

1504 /* This function supports releasing coalesed segments */
static FORCEINLINE int win32munmap(void* ptr, size_t size) {
    MEMORY_BASIC_INFORMATION minfo;
    char* cptr = (char*)ptr;
1508     while (size) {
        if (VirtualQuery(cptr, &minfo, sizeof(minfo)) == 0)
            return -1;
        if (minfo.BaseAddress != cptr || minfo.AllocationBase != cptr ||
1512         minfo.State != MEM_COMMIT || minfo.RegionSize > size)
            return -1;
        if (VirtualFree(cptr, 0, MEM_RELEASE) == 0)
            return -1;
1516         cptr += minfo.RegionSize;
        size -= minfo.RegionSize;
    }
    return 0;
1520 }

#define MMAP_DEFAULT(s)          win32mmap(s)
#define MUNMAP_DEFAULT(a, s)     win32munmap((a), (s))
1524 #define DIRECT_MMAP_DEFAULT(s)  win32direct_mmap(s)
#endif /* WIN32 */
#endif /* HAVE_MMAP */

1528 #if HAVE_MREMAP
#ifndef WIN32
#define MREMAP_DEFAULT(addr, osz, nsz, mv) mremap((addr), (osz), (nsz), (mv))
#endif /* WIN32 */
1532 #endif /* HAVE_MREMAP */

/**
1536  * Define CALL_MORECORE
  */
#if HAVE_MORECORE
    #ifdef MORECORE
1540         #define CALL_MORECORE(S)    MORECORE(S)
    #else /* MORECORE */
        #define CALL_MORECORE(S)    MORECORE_DEFAULT(S)
    
```

```

        #endif /* MORECORE */
1544 #else /* HAVE_MORECORE */
        #define CALL_MORECORE(S)      MFAIL
    #endif /* HAVE_MORECORE */

1548 /**
    * Define CALL_MMAP/CALL_MUNMAP/CALL_DIRECT_MMAP
    */
    #if HAVE_MMAP
1552        #define USE_MMAP_BIT          (SIZE_T_ONE)

        #ifdef MMAP
            #define CALL_MMAP(s)        MMAP(s)
1556 #else /* MMAP */
            #define CALL_MMAP(s)        MMAP_DEFAULT(s)
        #endif /* MMAP */
        #ifdef MUNMAP
1560            #define CALL_MUNMAP(a, s)   MUNMAP((a), (s))
        #else /* MUNMAP */
            #define CALL_MUNMAP(a, s)   MUNMAP_DEFAULT((a), (s))
        #endif /* MUNMAP */
1564 #ifdef DIRECT_MMAP
            #define CALL_DIRECT_MMAP(s)  DIRECT_MMAP(s)
        #else /* DIRECT_MMAP */
            #define CALL_DIRECT_MMAP(s)  DIRECT_MMAP_DEFAULT(s)
1568 #endif /* DIRECT_MMAP */
    #else /* HAVE_MMAP */
        #define USE_MMAP_BIT          (SIZE_T_ZERO)

1572        #define MMAP(s)                MFAIL
        #define MUNMAP(a, s)            (-1)
        #define DIRECT_MMAP(s)          MFAIL
        #define CALL_DIRECT_MMAP(s)      DIRECT_MMAP(s)
1576        #define CALL_MMAP(s)            MMAP(s)
        #define CALL_MUNMAP(a, s)        MUNMAP((a), (s))
    #endif /* HAVE_MMAP */

1580 /**
    * Define CALL_MREMAP
    */
    #if HAVE_MMAP && HAVE_MREMAP
1584        #ifdef MREMAP
            #define CALL_MREMAP(addr, osz, nsz, mv) MREMAP((addr), (osz), (nsz), (mv))
        #else /* MREMAP */
            #define CALL_MREMAP(addr, osz, nsz, mv) \
1588                MREMAP_DEFAULT((addr), (osz), (nsz), (mv))
        #endif /* MREMAP */
    #else /* HAVE_MMAP && HAVE_MREMAP */
        #define CALL_MREMAP(addr, osz, nsz, mv)      MFAIL
1592 #endif /* HAVE_MMAP && HAVE_MREMAP */

    /* mstate bit set if contiguous morecore disabled or failed */

```

```
#define USE_NONCONTIGUOUS_BIT (4U)

/* segment bit set in create_mspace_with_base */
#define EXTERN_BIT              (8U)
```

Chapter 8

Locks

8.1 Lock preliminaries

```
1600  /*
      When locks are defined, there is one global lock, plus
      one per-mspace lock.

1604  The global lock_ensures that mparams.magic and other unique
      mparams values are initialized only once. It also protects
      sequences of calls to MORECORE. In many cases sys_alloc requires
1608  two calls, that should not be interleaved with calls by other
      threads. This does not protect against direct calls to MORECORE
      by other threads not using this lock, so there is still code to
      cope the best we can on interference.

1612  Per-mspace locks surround calls to malloc, free, etc. To enable use
      in layered extensions, per-mspace locks are reentrant.

      Because lock-protected regions generally have bounded times, it is
1616  OK to use the supplied simple spinlocks in the custom versions for
      x86. Spinlocks are likely to improve performance for lightly
      contended applications, but worsen performance under heavy
      contention.

      If USE_LOCKS is > 1, the definitions of lock routines here are
      bypassed, in which case you will need to define the type MLOCK_T,
      and at least INITIAL_LOCK, ACQUIRE_LOCK, RELEASE_LOCK and possibly
1624  TRY_LOCK (which is not used in this malloc, but commonly needed in
      extensions.) You must also declare a
          static MLOCK_T malloc_global_mutex = { initialization values };.

1628  */

      #if USE_LOCKS == 1

1632  #if USE_SPIN_LOCKS && SPIN_LOCKS_AVAILABLE
      #ifndef WIN32

      /* Custom pthread-style spin locks on x86 and x64 for gcc */
```



```

1636 struct pthread_mlock_t {
        volatile unsigned int l;
        unsigned int c;
        pthread_t threadid;
1640 };
#define MLOCK_T          struct pthread_mlock_t
#define CURRENT_THREAD   pthread_self()
#define INITIAL_LOCK(sl) ((sl)->threadid = 0, (sl)->l = (sl)->c = 0, 0)
1644 #define ACQUIRE_LOCK(sl) pthread_acquire_lock(sl)
#define RELEASE_LOCK(sl) pthread_release_lock(sl)
#define TRY_LOCK(sl)      pthread_try_lock(sl)
#define SPINS_PER_YIELD   63

static MLOCK_T malloc_global_mutex = { 0, 0, 0};

static FORCEINLINE int pthread_acquire_lock (MLOCK_T *sl) {
1652     int spins = 0;
    volatile unsigned int* lp = &sl->l;
    for (;;) {
        if (*lp != 0) {
1656             if (sl->threadid == CURRENT_THREAD) {
                    ++sl->c;
                    return 0;
            }
1660         }
        else {
            /* place args to cmpxchgl in locals to evade oddities in some gccs */
            int cmp = 0;
1664             int val = 1;
            int ret;
            __asm__ __volatile__ ("lock; cmpxchgl %1, %2"
                                   : "=a" (ret)
1668                                   : "r" (val), "m" (*(lp)), "0" (cmp)
                                   : "memory", "cc");

            if (!ret) {
                assert(!sl->threadid);
1672                 sl->threadid = CURRENT_THREAD;
                sl->c = 1;
                return 0;
            }
1676         }
        if ((++spins & SPINS_PER_YIELD) == 0) {
            #if defined (__SVR4) && defined (__sun) /* solaris */
                thr_yield();
1680            #else
            #if defined (__linux__) || defined (__FreeBSD__) || defined (__APPLE__)
                sched_yield();
            #else /* no-op yield on unknown systems */
1684                ;
            #endif /* __linux__ || __FreeBSD__ || __APPLE__ */
            #endif /* solaris */
        }
    }
}

```

```

1688     }
    }

static FORCEINLINE void pthread_release_lock (MLOCK_T *sl) {
1692     volatile unsigned int* lp = &sl->l;
    assert(*lp != 0);
    assert(sl->threadid == CURRENT_THREAD);
    if (--sl->c == 0) {
1696         sl->threadid = 0;
        int prev = 0;
        int ret;
        __asm__ __volatile__ ("lock; xchgl %0, %1"
1700                                : "=r" (ret)
                                : "m" (*(lp)), "0"(prev)
                                : "memory");
    }
1704 }

static FORCEINLINE int pthread_try_lock (MLOCK_T *sl) {
    volatile unsigned int* lp = &sl->l;
1708     if (*lp != 0) {
        if (sl->threadid == CURRENT_THREAD) {
            ++sl->c;
            return 1;
1712         }
    }
    else {
        int cmp = 0;
        int val = 1;
1716         int ret;
        __asm__ __volatile__ ("lock; cmpxchgl %1, %2"
                                : "=a" (ret)
                                : "r" (val), "m" (*(lp)), "0"(cmp)
                                : "memory", "cc");

        if (!ret) {
            assert(!sl->threadid);
1724             sl->threadid = CURRENT_THREAD;
            sl->c = 1;
            return 1;
        }
1728     }
    return 0;
}

#else /* WIN32 */
/* Custom win32-style spin locks on x86 and x64 for MSC */
struct win32_mlock_t {
1736     volatile long l;
    unsigned int c;
    long threadid;
};

```

```

#define MLOCK_T          struct win32_mlock_t
#define CURRENT_THREAD   GetCurrentThreadId()
#define INITIAL_LOCK(sl) ((sl)->threadid = 0, (sl)->l = (sl)->c = 0, 0)
1744 #define ACQUIRE_LOCK(sl) win32_acquire_lock(sl)
#define RELEASE_LOCK(sl) win32_release_lock(sl)
#define TRY_LOCK(sl)      win32_try_lock(sl)
#define SPINS_PER_YIELD   63

static MLOCK_T malloc_global_mutex = { 0, 0, 0};

static FORCEINLINE int win32_acquire_lock (MLOCK_T *sl) {
1752   int spins = 0;
   for (;;) {
       if (sl->l != 0) {
           if (sl->threadid == CURRENT_THREAD) {
1756               ++sl->c;
               return 0;
           }
       }
       else {
1760           if (!interlockedexchange(&sl->l, 1)) {
               assert(!sl->threadid);
               sl->threadid = CURRENT_THREAD;
1764               sl->c = 1;
               return 0;
           }
       }
       if ((++spins & SPINS_PER_YIELD) == 0)
1768           SleepEx(0, FALSE);
   }
}

static FORCEINLINE void win32_release_lock (MLOCK_T *sl) {
   assert(sl->threadid == CURRENT_THREAD);
   assert(sl->l != 0);
1776   if (--sl->c == 0) {
       sl->threadid = 0;
       interlockedexchange (&sl->l, 0);
   }
1780 }

static FORCEINLINE int win32_try_lock (MLOCK_T *sl) {
   if (sl->l != 0) {
1784       if (sl->threadid == CURRENT_THREAD) {
           ++sl->c;
           return 1;
       }
   }
1788   else {
       if (!interlockedexchange(&sl->l, 1)){
           assert(!sl->threadid);

```

```

1792     sl->threadid = CURRENT_THREAD;
        sl->c = 1;
        return 1;
    }
1796 }
    return 0;
}

1800 #endif /* WIN32 */
    #else /* USE_SPIN_LOCKS */

        #ifndef WIN32
1804 /* pthreads-based locks */

            #define MLOCK_T          pthread_mutex_t
            #define CURRENT_THREAD    pthread_self()
1808 #define INITIAL_LOCK(sl)          pthread_init_lock(sl)
            #define ACQUIRE_LOCK(sl) pthread_mutex_lock(sl)
            #define RELEASE_LOCK(sl)  pthread_mutex_unlock(sl)
            #define TRY_LOCK(sl)      (!pthread_mutex_trylock(sl))

            static MLOCK_T malloc_global_mutex = PTHREAD_MUTEX_INITIALIZER;

            /* Cope with old-style linux recursive lock initialization by adding */
            /* skipped internal declaration from pthread.h */
            #ifdef linux
            #ifndef PTHREAD_MUTEX_RECURSIVE
            extern int pthread_mutexattr_setkind_np __P ((pthread_mutexattr_t *__attr,
1820                                     int __kind));
            #define PTHREAD_MUTEX_RECURSIVE PTHREAD_MUTEX_RECURSIVE_NP
            #define pthread_mutexattr_settype(x,y) pthread_mutexattr_setkind_np(x,y)
            #endif
1824 #endif

            static int pthread_init_lock (MLOCK_T *sl) {
                pthread_mutexattr_t attr;
1828     if (pthread_mutexattr_init(&attr)) return 1;
                if (pthread_mutexattr_settype(&attr, PTHREAD_MUTEX_RECURSIVE)) return 1;
                if (pthread_mutex_init(sl, &attr)) return 1;
                if (pthread_mutexattr_destroy(&attr)) return 1;
1832     return 0;
            }

            #else /* WIN32 */
1836 /* Win32 critical sections */
            #define MLOCK_T          CRITICAL_SECTION
            #define CURRENT_THREAD    GetCurrentThreadId()
            #define INITIAL_LOCK(s)   (!InitializeCriticalSectionAndSpinCount((s), \
1840                                     0x80000000|4000))
            #define ACQUIRE_LOCK(s) (EnterCriticalSection(sl), 0)
            #define RELEASE_LOCK(s)  LeaveCriticalSection(sl)
            #define TRY_LOCK(s)      TryEnterCriticalSection(sl)

```

```

1844 #define NEED_GLOBAL_LOCK_INIT

static MLOCK_T malloc_global_mutex;
static volatile long malloc_global_mutex_status;

/* Use spin loop to initialize global lock */
static void init_malloc_global_mutex() {
    for (;;) {
1852     long stat = malloc_global_mutex_status;
        if (stat > 0)
            return;
        /* transition to < 0 while initializing, then to > 0) */
1856     if (stat == 0 &&
        interlockedcompareexchange(&malloc_global_mutex_status, -1, 0) == 0) {
        InitializeCriticalSection(&malloc_global_mutex);
        interlockedexchange(&malloc_global_mutex_status, 1);
1860     return;
    }
    SleepEx(0, FALSE);
}
1864 }

#endif /* WIN32 */
#endif /* USE_SPIN_LOCKS */
1868 #endif /* USE_LOCKS == 1 */

```

8.2 User-defined locks

```

1872 #if USE_LOCKS > 1
/* Define your own lock implementation here */
/* #define INITIAL_LOCK(sl) ... */
/* #define ACQUIRE_LOCK(sl) ... */
1876 /* #define RELEASE_LOCK(sl) ... */
/* #define TRY_LOCK(sl) ... */
/* static MLOCK_T malloc_global_mutex = ... */
#endif /* USE_LOCKS > 1 */

```

8.3 Lock-based state

```

#if USE_LOCKS
1884 #define USE_LOCK_BIT (2U)
#else /* USE_LOCKS */
#define USE_LOCK_BIT (0U)
#define INITIAL_LOCK(1)
1888 #endif /* USE_LOCKS */

#if USE_LOCKS
#ifndef ACQUIRE_MALLOC_GLOBAL_LOCK

```

```
1892 #define ACQUIRE_MALLOC_GLOBAL_LOCK()  ACQUIRE_LOCK(&malloc_global_mutex);  
    #endif  
    #ifndef RELEASE_MALLOC_GLOBAL_LOCK  
    #define RELEASE_MALLOC_GLOBAL_LOCK()  RELEASE_LOCK(&malloc_global_mutex);  
1896 #endif  
    #else /* USE_LOCKS */  
    #define ACQUIRE_MALLOC_GLOBAL_LOCK()  
    #define RELEASE_MALLOC_GLOBAL_LOCK()  
1900 #endif /* USE_LOCKS */
```

Chapter 9

Chunks

9.1 Chunk representations

```
/*
  (The following includes lightly edited explanations by Colin Plumb.)
```

The `malloc_chunk` declaration below is misleading (but accurate and necessary). It declares a "view" into memory allowing access to necessary fields at known offsets from a given base.

Chunks of memory are maintained using a ‘boundary tag’ method as originally described by Knuth. (See the paper by Paul Wilson <ftp://ftp.cs.utexas.edu/pub/garbage/allocsrv.ps> for a survey of such techniques.) Sizes of free chunks are stored both in the front of each chunk and at the end. This makes consolidating fragmented chunks into bigger chunks fast. The head fields also hold bits representing whether chunks are free or in use.

Here are some pictures to make it clearer. They are "exploded" to show that the state of a chunk can be thought of as extending from the high 31 bits of the head field of its header through the prev_foot and PINUSE_BIT bit of the following chunk header.

A chunk that's in use looks like:

[illegible]

```

      |
      +-+-+-+-+-+-+-+-+
1940  +-+-+-+-+-+-+-+-+ |1|
      | Size of next chunk (may or may not be in use) | +-+
mem-> +-+-+-+-+-+-+-+-+

```

1944 And if it's free, it looks like this:

```

chunk-> +-
      | User payload (must be in use, or we would have merged!) |
1948  +-+-+-+-+-+-+-+-+
      +-+-+-+-+-+-+-+-+ |P|
      | Size of this chunk | 0 | +-+
mem-> +-+-+-+-+-+-+-+-+
1952  | Next pointer |
      +-+-+-+-+-+-+-+-+
      | Prev pointer |
      +-+-+-+-+-+-+-+-+
1956  |
      +- size - sizeof(struct chunk) unused bytes -+
      :
chunk-> +-+-+-+-+-+-+-+-+
1960  | Size of this chunk |
      +-+-+-+-+-+-+-+-+
      +-+-+-+-+-+-+-+-+ |0|
      | Size of next chunk (must be in use, or we would have merged)| +-+
1964 mem-> +-+-+-+-+-+-+-+-+
      |
      +- User payload -+
      :
1968  +-+-+-+-+-+-+-+-+
                                     |0|
                                     +-+

```

1972 Note that since we always merge adjacent free chunks, the chunks adjacent to a free chunk must be in use.

1976 Given a pointer to a chunk (which can be derived trivially from the payload pointer) we can, in $O(1)$ time, find out whether the adjacent chunks are free, and if so, unlink them from the lists that they are on and merge them with the current chunk.

1980 Chunks always begin on even word boundaries, so the mem portion (which is returned to the user) is also on an even word boundary, and thus at least double-word aligned.

1984 The P (PINUSE_BIT) bit, stored in the unused low-order bit of the chunk size (which is always a multiple of two words), is an in-use bit for the *previous* chunk. If that bit is *clear*, then the word before the current chunk size contains the previous chunk size, and can be used to find the front of the previous chunk.

1988 The very first chunk allocated always has this bit set, preventing access to non-existent (or non-owned) memory. If pinuse is set for

any given chunk, then you CANNOT determine the size of the previous chunk, and might even get a memory addressing fault when trying to do so.

The C (CINUSE_BIT) bit, stored in the unused second-lowest bit of the chunk size redundantly records whether the current chunk is inuse (unless the chunk is mmaped). This redundancy enables usage checks within free and realloc, and reduces indirection when freeing and consolidating chunks.

Each freshly allocated chunk must have both cinuse and pinuse set. That is, each allocated chunk borders either a previously allocated and still in-use chunk, or the base of its memory arena. This is ensured by making all allocations from the the 'lowest' part of any found chunk. Further, no free chunk physically borders another one, so each free chunk is known to be preceded and followed by either inuse chunks or the ends of memory.

Note that the 'foot' of the current chunk is actually represented as the prev_foot of the NEXT chunk. This makes it easier to deal with alignments etc but can be very confusing when trying to extend or adapt this code.

The exceptions to all this are

1. The special chunk 'top' is the top-most available chunk (i.e., the one bordering the end of available memory). It is treated specially. Top is never included in any bin, is used only if no other chunk is available, and is released back to the system if it is very large (see M_TRIM_THRESHOLD). In effect, the top chunk is treated as larger (and thus less well fitting) than any other available chunk. The top chunk doesn't update its trailing size field since there is no next contiguous chunk that would have to index off it. However, space is still allocated for it (TOP_FOOT_SIZE) to enable separation or merging when space is extended.

3. Chunks allocated via mmap, have both cinuse and pinuse bits cleared in their head fields. Because they are allocated one-by-one, each must carry its own prev_foot field, which is also used to hold the offset this chunk has within its mmaped region, which is needed to preserve alignment. Each mmaped chunk is trailed by the first two fields of a fake next-chunk for sake of usage checks.

*/

```
struct malloc_chunk {
    size_t      prev_foot; /* Size of previous chunk (if free). */
    size_t      head;      /* Size and inuse bits. */
    struct malloc_chunk* fd; /* double links -- used only if free. */
    struct malloc_chunk* bk;
```

```
};
```

```
2044 typedef struct malloc_chunk  mchunk;
      typedef struct malloc_chunk* mchunkptr;
      typedef struct malloc_chunk* sbinptr; /* The type of bins of chunks */
      typedef unsigned int bindex_t;        /* Described below */
2048 typedef unsigned int binmap_t;          /* Described below */
      typedef unsigned int flag_t;          /* The type of various bit flag sets */
```

9.2 Chunks sizes and alignments

```

#define MCHUNK_SIZE      (sizeof(mchunk))

#if FOOTERS
2056 #define CHUNK_OVERHEAD      (TWO_SIZE_T_SIZES)
#else /* FOOTERS */
#define CHUNK_OVERHEAD      (SIZE_T_SIZE)
#endif /* FOOTERS */

/* MMapped chunks need a second word of overhead ... */
#define MMAP_CHUNK_OVERHEAD (TWO_SIZE_T_SIZES)
/* ... and additional padding for fake next-chunk at foot */
2064 #define MMAP_FOOT_PAD      (FOUR_SIZE_T_SIZES)

/* The smallest size we can malloc is an aligned minimal chunk */
#define MIN_CHUNK_SIZE\
2068   ((MCHUNK_SIZE + CHUNK_ALIGN_MASK) & ~CHUNK_ALIGN_MASK)

/* conversion from malloc headers to user pointers, and back */
#define chunk2mem(p)        ((void*)((char*)(p)      + TWO_SIZE_T_SIZES))
2072 #define mem2chunk(mem)     ((mchunkptr)((char*)(mem) - TWO_SIZE_T_SIZES))
/* chunk associated with aligned address A */
#define align_as_chunk(A)   (mchunkptr)((A) + align_offset(chunk2mem(A)))

2076 /* Bounds on request (not chunk) sizes. */
#define MAX_REQUEST         ((-MIN_CHUNK_SIZE) << 2)
#define MIN_REQUEST         (MIN_CHUNK_SIZE - CHUNK_OVERHEAD - SIZE_T_ONE)

2080 /* pad request bytes into a usable size */
#define pad_request(req) \
    (((req) + CHUNK_OVERHEAD + CHUNK_ALIGN_MASK) & ~CHUNK_ALIGN_MASK)

2084 /* pad request, checking for minimum (but not maximum) */
#define request2size(req) \
    (((req) < MIN_REQUEST)? MIN_CHUNK_SIZE : pad_request(req))

```

9.3 Operations on head and foot fields

```

/*
2092 The head field of a chunk is or'ed with PINUSE_BIT when previous
      adjacent chunk in use, and or'ed with CINUSE_BIT if this chunk is in
      use, unless mmapped, in which case both bits are cleared.

2096 FLAG4_BIT is not used by this malloc, but might be useful in extensions.
*/

#define PINUSE_BIT          (SIZE_T_ONE)
2100 #define CINUSE_BIT        (SIZE_T_TWO)
#define FLAG4_BIT           (SIZE_T_FOUR)
#define INUSE_BITS          (PINUSE_BIT|CINUSE_BIT)
#define FLAG_BITS           (PINUSE_BIT|CINUSE_BIT|FLAG4_BIT)

/* Head value for fenceposts */
#define FENCEPOST_HEAD      (INUSE_BITS|SIZE_T_SIZE)

2108 /* extraction of fields from head words */
#define cinuse(p)            ((p)->head & CINUSE_BIT)
#define pinuse(p)            ((p)->head & PINUSE_BIT)
#define is_inuse(p)          (((p)->head & INUSE_BITS) != PINUSE_BIT)
2112 #define is_mmapped(p)      (((p)->head & INUSE_BITS) == 0)

#define chunksize(p)         ((p)->head & ~(FLAG_BITS))

2116 #define clear_pinuse(p)     ((p)->head &= ~PINUSE_BIT)

/* Treat space at ptr +/- offset as a chunk */
#define chunk_plus_offset(p, s) ((mchunkptr)(((char*)(p)) + (s)))
2120 #define chunk_minus_offset(p, s) ((mchunkptr)(((char*)(p)) - (s)))

/* Ptr to next or previous physical malloc_chunk. */
#define next_chunk(p) ((mchunkptr)((char*)(p) + ((p)->head & ~FLAG_BITS)))
2124 #define prev_chunk(p) ((mchunkptr)((char*)(p) - ((p)->prev_foot)))

/* extract next chunk's pinuse bit */
#define next_pinuse(p)       ((next_chunk(p)->head) & PINUSE_BIT)

/* Get/set size at footer */
#define get_foot(p, s)       (((mchunkptr)((char*)(p) + (s)))->prev_foot)
#define set_foot(p, s)       (((mchunkptr)((char*)(p) + (s)))->prev_foot = (s))

/* Set size, pinuse bit, and foot */
#define set_size_and_pinuse_of_free_chunk(p, s)\
    ((p)->head = (s|PINUSE_BIT), set_foot(p, s))

/* Set size, pinuse bit, foot, and clear next pinuse */
#define set_free_with_pinuse(p, s, n)\

```



```

                |           Pointer to left child (child[0])           |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
                |           Pointer to right child (child[1])          |
2192 +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
                |           Pointer to parent                          |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
                |           bin index of this chunk                    |
2196 +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
                |           Unused space                               .
                .                                                     |
nextchunk->+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
2200   'foot:' |           Size of chunk, in bytes                    |
                +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

```

Each tree holding treenodes is a tree of unique chunk sizes. Chunks of the same size are arranged in a circularly-linked list, with only the oldest chunk (the next to be used, in our FIFO ordering) actually in the tree. (Tree members are distinguished by a non-null parent pointer.) If a chunk with the same size as an existing node is inserted, it is linked off the existing node using pointers that work in the same way as fd/bk pointers of small chunks.

Each tree contains a power of 2 sized range of chunk sizes (the smallest is $0x100 \leq x < 0x180$), which is divided in half at each tree level, with the chunks in the smaller half of the range ($0x100 \leq x < 0x140$ for the top nose) in the left subtree and the larger half ($0x140 \leq x < 0x180$) in the right subtree. This is, of course, done by inspecting individual bits.

Using these rules, each node's left subtree contains all smaller sizes than its right subtree. However, the node at the root of each subtree has no particular ordering relationship to either. (The dividing line between the subtree sizes is based on trie relation.) If we remove the last chunk of a given size from the interior of the tree, we need to replace it with a leaf node. The tree ordering rules permit a node to be replaced by any leaf below it.

The smallest chunk in a tree (a common operation in a best-fit allocator) can be found by walking a path to the leftmost leaf in the tree. Unlike a usual binary tree, where we follow left child pointers until we reach a null, here we follow the right child pointer any time the left one is null, until we reach a leaf with both child pointers null. The smallest chunk in the tree will be somewhere along that path.

The worst case number of steps to add, find, or remove a node is bounded by the number of bits differentiating chunks within bins. Under current bin calculations, this ranges from 6 up to 21 (for 32 bit sizes) or up to 53 (for 64 bit sizes). The typical case is of course much better.

*/

```

struct malloc_tree_chunk {
    /* The first four fields must be compatible with malloc_chunk */
    size_t          prev_foot;
2244  size_t          head;
    struct malloc_tree_chunk* fd;
    struct malloc_tree_chunk* bk;

2248  struct malloc_tree_chunk* child[2];
    struct malloc_tree_chunk* parent;
    bindex_t          index;
};

typedef struct malloc_tree_chunk  tchunk;
typedef struct malloc_tree_chunk* tchunkptr;
typedef struct malloc_tree_chunk* tbinptr; /* The type of bins of trees */

/* A little helper macro for trees */
#define leftmost_child(t) ((t)->child[0] != 0? (t)->child[0] : (t)->child[1])

```

Chapter 10

Segments

```
/*
2260  Each malloc space may include non-contiguous segments, held in a
      list headed by an embedded malloc_segment record representing the
      top-most space. Segments also include flags holding properties of
2264  the space. Large chunks that are directly allocated by mmap are not
      included in this list. They are instead independently created and
      destroyed without otherwise keeping track of them.

      Segment management mainly comes into play for spaces allocated by
2268  MMAP. Any call to MMAP might or might not return memory that is
      adjacent to an existing segment. MORECORE normally contiguously
      extends the current space, so this space is almost always adjacent,
      which is simpler and faster to deal with. (This is why MORECORE is
2272  used preferentially to MMAP when both are available -- see
      sys_alloc.) When allocating using MMAP, we don't use any of the
      hinting mechanisms (inconsistently) supported in various
2276  implementations of unix mmap, or distinguish reserving from
      committing memory. Instead, we just ask for space, and exploit
      contiguity when we get it. It is probably possible to do
      better than this on some systems, but no general scheme seems
      to be significantly better.

      Management entails a simpler variant of the consolidation scheme
      used for chunks to reduce fragmentation -- new adjacent memory is
      normally prepended or appended to an existing segment. However,
2284  there are limitations compared to chunk consolidation that mostly
      reflect the fact that segment processing is relatively infrequent
      (occurring only when getting memory from system) and that we
      don't expect to have huge numbers of segments:

      * Segments are not indexed, so traversal requires linear scans. (It
        would be possible to index these, but is not worth the extra
        overhead and complexity for most programs on most platforms.)
2292  * New segments are only appended to old ones when holding top-most
        memory; if they cannot be prepended to others, they are held in
        different segments.

2296  Except for the top-most segment of an mstate, each segment record
      is kept at the tail of its segment. Segments are added by pushing
```

segment records onto the list headed by `&mstate.seg` for the containing `mstate`.

Segment flags control allocation/merge/deallocation policies:

- * If `EXTERN_BIT` set, then we did not allocate this segment, and so should not try to deallocate or merge with others. (This currently holds only for the initial segment passed into `create_mspace_with_base`.)
- * If `USE_MMAP_BIT` set, the segment may be merged with other surrounding mmapped segments and trimmed/de-allocated using `munmap`.
- * If neither bit is set, then the segment was obtained using `MORECORE` so can be merged with surrounding `MORECORE`'d segments and deallocated/trimmed using `MORECORE` with negative arguments.

2312 */

```
struct malloc_segment {
    char*      base;           /* base address */
    size_t     size;           /* allocated size */
    struct malloc_segment* next; /* ptr to next segment */
    flag_t     sflags;         /* mmap and extern flag */
};
```

2316

```
#define is_mmapped_segment(S) ((S)->sflags & USE_MMAP_BIT)
#define is_extern_segment(S)  ((S)->sflags & EXTERN_BIT)
```

2324

```
typedef struct malloc_segment msegment;
typedef struct malloc_segment* msegmentptr;
```


Chapter 11

State

11.1 malloc_state

```
/*
2328  A malloc_state holds all of the bookkeeping for a space.
    The main fields are:

    Top
2332  The topmost chunk of the currently active segment. Its size is
    cached in topsize. The actual size of topmost space is
    topsize+TOP_FOOT_SIZE, which includes space reserved for adding
2336  fenceposts and segment records if necessary when getting more
    space from the system. The size at which to autotrim top is
    cached from mparams in trim_check, except that it is disabled if
    an autotrim fails.

2340  Designated victim (dv)
    This is the preferred chunk for servicing small requests that
    don't have exact fits. It is normally the chunk split off most
    recently to service another small request. Its size is cached in
2344  dvsize. The link fields of this chunk are not maintained since it
    is not kept in a bin.

    SmallBins
2348  An array of bin headers for free chunks. These bins hold chunks
    with sizes less than MIN_LARGE_SIZE bytes. Each bin contains
    chunks of all the same size, spaced 8 bytes apart. To simplify
    use in double-linked lists, each bin header acts as a malloc_chunk
2352  pointing to the real first node, if it exists (else pointing to
    itself). This avoids special-casing for headers. But to avoid
    waste, we allocate only the fd/bk pointers of bins, and then use
    repositioning tricks to treat these as the fields of a chunk.

    TreeBins
    Treebins are pointers to the roots of trees holding a range of
    sizes. There are 2 equally spaced treebins for each power of two
2360  from TREE_SHIFT to TREE_SHIFT+16. The last bin holds anything
    larger.
```

Bin maps

2364 There is one bit map for small bins ("smallmap") and one for
 treebins ("treemap"). Each bin sets its bit when non-empty, and
 clears the bit when empty. Bit operations are then used to avoid
 bin-by-bin searching -- nearly all "search" is done without ever
 2368 looking at bins that won't be selected. The bit maps
 conservatively use 32 bits per map word, even if on 64bit system.
 For a good description of some of the bit-based techniques used
 here, see Henry S. Warren Jr's book "Hacker's Delight" (and
 2372 supplement at <http://hackersdelight.org/>). Many of these are
 intended to reduce the branchiness of paths through malloc etc, as
 well as to reduce the number of memory locations read or written.

Segments

2376 A list of segments headed by an embedded malloc_segment record
 representing the initial space.

Address check support

2380 The least_addr field is the least address ever obtained from
 MORECORE or MMAP. Attempted frees and reallocs of any address less
 than this are trapped (unless INSECURE is defined).

Magic tag

A cross-check field that should always hold same value as mparams.magic.

Flags

2388 Bits recording whether to use MMAP, locks, or contiguous MORECORE

Statistics

2392 Each space keeps track of current and maximum system memory
 obtained via MORECORE or MMAP.

Trim support

2396 Fields holding the amount of unused topmost memory that should trigger
 timing, and a counter to force periodic scanning to release unused
 non-topmost segments.

Locking

2400 If USE_LOCKS is defined, the "mutex" lock is acquired and released
 around every public call using this mspace.

Extension support

2404 A void* pointer and a size_t field that can be used to help implement
 extensions to this malloc.

*/

/* Bin types, widths and sizes */

```
#define NSMALLBINS      (32U)
#define NTREEBINS      (32U)
2412 #define SMALLBIN_SHIFT (3U)
#define SMALLBIN_WIDTH  (SIZE_T_ONE << SMALLBIN_SHIFT)
#define TREEBIN_SHIFT   (8U)
```

```

#define MIN_LARGE_SIZE      (SIZE_T_ONE << TREEBIN_SHIFT)
2416 #define MAX_SMALL_SIZE    (MIN_LARGE_SIZE - SIZE_T_ONE)
#define MAX_SMALL_REQUEST  (MAX_SMALL_SIZE - CHUNK_ALIGN_MASK - CHUNK_OVERHEAD)

struct malloc_state {
2420     binmap_t    smallmap;
        binmap_t    treemap;
        size_t     dvsize;
        size_t     topsize;
2424     char*       least_addr;
        mchunkptr  dv;
        mchunkptr  top;
        size_t     trim_check;
2428     size_t     release_checks;
        size_t     magic;
        mchunkptr  smallbins[(NSMALLBINS+1)*2];
        tbinptr    treebins[NTREEBINS];
2432     size_t     footprint;
        size_t     max_footprint;
        flag_t     mflags;
#ifdef USE_LOCKS
2436     MLOCK_T     mutex;      /* locate lock among fields that rarely change */
#endif /* USE_LOCKS */
        msegment    seg;
        void*       extp;      /* Unused but available for extensions */
2440     size_t     exts;
};

typedef struct malloc_state*    mstate;

```

11.2 Global malloc_state and malloc_params

```

/*
2448     malloc_params holds global properties, including those that can be
        dynamically set using mallopt. There is a single instance, mparams,
        initialized in init_mparams. Note that the non-zerosness of "magic"
        also serves as an initialization flag.
2452 */

struct malloc_params {
        volatile size_t magic;
2456     size_t page_size;
        size_t granularity;
        size_t mmap_threshold;
        size_t trim_threshold;
2460     flag_t default_mflags;
};

static struct malloc_params mparams;

```

```
/* Ensure mparams initialized */
#define ensure_initialization() (void)(mparams.magic != 0 || init_mparams())

2468 #if !ONLY_MSPACES

/* The global malloc_state used for all non-"mspace" calls */
static struct malloc_state _gm_;
2472 #define gm          (&_gm_)
#define is_global(M)  ((M) == &_gm_)

#endif /* !ONLY_MSPACES */

#define is_initialized(M) ((M)->top != 0)
```

Chapter 12

System alloc setup

```
/* Operations on mflags */

2480 #define use_lock(M)          ((M)->mflags &  USE_LOCK_BIT)
      #define enable_lock(M)    ((M)->mflags |=  USE_LOCK_BIT)
      #define disable_lock(M)   ((M)->mflags &= ~USE_LOCK_BIT)

2484 #define use_mmap(M)          ((M)->mflags &  USE_MMAP_BIT)
      #define enable_mmap(M)    ((M)->mflags |=  USE_MMAP_BIT)
      #define disable_mmap(M)   ((M)->mflags &= ~USE_MMAP_BIT)

2488 #define use_noncontiguous(M) ((M)->mflags &  USE_NONCONTIGUOUS_BIT)
      #define disable_contiguous(M) ((M)->mflags |=  USE_NONCONTIGUOUS_BIT)

      #define set_lock(M,L)\
2492   ((M)->mflags = (L)?\
        ((M)->mflags | USE_LOCK_BIT) :\
        ((M)->mflags & ~USE_LOCK_BIT))

2496 /* page-align a size */
      #define page_align(S)\
        (((S) + (mparams.page_size - SIZE_T_ONE)) & ~(mparams.page_size - SIZE_T_ONE))

2500 /* granularity-align a size */
      #define granularity_align(S)\
        (((S) + (mparams.granularity - SIZE_T_ONE))\
        & ~(mparams.granularity - SIZE_T_ONE))

/* For mmap, use granularity alignment on windows, else page-align */
#ifdef WIN32
2508 #define mmap_align(S) granularity_align(S)
      #else
      #define mmap_align(S) page_align(S)
      #endif

/* For sys_alloc, enough padding to ensure can malloc request on success */
#define SYS_ALLOC_PADDING (TOP_FOOT_SIZE + MALLOC_ALIGNMENT)

2516 #define is_page_aligned(S)\
```


Chapter 13

Hooks

```
/*
  PREACTION should be defined to return 0 on success, and nonzero on
2564 failure. If you are not using locking, you can redefine these to do
  anything you like.
*/

2568 #if USE_LOCKS

#define PREACTION(M) ((use_lock(M))? ACQUIRE_LOCK(&(M)->mutex) : 0)
#define POSTACTION(M) { if (use_lock(M)) RELEASE_LOCK(&(M)->mutex); }
2572 #else /* USE_LOCKS */

#ifndef PREACTION
#define PREACTION(M) (0)
2576 #endif /* PREACTION */

#ifndef POSTACTION
#define POSTACTION(M)
2580 #endif /* POSTACTION */

#endif /* USE_LOCKS */

2584 /*
  CORRUPTION_ERROR_ACTION is triggered upon detected bad addresses.
  USAGE_ERROR_ACTION is triggered on detected bad frees and
  reallocs. The argument p is an address that might have triggered the
2588 fault. It is ignored by the two predefined actions, but might be
  useful in custom actions that try to help diagnose errors.
*/

2592 #if PROCEED_ON_ERROR

/* A count of the number of corruption errors causing resets */
int malloc_corruption_error_count;

/* default corruption action */
static void reset_on_error(mstate m);

2600 #define CORRUPTION_ERROR_ACTION(m) reset_on_error(m)
```

```
#define USAGE_ERROR_ACTION(m, p)

#else /* PROCEED_ON_ERROR */

#ifndef CORRUPTION_ERROR_ACTION
#define CORRUPTION_ERROR_ACTION(m) ABORT
#endif /* CORRUPTION_ERROR_ACTION */

#ifndef USAGE_ERROR_ACTION
#define USAGE_ERROR_ACTION(m,p) ABORT
#endif /* USAGE_ERROR_ACTION */

#endif /* PROCEED_ON_ERROR */
```


Chapter 14

Debugging setup

```
#if ! DEBUG

2616 #define check_free_chunk(M,P)
      #define check_inuse_chunk(M,P)
      #define check_malloced_chunk(M,P,N)
      #define check_mmapped_chunk(M,P)
2620 #define check_malloc_state(M)
      #define check_top_chunk(M,P)

      #else /* DEBUG */

2624 #define check_free_chunk(M,P)      do_check_free_chunk(M,P)
      #define check_inuse_chunk(M,P)  do_check_inuse_chunk(M,P)
      #define check_top_chunk(M,P)    do_check_top_chunk(M,P)
      #define check_malloced_chunk(M,P,N) do_check_malloced_chunk(M,P,N)
2628 #define check_mmapped_chunk(M,P)    do_check_mmapped_chunk(M,P)
      #define check_malloc_state(M)    do_check_malloc_state(M)

      static void do_check_any_chunk(mstate m, mchunkptr p);
2632 static void do_check_top_chunk(mstate m, mchunkptr p);
      static void do_check_mmapped_chunk(mstate m, mchunkptr p);
      static void do_check_inuse_chunk(mstate m, mchunkptr p);
      static void do_check_free_chunk(mstate m, mchunkptr p);
2636 static void do_check_malloced_chunk(mstate m, void* mem, size_t s);
      static void do_check_tree(mstate m, tchunkptr t);
      static void do_check_treebin(mstate m, bindex_t i);
      static void do_check_smallbin(mstate m, bindex_t i);
2640 static void do_check_malloc_state(mstate m);
      static int bin_find(mstate m, mchunkptr x);
      static size_t traverse_and_check(mstate m);
      #endif /* DEBUG */
```

Chapter 15

Bins

2644

15.1 Indexing Bins

```
#define is_small(s)          (((s) >> SMALLBIN_SHIFT) < NSMALLBINS)
#define small_index(s)      ((s) >> SMALLBIN_SHIFT)
2648 #define small_index2size(i) ((i) << SMALLBIN_SHIFT)
#define MIN_SMALL_INDEX     (small_index(MIN_CHUNK_SIZE))

/* addressing by index. See above about smallbin repositioning */
2652 #define smallbin_at(M, i)  ((sbinptr)((char*)&((M)->smallbins[(i)<<1])))
#define treebin_at(M,i)     (&((M)->treebins[i]))

/* assign tree index for size S to variable I. Use x86 asm if possible */
2656 #if defined(__GNUC__) && (defined(__i386__) || defined(__x86_64__))
#define compute_tree_index(S, I)\
{\
    unsigned int X = S >> TREEBIN_SHIFT;\
2660     if (X == 0)\
        I = 0;\
    else if (X > 0xFFFF)\
        I = NTREEBINS-1;\
2664     else {\
        unsigned int K;\
        __asm__("bsrl\t%1, %0\n\t" : "=r" (K) : "g" (X));\
        I = (bindex_t)((K << 1) + ((S >> (K + (TREEBIN_SHIFT-1)) & 1))); \
2668     }\
}

#elif defined (__INTEL_COMPILER)
2672 #define compute_tree_index(S, I)\
{\
    size_t X = S >> TREEBIN_SHIFT;\
    if (X == 0)\
2676     I = 0;\
    else if (X > 0xFFFF)\
        I = NTREEBINS-1;\
    else {\
2680     unsigned int K = _bit_scan_reverse (X); \
```

```

        I = (bindex_t)((K << 1) + ((S >> (K + (TREEBIN_SHIFT-1)) & 1))); \
    } \
}

#elif defined(_MSC_VER) && _MSC_VER>=1300
#define compute_tree_index(S, I) \
{ \
2688     size_t X = S >> TREEBIN_SHIFT; \
        if (X == 0) \
            I = 0; \
        else if (X > 0xFFFF) \
2692             I = NTREEBINS-1; \
        else { \
            unsigned int K; \
            _BitScanReverse((DWORD *) &K, X); \
2696             I = (bindex_t)((K << 1) + ((S >> (K + (TREEBIN_SHIFT-1)) & 1))); \
        } \
    }

2700 #else /* GNUC */
#define compute_tree_index(S, I) \
{ \
    size_t X = S >> TREEBIN_SHIFT; \
2704     if (X == 0) \
        I = 0; \
        else if (X > 0xFFFF) \
            I = NTREEBINS-1; \
2708     else { \
        unsigned int Y = (unsigned int)X; \
        unsigned int N = ((Y - 0x100) >> 16) & 8; \
        unsigned int K = (((Y <= N) - 0x1000) >> 16) & 4; \
2712         N += K; \
        N += K = (((Y <= K) - 0x4000) >> 16) & 2; \
        K = 14 - N + ((Y <= K) >> 15); \
        I = (K << 1) + ((S >> (K + (TREEBIN_SHIFT-1)) & 1)); \
2716     } \
    }
#endif /* GNUC */

2720 /* Bit representing maximum resolved size in a treebin at i */
#define bit_for_tree_index(i) \
    (i == NTREEBINS-1) ? (SIZE_T_BITSIZE-1) : (((i) >> 1) + TREEBIN_SHIFT - 2)

2724 /* Shift placing maximum resolved bit in a treebin at i as sign bit */
#define leftshift_for_tree_index(i) \
    ((i == NTREEBINS-1) ? 0 : \
        ((SIZE_T_BITSIZE-SIZE_T_ONE) - (((i) >> 1) + TREEBIN_SHIFT - 2)))

/* The size of the smallest chunk held in bin with index i */
#define minsize_for_tree_index(i) \
    ((SIZE_T_ONE << (((i) >> 1) + TREEBIN_SHIFT)) | \
2732     (((size_t)((i) & SIZE_T_ONE)) << (((i) >> 1) + TREEBIN_SHIFT - 1)))

```

15.2 Operations on bin maps

```

/* bit corresponding to given index */
#define idx2bit(i)          ((binmap_t)(1) << (i))

2740 /* Mark/Clear bits with given index */
#define mark_smallmap(M,i)   ((M)->smallmap |=  idx2bit(i))
#define clear_smallmap(M,i) ((M)->smallmap &= ~idx2bit(i))
#define smallmap_is_marked(M,i) ((M)->smallmap &  idx2bit(i))

#define mark_treemap(M,i)   ((M)->treemap |=  idx2bit(i))
#define clear_treemap(M,i) ((M)->treemap &= ~idx2bit(i))
#define treemap_is_marked(M,i) ((M)->treemap &  idx2bit(i))

/* isolate the least set bit of a bitmap */
#define least_bit(x)        ((x) & -(x))

2752 /* mask with all bits to left of least bit of x on */
#define left_bits(x)        ((x<<1) | -(x<<1))

/* mask with all bits to left of or equal to least bit of x on */
2756 #define same_or_left_bits(x) ((x) | -(x))

/* index corresponding to given bit. Use x86 asm if possible */

2760 #if defined(__GNUC__) && (defined(__i386__) || defined(__x86_64__))
#define compute_bit2idx(X, I)\
{\
    unsigned int J;\
2764    __asm__ ("bsfl\t%1, %0\n\t" : "=r" (J) : "g" (X));\
    I = (bindex_t)J;\
}

2768 #elif defined (__INTEL_COMPILER)
#define compute_bit2idx(X, I)\
{\
    unsigned int J;\
2772    J = _bit_scan_forward (X); \
    I = (bindex_t)J;\
}

2776 #elif defined(_MSC_VER) && _MSC_VER>=1300
#define compute_bit2idx(X, I)\
{\
    unsigned int J;\
2780    _BitScanForward((DWORD *) &J, X);\
    I = (bindex_t)J;\
}

```

```

2784 #elif USE_BUILTIN_FFS
    #define compute_bit2idx(X, I) I = ffs(X)-1

    #else
2788 #define compute_bit2idx(X, I)\
    {\
        unsigned int Y = X - 1;\
        unsigned int K = Y >> (16-4) & 16;\
2792 unsigned int N = K;      Y >>= K;\
        N += K = Y >> (8-3) & 8;  Y >>= K;\
        N += K = Y >> (4-2) & 4;  Y >>= K;\
        N += K = Y >> (2-1) & 2;  Y >>= K;\
2796 N += K = Y >> (1-0) & 1;  Y >>= K;\
        I = (bindex_t)(N + Y);\
    }
    #endif /* GNUC */

```

Chapter 16

Runtime Check Support

```
2800  /*
    For security, the main invariant is that malloc/free/etc never
    writes to a static address other than malloc_state, unless static
    malloc_state itself has been corrupted, which cannot occur via
2804  malloc (because of these checks). In essence this means that we
    believe all pointers, sizes, maps etc held in malloc_state, but
    check all of those linked or offsetted from other embedded data
    structures. These checks are interspersed with main code in a way
2808  that tends to minimize their run-time cost.

    When FOOTERS is defined, in addition to range checking, we also
    verify footer fields of inuse chunks, which can be used guarantee
2812  that the mstate controlling malloc/free is intact. This is a
    streamlined version of the approach described by William Robertson
    et al in "Run-time Detection of Heap-based Overflows" LISA'03
    http://www.usenix.org/events/lisa03/tech/robertson.html The footer
2816  of an inuse chunk holds the xor of its mstate and a random seed,
    that is checked upon calls to free() and realloc(). This is
    (probablistically) unguessable from outside the program, but can be
    computed by any code successfully malloc'ing any chunk, so does not
2820  itself provide protection against code that has already broken
    security through some other means. Unlike Robertson et al, we
    always dynamically check addresses of all offset chunks (previous,
    next, etc). This turns out to be cheaper than relying on hashes.
2824  */

    #if !INSECURE
    /* Check if address a is at least as high as any from MORECORE or MMAP */
2828  #define ok_address(M, a) ((char*)(a) >= (M)->least_addr)
    /* Check if address of next chunk n is higher than base chunk p */
    #define ok_next(p, n)    ((char*)(p) < (char*)(n))
    /* Check if p has inuse status */
    #define ok_inuse(p)      is_inuse(p)
2832  /* Check if p has its pinuse bit on */
    #define ok_pinuse(p)     pinuse(p)

2836  #else /* !INSECURE */
    #define ok_address(M, a) (1)
    #define ok_next(b, n)    (1)
```

```

#define ok_inuse(p)      (1)
2840 #define ok_pinuse(p)   (1)
#define /* !INSECURE */

#if (FOOTERS && !INSECURE)
2844 /* Check if (alleged) mstate m has expected magic field */
#define ok_magic(M)      ((M)->magic == mparams.magic)
#define /* (FOOTERS && !INSECURE) */
#define ok_magic(M)      (1)
2848 #endif /* (FOOTERS && !INSECURE) */

/* In gcc, use __builtin_expect to minimize impact of checks */
2852 #if !INSECURE
#define __builtin_expect(e, 1)
#define RTCHECK(e)
#else /* GNUC */
2856 #define RTCHECK(e) (e)
#define /* GNUC */
#define /* !INSECURE */
#define RTCHECK(e) (1)
2860 #endif /* !INSECURE */

/* macros to set up inuse chunks with or without footers */

2864 #if !FOOTERS

#define mark_inuse_foot(M,p,s)

2868 /* Macros for setting head/foot of non-mmapped chunks */

/* Set cinuse bit and pinuse bit of next chunk */
#define set_inuse(M,p,s)\
2872 ((p)->head = (((p)->head & PINUSE_BIT)|s|CINUSE_BIT),\
((mchunkptr)(((char*)(p)) + (s)))->head |= PINUSE_BIT)

/* Set cinuse and pinuse of this chunk and pinuse of next chunk */
2876 #define set_inuse_and_pinuse(M,p,s)\
((p)->head = (s|PINUSE_BIT|CINUSE_BIT),\
((mchunkptr)(((char*)(p)) + (s)))->head |= PINUSE_BIT)

2880 /* Set size, cinuse and pinuse bit of this chunk */
#define set_size_and_pinuse_of_inuse_chunk(M, p, s)\
((p)->head = (s|PINUSE_BIT|CINUSE_BIT))

2884 #else /* FOOTERS */

/* Set foot of inuse chunk to be xor of mstate and seed */
#define mark_inuse_foot(M,p,s)\
2888 (((mchunkptr)(((char*)(p)) + (s)))->prev_foot = ((size_t)(M) ^ mparams.magic))

#define get_mstate_for(p)\

```

```

((mstate)(((mchunkptr)((char*)(p) +\
2892     (chunksize(p))))->prev_foot ^ mparams.magic))

#define set_inuse(M,p,s)\
    ((p)->head = (((p)->head & PINUSE_BIT)|s|CINUSE_BIT),\
2896     (((mchunkptr)(((char*)(p)) + (s)))->head |= PINUSE_BIT), \
    mark_inuse_foot(M,p,s))

#define set_inuse_and_pinuse(M,p,s)\
2900     ((p)->head = (s|PINUSE_BIT|CINUSE_BIT),\
    (((mchunkptr)(((char*)(p)) + (s)))->head |= PINUSE_BIT),\
    mark_inuse_foot(M,p,s))

2904 #define set_size_and_pinuse_of_inuse_chunk(M, p, s)\
    ((p)->head = (s|PINUSE_BIT|CINUSE_BIT),\
    mark_inuse_foot(M, p, s))

2908 #endif /* !FOOTERS */

```


Chapter 17

Setting mparams

```
/* Initialize mparams */
static int init_mparams(void) {
#ifdef NEED_GLOBAL_LOCK_INIT
2912   if (malloc_global_mutex_status <= 0)
        init_malloc_global_mutex();
#endif

2916   ACQUIRE_MALLOC_GLOBAL_LOCK();
   if (mparams.magic == 0) {
       size_t magic;
       size_t psize;
2920       size_t gsize;

#ifdef WIN32
       psize = malloc_getpagesize;
2924       gsize = ((DEFAULT_GRANULARITY != 0)? DEFAULT_GRANULARITY : psize);
#else /* WIN32 */
       {
           SYSTEM_INFO system_info;
2928           GetSystemInfo(&system_info);
           psize = system_info.dwPageSize;
           gsize = ((DEFAULT_GRANULARITY != 0)?
                   DEFAULT_GRANULARITY : system_info.dwAllocationGranularity);
2932       }
#endif /* WIN32 */

       /* Sanity-check configuration:
2936       size_t must be unsigned and as wide as pointer type.
           ints must be at least 4 bytes.
           alignment must be at least 8.
           Alignment, min chunk size, and page size must all be powers of 2.
2940       */
       if ((sizeof(size_t) != sizeof(char*)) ||
           (MAX_SIZE_T < MIN_CHUNK_SIZE) ||
           (sizeof(int) < 4) ||
2944           (MALLOC_ALIGNMENT < (size_t)8U) ||
           ((MALLOC_ALIGNMENT & (MALLOC_ALIGNMENT-SIZE_T_ONE)) != 0) ||
           ((MCHUNK_SIZE & (MCHUNK_SIZE-SIZE_T_ONE)) != 0) ||
           ((gsize & (gsize-SIZE_T_ONE)) != 0) ||
```

```

2948         ((psize                & (psize-SIZE_T_ONE))                != 0))
        ABORT;

        mparams.granularity = gsize;
2952     mparams.page_size = psize;
        mparams.mmap_threshold = DEFAULT_MMAP_THRESHOLD;
        mparams.trim_threshold = DEFAULT_TRIM_THRESHOLD;
    #if MORECORE_CONTIGUOUS
2956     mparams.default_mflags = USE_LOCK_BIT|USE_MMAP_BIT;
    #else /* MORECORE_CONTIGUOUS */
        mparams.default_mflags = USE_LOCK_BIT|USE_MMAP_BIT|USE_NONCONTIGUOUS_BIT;
    #endif /* MORECORE_CONTIGUOUS */

    #if !ONLY_MSPACES
        /* Set up lock for main malloc area */
        gm->mflags = mparams.default_mflags;
2964     INITIAL_LOCK(&gm->mutex);
    #endif

    {
2968 #if USE_DEV_RANDOM
        int fd;
        unsigned char buf[sizeof(size_t)];
        /* Try to use /dev/urandom, else fall back on using time */
2972     if ((fd = open("/dev/urandom", O_RDONLY)) >= 0 &&
        read(fd, buf, sizeof(buf)) == sizeof(buf)) {
        magic = *((size_t *) buf);
        close(fd);
2976     }
        else
    #endif /* USE_DEV_RANDOM */
    #ifdef WIN32
2980     magic = (size_t)(GetTickCount() ^ (size_t)0x55555555U);
    #else
        magic = (size_t)(time(0) ^ (size_t)0x55555555U);
    #endif
2984     magic |= (size_t)8U;    /* ensure nonzero */
        magic &= ~(size_t)7U; /* improve chances of fault for bad values */
        mparams.magic = magic;
    }
2988 }

    RELEASE_MALLOC_GLOBAL_LOCK();
    return 1;
2992 }

/* support for mallopt */
static int change_mparam(int param_number, int value) {
2996     size_t val;
    ensure_initialization();
    val = (value == -1)? MAX_SIZE_T : (size_t)value;
    switch(param_number) {

```

```
3000  case M_TRIM_THRESHOLD:
        mparams.trim_threshold = val;
        return 1;
case M_GRANULARITY:
3004  if (val >= mparams.page_size && ((val & (val-1)) == 0)) {
        mparams.granularity = val;
        return 1;
    }
3008  else
        return 0;
case M_MMAP_THRESHOLD:
        mparams.mmap_threshold = val;
3012  return 1;
default:
        return 0;
    }
3016 }
```

Chapter 18

Debugging Support

```
#if DEBUG

/* Check properties of any chunk, whether free, inuse, mmaped etc */
3020 static void do_check_any_chunk(mstate m, mchunkptr p) {
    assert((is_aligned(chunk2mem(p))) || (p->head == FENCEPOST_HEAD));
    assert(ok_address(m, p));
}

/* Check properties of top chunk */
static void do_check_top_chunk(mstate m, mchunkptr p) {
    msegmentptr sp = segment_holding(m, (char*)p);
3028 size_t sz = p->head & ~INUSE_BITS; /* third-lowest bit can be set! */
    assert(sp != 0);
    assert((is_aligned(chunk2mem(p))) || (p->head == FENCEPOST_HEAD));
    assert(ok_address(m, p));
3032 assert(sz == m->topsize);
    assert(sz > 0);
    assert(sz == ((sp->base + sp->size) - (char*)p) - TOP_FOOT_SIZE);
    assert(pinuse(p));
3036 assert(!pinuse(chunk_plus_offset(p, sz)));
}

/* Check properties of (inuse) mmaped chunks */
3040 static void do_check_mmaped_chunk(mstate m, mchunkptr p) {
    size_t sz = chunksize(p);
    size_t len = (sz + (p->prev_foot) + MMAP_FOOT_PAD);
    assert(is_mmaped(p));
3044 assert(use_mmap(m));
    assert((is_aligned(chunk2mem(p))) || (p->head == FENCEPOST_HEAD));
    assert(ok_address(m, p));
    assert(!is_small(sz));
3048 assert((len & (mparams.page_size-SIZE_T_ONE)) == 0);
    assert(chunk_plus_offset(p, sz)->head == FENCEPOST_HEAD);
    assert(chunk_plus_offset(p, sz+SIZE_T_SIZE)->head == 0);
}

/* Check properties of inuse chunks */
static void do_check_inuse_chunk(mstate m, mchunkptr p) {
    do_check_any_chunk(m, p);
}
```

```

3056     assert(is_inuse(p));
        assert(next_pinuse(p));
        /* If not pinuse and not mmapped, previous chunk has OK offset */
        assert(is_mmapped(p) || pinuse(p) || next_chunk(prev_chunk(p)) == p);
3060     if (is_mmapped(p))
        do_check_mmapped_chunk(m, p);
    }

3064     /* Check properties of free chunks */
    static void do_check_free_chunk(mstate m, mchunkptr p) {
        size_t sz = chunksize(p);
        mchunkptr next = chunk_plus_offset(p, sz);
3068     do_check_any_chunk(m, p);
        assert(!is_inuse(p));
        assert(!next_pinuse(p));
        assert(!is_mmapped(p));
3072     if (p != m->dv && p != m->top) {
        if (sz >= MIN_CHUNK_SIZE) {
            assert((sz & CHUNK_ALIGN_MASK) == 0);
            assert(is_aligned(chunk2mem(p)));
3076            assert(next->prev_foot == sz);
            assert(pinuse(p));
            assert(next == m->top || is_inuse(next));
            assert(p->fd->bk == p);
3080            assert(p->bk->fd == p);
        }
        else /* markers are always of size SIZE_T_SIZE */
            assert(sz == SIZE_T_SIZE);
3084     }
    }

    /* Check properties of malloced chunks at the point they are malloced */
3088     static void do_check_malloced_chunk(mstate m, void* mem, size_t s) {
        if (mem != 0) {
            mchunkptr p = mem2chunk(mem);
            size_t sz = p->head & ~INUSE_BITS;
3092            do_check_inuse_chunk(m, p);
            assert((sz & CHUNK_ALIGN_MASK) == 0);
            assert(sz >= MIN_CHUNK_SIZE);
            assert(sz >= s);
3096            /* unless mmapped, size is less than MIN_CHUNK_SIZE more than request */
            assert(is_mmapped(p) || sz < (s + MIN_CHUNK_SIZE));
        }
    }

    /* Check a tree and its subtrees. */
    static void do_check_tree(mstate m, tchunkptr t) {
        tchunkptr head = 0;
3104     tchunkptr u = t;
        bindex_t tindex = t->index;
        size_t tsize = chunksize(t);
        bindex_t idx;

```

```

3108 compute_tree_index(tsize, idx);
    assert(tindex == idx);
    assert(tsize >= MIN_LARGE_SIZE);
    assert(tsize >= minsize_for_tree_index(idx));
3112 assert((idx == NTREEBINS-1) || (tsize < minsize_for_tree_index((idx+1))));

    do { /* traverse through chain of same-sized nodes */
        do_check_any_chunk(m, ((mchunkptr)u));
3116 assert(u->index == tindex);
        assert(chunksize(u) == tsize);
        assert(!is_inuse(u));
        assert(!next_pinuse(u));
3120 assert(u->fd->bk == u);
        assert(u->bk->fd == u);
        if (u->parent == 0) {
            assert(u->child[0] == 0);
3124 assert(u->child[1] == 0);
        }
        else {
            assert(head == 0); /* only one node on chain has parent */
3128 head = u;
            assert(u->parent != u);
            assert (u->parent->child[0] == u ||
                    u->parent->child[1] == u ||
3132 *((tbinptr*)(u->parent)) == u);
            if (u->child[0] != 0) {
                assert(u->child[0]->parent == u);
                assert(u->child[0] != u);
3136 do_check_tree(m, u->child[0]);
            }
            if (u->child[1] != 0) {
                assert(u->child[1]->parent == u);
3140 assert(u->child[1] != u);
                do_check_tree(m, u->child[1]);
            }
            if (u->child[0] != 0 && u->child[1] != 0) {
3144 assert(chunksize(u->child[0]) < chunksize(u->child[1]));
            }
        }
        u = u->fd;
3148 } while (u != t);
    assert(head != 0);
}

3152 /* Check all the chunks in a treebin. */
static void do_check_treebin(mstate m, bindex_t i) {
    tbinptr* tb = treebin_at(m, i);
    tchunkptr t = *tb;
3156 int empty = (m->treemap & (1U << i)) == 0;
    if (t == 0)
        assert(empty);
    if (!empty)

```

```

3160     do_check_tree(m, t);
    }

    /* Check all the chunks in a smallbin. */
3164 static void do_check_smallbin(mstate m, bindex_t i) {
    sbinptr b = smallbin_at(m, i);
    mchunkptr p = b->bk;
    unsigned int empty = (m->smallmap & (1U << i)) == 0;
3168     if (p == b)
        assert(empty);
    if (!empty) {
        for (; p != b; p = p->bk) {
3172             size_t size = chunksize(p);
            mchunkptr q;
            /* each chunk claims to be free */
            do_check_free_chunk(m, p);
3176             /* chunk belongs in bin */
            assert(small_index(size) == i);
            assert(p->bk == b || chunksize(p->bk) == chunksize(p));
            /* chunk is followed by an inuse chunk */
3180             q = next_chunk(p);
            if (q->head != FENCEPOST_HEAD)
                do_check_inuse_chunk(m, q);
        }
3184     }
}

/* Find x in a bin. Used in other check functions. */
3188 static int bin_find(mstate m, mchunkptr x) {
    size_t size = chunksize(x);
    if (is_small(size)) {
        bindex_t sidx = small_index(size);
3192         sbinptr b = smallbin_at(m, sidx);
        if (smallmap_is_marked(m, sidx)) {
            mchunkptr p = b;
            do {
3196                 if (p == x)
                    return 1;
            } while ((p = p->fd) != b);
        }
3200     }
    else {
        bindex_t tidx;
        compute_tree_index(size, tidx);
3204         if (treemap_is_marked(m, tidx)) {
            tchunkptr t = *treebin_at(m, tidx);
            size_t sizebits = size << leftshift_for_tree_index(tidx);
            while (t != 0 && chunksize(t) != size) {
3208                 t = t->child[(sizebits >> (SIZE_T_BITSIZE-SIZE_T_ONE)) & 1];
                sizebits <=> 1;
            }
            if (t != 0) {

```

```

3212     tchunkptr u = t;
        do {
            if (u == (tchunkptr)x)
                return 1;
3216     } while ((u = u->fd) != t);
    }
}
}
3220 return 0;
}

/* Traverse each chunk and check it; return total */
3224 static size_t traverse_and_check(mstate m) {
    size_t sum = 0;
    if (is_initialized(m)) {
        msegmentptr s = &m->seg;
3228     sum += m->topsize + TOP_FOOT_SIZE;
        while (s != 0) {
            mchunkptr q = align_as_chunk(s->base);
            mchunkptr lastq = 0;
3232     assert(pinuse(q));
            while (segment_holds(s, q) &&
                    q != m->top && q->head != FENCEPOST_HEAD) {
                sum += chunksize(q);
3236     if (is_inuse(q)) {
                    assert(!bin_find(m, q));
                    do_check_inuse_chunk(m, q);
                }
3240     else {
                    assert(q == m->dv || bin_find(m, q));
                    assert(lastq == 0 || is_inuse(lastq)); /* Not 2 consecutive free */
                    do_check_free_chunk(m, q);
3244     }
                    lastq = q;
                    q = next_chunk(q);
                }
3248     s = s->next;
    }
}
return sum;
3252 }

/* Check all properties of malloc_state. */
static void do_check_malloc_state(mstate m) {
3256     bindex_t i;
    size_t total;
    /* check bins */
    for (i = 0; i < NSMALLBINS; ++i)
        do_check_smallbin(m, i);
3260     for (i = 0; i < NTREEBINS; ++i)
        do_check_treebin(m, i);

```



```

3264     if (m->dvsizе != 0) { /* check dv chunk */
        do_check_any_chunk(m, m->dv);
        assert(m->dvsizе == chunksize(m->dv));
        assert(m->dvsizе >= MIN_CHUNK_SIZE);
3268     assert(bin_find(m, m->dv) == 0);
    }

    if (m->top != 0) { /* check top chunk */
3272     do_check_top_chunk(m, m->top);
        /*assert(m->topsize == chunksize(m->top)); redundant */
        assert(m->topsize > 0);
        assert(bin_find(m, m->top) == 0);
3276     }

    total = traverse_and_check(m);
    assert(total <= m->footprint);
3280     assert(m->footprint <= m->max_footprint);
    }
#endif /* DEBUG */

```

Chapter 19

Statistics

```

    #if !NO_MALLINFO
3284 static struct mallinfo internal_mallinfo(mstate m) {
    struct mallinfo nm = { 0, 0, 0, 0, 0, 0, 0, 0, 0, 0 };
    ensure_initialization();
    if (!PREACTION(m)) {
3288     check_malloc_state(m);
    if (is_initialized(m)) {
        size_t nfree = SIZE_T_ONE; /* top always free */
        size_t mfree = m->topsize + TOP_FOOT_SIZE;
3292     size_t sum = mfree;
        msegmentptr s = &m->seg;
        while (s != 0) {
            mchunkptr q = align_as_chunk(s->base);
3296     while (segment_holds(s, q) &&
                q != m->top && q->head != FENCEPOST_HEAD) {
                size_t sz = chunksize(q);
                sum += sz;
3300                if (!is_inuse(q)) {
                    mfree += sz;
                    ++nfree;
                }
3304                q = next_chunk(q);
            }
            s = s->next;
        }

        nm.arena      = sum;
        nm.ordblks     = nfree;
        nm.hblkhd      = m->footprint - sum;
3312     nm.usmblks      = m->max_footprint;
        nm.uordblks    = m->footprint - mfree;
        nm.fordblks    = mfree;
        nm.keepcost    = m->topsize;
3316     }

    POSTACTION(m);
    }
3320 return nm;
    }

```

```

#endif /* !NO_MALLINFO */

3324 static void internal_malloc_stats(mstate m) {
    ensure_initialization();
    if (!PREACTION(m)) {
        size_t maxfp = 0;
3328         size_t fp = 0;
        size_t used = 0;
        check_malloc_state(m);
        if (is_initialized(m)) {
3332             msegmentptr s = &m->seg;
            maxfp = m->max_footprint;
            fp = m->footprint;
            used = fp - (m->topsize + TOP_FOOT_SIZE);

            while (s != 0) {
                mchunkptr q = align_as_chunk(s->base);
                while (segment_holds(s, q) &&
3340                    q != m->top && q->head != FENCEPOST_HEAD) {
                    if (!is_inuse(q))
                        used -= chunksize(q);
                    q = next_chunk(q);
3344                }
                s = s->next;
            }
        }

        fprintf(stderr, "max system bytes = %10lu\n", (unsigned long)(maxfp));
        fprintf(stderr, "system bytes      = %10lu\n", (unsigned long)(fp));
        fprintf(stderr, "in use bytes       = %10lu\n", (unsigned long)(used));

        POSTACTION(m);
    }
}

```

Chapter 20

Operations on smallbins and trees

```
3356  /*
      Various forms of linking and unlinking are defined as macros.  Even
      the ones for trees, which are very long but have very short typical
      paths.  This is ugly but reduces reliance on inlining support of
3360  compilers.
      */
```

20.1 Operations on smallbins

```
      /* Link a free chunk into a smallbin */
3364  #define insert_small_chunk(M, P, S) {\
      bindex_t I = small_index(S);\
      mchunkptr B = smallbin_at(M, I);\
      mchunkptr F = B;\
3368  assert(S >= MIN_CHUNK_SIZE);\
      if (!smallmap_is_marked(M, I))\
          mark_smallmap(M, I);\
      else if (RTCHECK(ok_address(M, B->fd)))\
3372  F = B->fd;\
      else {\
          CORRUPTION_ERROR_ACTION(M);\
      }\
3376  B->fd = P;\
      F->bk = P;\
      P->fd = F;\
      P->bk = B;\
3380  }

      /* Unlink a chunk from a smallbin */
      #define unlink_small_chunk(M, P, S) {\
3384  mchunkptr F = P->fd;\
      mchunkptr B = P->bk;\
      bindex_t I = small_index(S);\
      assert(P != B);\
3388  assert(P != F);\
      assert(chunksize(P) == small_index2size(I));\
      if (F == B)\
          clear_smallmap(M, I);\
3392  else if (RTCHECK((F == smallbin_at(M, I) || ok_address(M, F)) &&\
```

```

        (B == smallbin_at(M,I) || ok_address(M, B)))) {\
    F->bk = B;\
    B->fd = F;\
3396 } \
    else {\
        CORRUPTION_ERROR_ACTION(M);\
    } \
3400 }

/* Unlink the first chunk from a smallbin */
#define unlink_first_small_chunk(M, B, P, I) {\
3404   mchunkptr F = P->fd;\
    assert(P != B);\
    assert(P != F);\
    assert(chunksize(P) == small_index2size(I));\
3408   if (B == F)\
        clear_smallmap(M, I);\
    else if (RTCHECK(ok_address(M, F))) {\
        B->fd = F;\
3412   F->bk = B;\
    } \
    else {\
        CORRUPTION_ERROR_ACTION(M);\
3416 } \
}

/* Replace dv node, binning the old one */
/* Used only when dvsize known to be small */
#define replace_dv(M, P, S) {\
3424   size_t DVS = M->dvsize;\
    if (DVS != 0) {\
        mchunkptr DV = M->dv;\
        assert(is_small(DVS));\
3428   insert_small_chunk(M, DV, DVS);\
    } \
    M->dvsize = S;\
    M->dv = P;\
3432 }

```

20.2 Operations on trees

```

/* Insert chunk into tree */
3436 #define insert_large_chunk(M, X, S) {\
    tbinptr* H;\
    bindex_t I;\
    compute_tree_index(S, I);\
3440   H = treebin_at(M, I);\
    X->index = I;\
    X->child[0] = X->child[1] = 0;\
}

```

```

if (!treemap_is_marked(M, I)) {\
3444   mark_treemap(M, I);\
   *H = X;\
   X->parent = (tchunkptr)H;\
   X->fd = X->bk = X;\
3448 }\
else {\
   tchunkptr T = *H;\
   size_t K = S << leftshift_for_tree_index(I);\
3452   for (;;) {\
       if (chunksize(T) != S) {\
           tchunkptr* C = &(T->child[(K >> (SIZE_T_BITSIZE-SIZE_T_ONE)) & 1]);\
           K <= 1;\
3456           if (*C != 0)\
               T = *C;\
           else if (RTCHECK(ok_address(M, C))) {\
               *C = X;\
3460               X->parent = T;\
               X->fd = X->bk = X;\
               break;\
           }\
3464           else {\
               CORRUPTION_ERROR_ACTION(M);\
               break;\
           }\
3468   }\
   else {\
       tchunkptr F = T->fd;\
       if (RTCHECK(ok_address(M, T) && ok_address(M, F))) {\
3472           T->fd = F->bk = X;\
           X->fd = F;\
           X->bk = T;\
           X->parent = 0;\
3476           break;\
       }\
       else {\
           CORRUPTION_ERROR_ACTION(M);\
3480           break;\
       }\
   }\
}\
3484 }\
}

/*
3488   Unlink steps:

```

1. If *x* is a chained node, unlink it from its same-sized fd/bk links and choose its bk node as its replacement.
- 3492 2. If *x* was the last node of its size, but not a leaf node, it must be replaced with a leaf node (not merely one with an open left or right), to make sure that lefts and rights of descendants

```

correspond properly to bit masks. We use the rightmost descendent
3496 of x. We could use any other leaf, but this is easy to locate and
tends to counteract removal of leftmosts elsewhere, and so keeps
paths shorter than minimally guaranteed. This doesn't loop much
because on average a node in a tree is near the bottom.
3500 3. If x is the base of a chain (i.e., has parent links) relink
x's parent and children to x's replacement (or null if none).
*/

3504 #define unlink_large_chunk(M, X) {\
    tchunkptr XP = X->parent;\
    tchunkptr R;\
    if (X->bk != X) {\
3508     tchunkptr F = X->fd;\
    R = X->bk;\
    if (RTCHECK(ok_address(M, F))) {\
        F->bk = R;\
3512     R->fd = F;\
    }\
    else {\
        CORRUPTION_ERROR_ACTION(M);\
3516    }\
    }\
    else {\
        tchunkptr* RP;\
3520     if (((R = *(RP = &(X->child[1]))) != 0) ||\
        ((R = *(RP = &(X->child[0]))) != 0)) {\
        tchunkptr* CP;\
        while ((*CP = &(R->child[1])) != 0) ||\
3524         ((*CP = &(R->child[0])) != 0)) {\
            R = *(RP = CP);\
        }\
        if (RTCHECK(ok_address(M, RP)))\
3528         *RP = 0;\
        else {\
            CORRUPTION_ERROR_ACTION(M);\
        }\
3532    }\
    }\
    if (XP != 0) {\
        tbinptr* H = treebin_at(M, X->index);\
3536     if (X == *H) {\
        if ((*H = R) == 0) \
            clear_treemap(M, X->index);\
    }\
    else if (RTCHECK(ok_address(M, XP))) {\
        if (XP->child[0] == X) \
            XP->child[0] = R;\
        else \
3544         XP->child[1] = R;\
    }\
    else\

```

```

CORRUPTION_ERROR_ACTION(M);\
3548 if (R != 0) {\
    if (RTCHECK(ok_address(M, R))) {\
        tchunkptr C0, C1;\
        R->parent = XP;\
3552    if ((C0 = X->child[0]) != 0) {\
        if (RTCHECK(ok_address(M, C0))) {\
            R->child[0] = C0;\
            C0->parent = R;\
3556        }\
        else\
            CORRUPTION_ERROR_ACTION(M);\
    }\
3560    if ((C1 = X->child[1]) != 0) {\
        if (RTCHECK(ok_address(M, C1))) {\
            R->child[1] = C1;\
            C1->parent = R;\
3564        }\
        else\
            CORRUPTION_ERROR_ACTION(M);\
    }\
3568    }\
    else\
        CORRUPTION_ERROR_ACTION(M);\
    }\
3572 }\
}

/* Relays to large vs small bin operations */

#define insert_chunk(M, P, S)\
    if (is_small(S)) insert_small_chunk(M, P, S)\
    else { tchunkptr TP = (tchunkptr)(P); insert_large_chunk(M, TP, S); }

#define unlink_chunk(M, P, S)\
    if (is_small(S)) unlink_small_chunk(M, P, S)\
    else { tchunkptr TP = (tchunkptr)(P); unlink_large_chunk(M, TP); }

/* Relays to internal calls to malloc/free from realloc, memalign etc */

3588 #if ONLY_MSPACES
#define internal_malloc(m, b) mspace_malloc(m, b)
#define internal_free(m, mem) mspace_free(m, mem);
#else /* ONLY_MSPACES */
3592 #if MSPACES
#define internal_malloc(m, b)\
    (m == gm)? dlmalloc(b) : mspace_malloc(m, b)
#define internal_free(m, mem)\
3596    if (m == gm) dlfree(mem); else mspace_free(m, mem);
#else /* MSPACES */
#define internal_malloc(m, b) dlmalloc(b)

```



```
    #define internal_free(m, mem) dlfree(mem)
3600 #endif /* MSPACES */
    #endif /* ONLY_MSPACES */
```

Chapter 21

Direct-mmapping chunks

```
/*
  Directly mmaped chunks are set up with an offset to the start of
3604 the mmaped region stored in the prev_foot field of the chunk. This
    allows reconstruction of the required argument to MUNMAP when freed,
    and also allows adjustment of the returned chunk to meet alignment
    requirements (especially in memalign).
3608 */

/* Malloc using mmap */
static void* mmap_alloc(mstate m, size_t nb) {
3612   size_t mmsize = mmap_align(nb + SIX_SIZE_T_SIZES + CHUNK_ALIGN_MASK);
    if (mmsize > nb) { /* Check for wrap around 0 */
        char* mm = (char*)(CALL_DIRECT_MMAP(mmsize));
        if (mm != CMFAIL) {
3616             size_t offset = align_offset(chunk2mem(mm));
                size_t psize = mmsize - offset - MMAP_FOOT_PAD;
                mchunkptr p = (mchunkptr)(mm + offset);
                p->prev_foot = offset;
3620             p->head = psize;
                mark_inuse_foot(m, p, psize);
                chunk_plus_offset(p, psize)->head = FENCEPOST_HEAD;
                chunk_plus_offset(p, psize+SIZE_T_SIZE)->head = 0;

                if (m->least_addr == 0 || mm < m->least_addr)
                    m->least_addr = mm;
                if ((m->footprint += mmsize) > m->max_footprint)
3628                 m->max_footprint = m->footprint;
                assert(is_aligned(chunk2mem(p)));
                check_mmapped_chunk(m, p);
                return chunk2mem(p);
3632         }
    }
    return 0;
}

/* Realloc using mmap */
static mchunkptr mmap_resize(mstate m, mchunkptr oldp, size_t nb) {
    size_t oldsize = chunksize(oldp);
3640     if (is_small(nb)) /* Can't shrink mmap regions below small size */
```

```

    return 0;
    /* Keep old chunk if big enough but not too big */
    if (oldsize >= nb + SIZE_T_SIZE &&
3644      (oldsize - nb) <= (mparams.granularity << 1))
        return oldp;
    else {
        size_t offset = oldp->prev_foot;
3648      size_t oldmmsize = oldsize + offset + MMAP_FOOT_PAD;
        size_t newmmsize = mmap_align(nb + SIX_SIZE_T_SIZES + CHUNK_ALIGN_MASK);
        char* cp = (char*)CALL_MREMAP((char*)oldp - offset,
                                      oldmmsize, newmmsize, 1);
3652      if (cp != CMFAIL) {
          mchunkptr newp = (mchunkptr)(cp + offset);
          size_t psize = newmmsize - offset - MMAP_FOOT_PAD;
          newp->head = psize;
3656          mark_inuse_foot(m, newp, psize);
          chunk_plus_offset(newp, psize)->head = FENCEPOST_HEAD;
          chunk_plus_offset(newp, psize+SIZE_T_SIZE)->head = 0;

3660          if (cp < m->least_addr)
              m->least_addr = cp;
          if ((m->footprint += newmmsize - oldmmsize) > m->max_footprint)
              m->max_footprint = m->footprint;
3664          check_mmapped_chunk(m, newp);
          return newp;
      }
    }
3668  return 0;
}

```

Chapter 22

mspace management

```
/* Initialize top chunk and its size */
static void init_top(mstate m, mchunkptr p, size_t psize) {
3672  /* Ensure alignment */
    size_t offset = align_offset(chunk2mem(p));
    p = (mchunkptr)((char*)p + offset);
    psize -= offset;

    m->top = p;
    m->topsize = psize;
    p->head = psize | PINUSE_BIT;
3680  /* set size of fake trailing chunk holding overhead space only once */
    chunk_plus_offset(p, psize)->head = TOP_FOOT_SIZE;
    m->trim_check = mparams.trim_threshold; /* reset on each update */
}

/* Initialize bins for a new mstate that is otherwise zeroed out */
static void init_bins(mstate m) {
    /* Establish circular links for smallbins */
3688  bindex_t i;
    for (i = 0; i < NSMALLBINS; ++i) {
        sbindex_t bin = smallbin_at(m,i);
        bin->fd = bin->bk = bin;
3692  }
}

#if PROCEED_ON_ERROR

/* default corruption action */
static void reset_on_error(mstate m) {
    int i;
3700  ++malloc_corruption_error_count;
    /* Reinitialize fields to forget about all memory */
    m->smallbins = m->treebins = 0;
    m->dvsize = m->topsize = 0;
3704  m->seg.base = 0;
    m->seg.size = 0;
    m->seg.next = 0;
    m->top = m->dv = 0;
3708  for (i = 0; i < NTREEBINS; ++i)
```

```

        *treebin_at(m, i) = 0;
        init_bins(m);
    }
3712 #endif /* PROCEED_ON_ERROR */

/* Allocate chunk and prepend remainder with chunk in successor base. */
static void* prepend_alloc(mstate m, char* newbase, char* oldbase,
3716                          size_t nb) {
    mchunkptr p = align_as_chunk(newbase);
    mchunkptr oldfirst = align_as_chunk(oldbase);
    size_t psize = (char*)oldfirst - (char*)p;
3720 mchunkptr q = chunk_plus_offset(p, nb);
    size_t qsize = psize - nb;
    set_size_and_pinuse_of_inuse_chunk(m, p, nb);

3724 assert((char*)oldfirst > (char*)q);
    assert(pinuse(oldfirst));
    assert(qsize >= MIN_CHUNK_SIZE);

3728 /* consolidate remainder with first chunk of old base */
    if (oldfirst == m->top) {
        size_t tsize = m->topsize += qsize;
        m->top = q;
3732 q->head = tsize | PINUSE_BIT;
        check_top_chunk(m, q);
    }
    else if (oldfirst == m->dv) {
3736 size_t dsize = m->dvsizes += qsize;
        m->dv = q;
        set_size_and_pinuse_of_free_chunk(q, dsize);
    }
    else {
3740 if (!is_inuse(oldfirst)) {
        size_t nsize = chunksize(oldfirst);
        unlink_chunk(m, oldfirst, nsize);
3744 oldfirst = chunk_plus_offset(oldfirst, nsize);
        qsize += nsize;
    }
    set_free_with_pinuse(q, qsize, oldfirst);
3748 insert_chunk(m, q, qsize);
    check_free_chunk(m, q);
}

3752 check_mallocated_chunk(m, chunk2mem(p), nb);
    return chunk2mem(p);
}

3756 /* Add a segment to hold a new noncontiguous region */
static void add_segment(mstate m, char* tbase, size_t tsize, flag_t mmapped) {
    /* Determine locations and sizes of segment, fenceposts, old top */
    char* old_top = (char*)m->top;
3760 msegmentptr oldsp = segment_holding(m, old_top);

```

```

char* old_end = oldsp->base + oldsp->size;
size_t ssize = pad_request(sizeof(struct malloc_segment));
char* rawsp = old_end - (ssize + FOUR_SIZE_T_SIZES + CHUNK_ALIGN_MASK);
3764 size_t offset = align_offset(chunk2mem(rawsp));
char* asp = rawsp + offset;
char* csp = (asp < (old_top + MIN_CHUNK_SIZE)) ? old_top : asp;
mchunkptr sp = (mchunkptr)csp;
3768 msegmentptr ss = (msegmentptr)(chunk2mem(sp));
mchunkptr tnext = chunk_plus_offset(sp, ssize);
mchunkptr p = tnext;
int nfences = 0;

/* reset top to new space */
init_top(m, (mchunkptr)tbase, tsize - TOP_FOOT_SIZE);

3776 /* Set up segment record */
assert(is_aligned(ss));
set_size_and_pinuse_of_inuse_chunk(m, sp, ssize);
*ss = m->seg; /* Push current record */
3780 m->seg.base = tbase;
m->seg.size = tsize;
m->seg.sflags = mmapped;
m->seg.next = ss;

/* Insert trailing fenceposts */
for (;;) {
    mchunkptr nextp = chunk_plus_offset(p, SIZE_T_SIZE);
3788 p->head = FENCEPOST_HEAD;
    ++nfences;
    if ((char*)&(nextp->head) < old_end)
        p = nextp;
3792 else
        break;
}
assert(nfences >= 2);

/* Insert the rest of old top into a bin as an ordinary free chunk */
if (csp != old_top) {
    mchunkptr q = (mchunkptr)old_top;
3800 size_t psize = csp - old_top;
    mchunkptr tn = chunk_plus_offset(q, psize);
    set_free_with_pinuse(q, psize, tn);
    insert_chunk(m, q, psize);
3804 }

check_top_chunk(m, m->top);
}

```

Chapter 23

System allocation and deallocation

3808

23.1 System allocation

```
/* Get memory from system using MORECORE or MMAP */
static void* sys_alloc(mstate m, size_t nb) {
    char* tbase = CMFAIL;
3812    size_t tsize = 0;
    flag_t mmap_flag = 0;

    ensure_initialization();

    /* Directly map large chunks, but only if already initialized */
    if (use_mmap(m) && nb >= mparams.mmap_threshold && m->topsize != 0) {
        void* mem = mmap_alloc(m, nb);
3820        if (mem != 0)
            return mem;
    }

3824    /*
    Try getting memory in any of three ways (in most-preferred to
    least-preferred order):
    1. A call to MORECORE that can normally contiguously extend memory.
    (disabled if not MORECORE_CONTIGUOUS or not HAVE_MORECORE or
    or main space is mmapped or a previous contiguous call failed)
    2. A call to MMAP new space (disabled if not HAVE_MMAP).
    Note that under the default settings, if MORECORE is unable to
    fulfill a request, and HAVE_MMAP is true, then mmap is
3832    used as a noncontiguous system allocator. This is a useful backup
    strategy for systems with holes in address spaces -- in this case
    sbrk cannot contiguously expand the heap, but mmap may be able to
    find space.
3836    3. A call to MORECORE that cannot usually contiguously extend memory.
    (disabled if not HAVE_MORECORE)

3840    In all cases, we need to request enough bytes from system to ensure
    we can malloc nb bytes upon success, so pad with enough space for
    top_foot, plus alignment-pad to make sure we don't lose bytes if
    not on boundary, and round this up to a granularity unit.
3844    */
```

```

if (MORECORE_CONTIGUOUS && !use_noncontiguous(m)) {
    char* br = CMFAIL;
3848    msegmentptr ss = (m->top == 0)? 0 : segment_holding(m, (char*)m->top);
    size_t asize = 0;
    ACQUIRE_MALLOC_GLOBAL_LOCK();

3852    if (ss == 0) { /* First time through or recovery */
        char* base = (char*)CALL_MORECORE(0);
        if (base != CMFAIL) {
            asize = granularity_align(nb + SYS_ALLOC_PADDING);
3856            /* Adjust to end on a page boundary */
            if (!is_page_aligned(base))
                asize += (page_align((size_t)base) - (size_t)base);
            /* Can't call MORECORE if size is negative when treated as signed */
3860            if (asize < HALF_MAX_SIZE_T &&
                (br = (char*)(CALL_MORECORE(asize))) == base) {
                tbase = base;
                tsize = asize;
3864            }
        }
    }
    else {
3868        /* Subtract out existing available top space from MORECORE request. */
        asize = granularity_align(nb - m->topsize + SYS_ALLOC_PADDING);
        /* Use mem here only if it did continuously extend old space */
        if (asize < HALF_MAX_SIZE_T &&
3872            (br = (char*)(CALL_MORECORE(asize))) == ss->base+ss->size) {
            tbase = br;
            tsize = asize;
        }
3876    }

    if (tbase == CMFAIL) { /* Cope with partial failure */
        if (br != CMFAIL) { /* Try to use/extend the space we did get */
3880            if (asize < HALF_MAX_SIZE_T &&
                asize < nb + SYS_ALLOC_PADDING) {
                size_t esize = granularity_align(nb + SYS_ALLOC_PADDING - asize);
                if (esize < HALF_MAX_SIZE_T) {
3884                    char* end = (char*)CALL_MORECORE(esize);
                    if (end != CMFAIL)
                        asize += esize;
                    else { /* Can't use; try to release */
3888                        (void) CALL_MORECORE(-asize);
                        br = CMFAIL;
                    }
                }
            }
        }
3892    }
    if (br != CMFAIL) { /* Use the space we did get */
        tbase = br;
3896        tsize = asize;
    }
}

```



```

    }
    else
        disable_contiguous(m); /* Don't try contiguous path in the future */
3900 }

    RELEASE_MALLOC_GLOBAL_LOCK();
}

if (HAVE_MMAP && tbase == CMFAIL) { /* Try MMAP */
    size_t rsize = granularity_align(nb + SYS_ALLOC_PADDING);
    if (rsize > nb) { /* Fail if wraps around zero */
3908     char* mp = (char*)(CALL_MMAP(rsize));
        if (mp != CMFAIL) {
            tbase = mp;
            tsize = rsize;
3912     mmap_flag = USE_MMAP_BIT;
        }
    }
}

if (HAVE_MORECORE && tbase == CMFAIL) { /* Try noncontiguous MORECORE */
    size_t asize = granularity_align(nb + SYS_ALLOC_PADDING);
    if (asize < HALF_MAX_SIZE_T) {
3920     char* br = CMFAIL;
        char* end = CMFAIL;
        ACQUIRE_MALLOC_GLOBAL_LOCK();
        br = (char*)(CALL_MORECORE(asize));
3924     end = (char*)(CALL_MORECORE(0));
        RELEASE_MALLOC_GLOBAL_LOCK();
        if (br != CMFAIL && end != CMFAIL && br < end) {
            size_t ssize = end - br;
3928     if (ssize > nb + TOP_FOOT_SIZE) {
                tbase = br;
                tsize = ssize;
            }
3932     }
    }
}

3936 if (tbase != CMFAIL) {

    if ((m->footprint += tsize) > m->max_footprint)
        m->max_footprint = m->footprint;

    if (!is_initialized(m)) { /* first-time initialization */
        if (m->least_addr == 0 || tbase < m->least_addr)
            m->least_addr = tbase;
3944     m->seg.base = tbase;
        m->seg.size = tsize;
        m->seg.sflags = mmap_flag;
        m->magic = mparams.magic;
3948     m->release_checks = MAX_RELEASE_CHECK_RATE;

```

```

        init_bins(m);
    #if !ONLY_MSPACES
        if (is_global(m))
3952         init_top(m, (mchunkptr)tbase, tsize - TOP_FOOT_SIZE);
        else
    #endif
    {
3956         /* Offset top by embedded malloc_state */
        mchunkptr mn = next_chunk(mem2chunk(m));
        init_top(m, mn, (size_t)((tbase + tsize) - (char*)mn) - TOP_FOOT_SIZE);
    }
3960 }

    else {
        /* Try to merge with an existing segment */
3964 msegmentptr sp = &m->seg;
        /* Only consider most recent segment if traversal suppressed */
        while (sp != 0 && tbase != sp->base + sp->size)
            sp = (NO_SEGMENT_TRAVERSAL) ? 0 : sp->next;
3968 if (sp != 0 &&
            !is_extern_segment(sp) &&
            (sp->sflags & USE_MMAP_BIT) == mmap_flag &&
            segment_holds(sp, m->top)) { /* append */
3972     sp->size += tsize;
            init_top(m, m->top, m->topsize + tsize);
        }
        else {
3976         if (tbase < m->least_addr)
            m->least_addr = tbase;
            sp = &m->seg;
            while (sp != 0 && sp->base != tbase + tsize)
3980             sp = (NO_SEGMENT_TRAVERSAL) ? 0 : sp->next;
            if (sp != 0 &&
                !is_extern_segment(sp) &&
                (sp->sflags & USE_MMAP_BIT) == mmap_flag) {
3984             char* oldbase = sp->base;
                sp->base = tbase;
                sp->size += tsize;
                return prepend_alloc(m, tbase, oldbase, nb);
3988             }
            else
                add_segment(m, tbase, tsize, mmap_flag);
        }
3992 }

    if (nb < m->topsize) { /* Allocate from new or extended top space */
        size_t rsize = m->topsize - nb;
3996 mchunkptr p = m->top;
        mchunkptr r = m->top = chunk_plus_offset(p, nb);
        r->head = rsize | PINUSE_BIT;
        set_size_and_pinuse_of_inuse_chunk(m, p, nb);
4000 check_top_chunk(m, m->top);

```

```

        check_malloced_chunk(m, chunk2mem(p), nb);
        return chunk2mem(p);
    }
4004 }

    MALLOC_FAILURE_ACTION;
    return 0;
4008 }

```

23.2 System deallocation

```

/* Unmap and unlink any mmaped segments that don't contain used chunks */
4012 static size_t release_unused_segments(mstate m) {
    size_t released = 0;
    int nsegs = 0;
    msegmentptr pred = &m->seg;
4016 msegmentptr sp = pred->next;
    while (sp != 0) {
        char* base = sp->base;
        size_t size = sp->size;
4020 msegmentptr next = sp->next;
        ++nsegs;
        if (is_mmaped_segment(sp) && !is_extern_segment(sp)) {
            mchunkptr p = align_as_chunk(base);
4024 size_t psize = chunksize(p);
            /* Can unmap if first chunk holds entire segment and not pinned */
            if (!is_inuse(p) && (char*)p + psize >= base + size - TOP_FOOT_SIZE) {
                tchunkptr tp = (tchunkptr)p;
4028 assert(segment_holds(sp, (char*)sp));
                if (p == m->dv) {
                    m->dv = 0;
                    m->dvsizesize = 0;
4032 }
                else {
                    unlink_large_chunk(m, tp);
                }
4036 if (CALL_MUNMAP(base, size) == 0) {
                    released += size;
                    m->footprint -= size;
                    /* unlink obsoleted record */
4040 sp = pred;
                    sp->next = next;
                }
                else { /* back out if cannot unmap */
4044 insert_large_chunk(m, tp, psize);
                }
            }
        }
    }
4048 if (NO_SEGMENT_TRAVERSAL) /* scan only first segment */
    break;
    pred = sp;

```

```

    sp = next;
4052 }
    /* Reset check counter */
    m->release_checks = ((nsecs > MAX_RELEASE_CHECK_RATE)?
                        nsecs : MAX_RELEASE_CHECK_RATE);
4056 return released;
}

static int sys_trim(mstate m, size_t pad) {
4060 size_t released = 0;
    ensure_initialization();
    if (pad < MAX_REQUEST && is_initialized(m)) {
        pad += TOP_FOOT_SIZE; /* ensure enough room for segment overhead */

        if (m->topsize > pad) {
            /* Shrink top space in granularity-size units, keeping at least one */
            size_t unit = mparams.granularity;
4068 size_t extra = ((m->topsize - pad + (unit - SIZE_T_ONE)) / unit -
                    SIZE_T_ONE) * unit;
            msegmentptr sp = segment_holding(m, (char*)m->top);

            if (!is_extern_segment(sp)) {
                if (is_mmapped_segment(sp)) {
                    if (HAVE_MMAP &&
4076 sp->size >= extra &&
                        !has_segment_link(m, sp)) { /* can't shrink if pinned */
                            size_t newsize = sp->size - extra;
                            /* Prefer mremap, fall back to munmap */
                            if (((CALL_MREMAP(sp->base, sp->size, newsize, 0) != MFAIL) ||
4080 (CALL_MUNMAP(sp->base + newsize, extra) == 0)) {
                                released = extra;
                            }
                        }
                    }
                }
4084 }
            else if (HAVE_MORECORE) {
                if (extra >= HALF_MAX_SIZE_T) /* Avoid wrapping negative */
                    extra = (HALF_MAX_SIZE_T) + SIZE_T_ONE - unit;
                ACQUIRE_MALLOC_GLOBAL_LOCK();
                {
                    /* Make sure end of memory is where we last set it. */
                    char* old_br = (char*)(CALL_MORECORE(0));
4092 if (old_br == sp->base + sp->size) {
                        char* rel_br = (char*)(CALL_MORECORE(-extra));
                        char* new_br = (char*)(CALL_MORECORE(0));
                        if (rel_br != CMFAIL && new_br < old_br)
4096 released = old_br - new_br;
                    }
                }
                RELEASE_MALLOC_GLOBAL_LOCK();
4100 }
    }
}

```

```
        if (released != 0) {
4104         sp->size -= released;
            m->footprint -= released;
            init_top(m, m->top, m->topsize - released);
            check_top_chunk(m, m->top);
4108     }
    }

    /* Unmap any unused mmapped segments */
4112     if (HAVE_MMAP)
        released += release_unused_segments(m);

    /* On failure, disable autotrim to avoid repeated failed future calls */
4116     if (released == 0 && m->topsize > m->trim_check)
        m->trim_check = MAX_SIZE_T;
    }

4120     return (released != 0)? 1 : 0;
}
```

Chapter 24

Support for public routines

24.1 malloc support

```
4124 /* allocate a large request from the best fitting chunk in a treebin */
static void* tmalloc_large(mstate m, size_t nb) {
    tchunkptr v = 0;
    size_t rsize = -nb; /* Unsigned negation */
4128 tchunkptr t;
    bindindex_t idx;
    compute_tree_index(nb, idx);
    if ((t = *treebin_at(m, idx)) != 0) {
4132 /* Traverse tree for this bin looking for node with size == nb */
        size_t sizebits = nb << leftshift_for_tree_index(idx);
        tchunkptr rst = 0; /* The deepest untaken right subtree */
        for (;;) {
4136 tchunkptr rt;
            size_t trem = chunksize(t) - nb;
            if (trem < rsize) {
                v = t;
4140 if ((rsize = trem) == 0)
                    break;
            }
            rt = t->child[1];
4144 t = t->child[(sizebits >> (SIZE_T_BITSIZE-SIZE_T_ONE)) & 1];
            if (rt != 0 && rt != t)
                rst = rt;
            if (t == 0) {
4148 t = rst; /* set t to least subtree holding sizes > nb */
                break;
            }
            sizebits <=> 1;
4152 }
        }
    }
    if (t == 0 && v == 0) { /* set t to root of next non-empty treebin */
        binmap_t leftbits = left_bits(idx2bit(idx)) & m->treemap;
4156 if (leftbits != 0) {
            bindindex_t i;
            binmap_t leastbit = least_bit(leftbits);
            compute_bit2idx(leastbit, i);
```

```

4160     t = *treebin_at(m, i);
    }
}

4164 while (t != 0) { /* find smallest of tree or subtree */
    size_t trem = chunksize(t) - nb;
    if (trem < rsize) {
        rsize = trem;
4168     v = t;
    }
    t = leftmost_child(t);
}

/* If dv is a better fit, return 0 so malloc will use it */
if (v != 0 && rsize < (size_t)(m->dvsizes - nb)) {
    if (RTCHECK(ok_address(m, v))) { /* split */
4176     mchunkptr r = chunk_plus_offset(v, nb);
    assert(chunksize(v) == rsize + nb);
    if (RTCHECK(ok_next(v, r))) {
        unlink_large_chunk(m, v);
4180     if (rsize < MIN_CHUNK_SIZE)
        set_inuse_and_pinuse(m, v, (rsize + nb));
    else {
        set_size_and_pinuse_of_inuse_chunk(m, v, nb);
4184     set_size_and_pinuse_of_free_chunk(r, rsize);
        insert_chunk(m, r, rsize);
    }
    return chunk2mem(v);
4188 }
}
CORRUPTION_ERROR_ACTION(m);
}
4192 return 0;
}

/* allocate a small request from the best fitting chunk in a treebin */
4196 static void* tmalloc_small(mstate m, size_t nb) {
    tchunkptr t, v;
    size_t rsize;
    bindindex_t i;
4200 binmap_t leastbit = least_bit(m->treemap);
    compute_bit2idx(leastbit, i);
    v = t = *treebin_at(m, i);
    rsize = chunksize(t) - nb;

    while ((t = leftmost_child(t)) != 0) {
        size_t trem = chunksize(t) - nb;
        if (trem < rsize) {
4208     rsize = trem;
            v = t;
        }
    }
}

```

```

    if (RTCHECK(ok_address(m, v))) {
        mchunkptr r = chunk_plus_offset(v, nb);
        assert(chunksize(v) == rsize + nb);
4216    if (RTCHECK(ok_next(v, r))) {
            unlink_large_chunk(m, v);
            if (rsize < MIN_CHUNK_SIZE)
                set_inuse_and_pinuse(m, v, (rsize + nb));
4220            else {
                set_size_and_pinuse_of_inuse_chunk(m, v, nb);
                set_size_and_pinuse_of_free_chunk(r, rsize);
                replace_dv(m, r, rsize);
4224            }
            return chunk2mem(v);
        }
    }

    CORRUPTION_ERROR_ACTION(m);
    return 0;
}

```

24.2 realloc support

```

static void* internal_realloc(mstate m, void* oldmem, size_t bytes) {
    if (bytes >= MAX_REQUEST) {
4236        MALLOC_FAILURE_ACTION;
        return 0;
    }
    if (!PREACTION(m)) {
4240        mchunkptr oldp = mem2chunk(oldmem);
        size_t oldsize = chunksize(oldp);
        mchunkptr next = chunk_plus_offset(oldp, oldsize);
        mchunkptr newp = 0;
4244        void* extra = 0;

        /* Try to either shrink or extend into top. Else malloc-copy-free */

4248        if (RTCHECK(ok_address(m, oldp) && ok_inuse(oldp) &&
                        ok_next(oldp, next) && ok_pinuse(next))) {
            size_t nb = request2size(bytes);
            if (is_mmapped(oldp))
4252                newp = mmap_resize(m, oldp, nb);
            else if (oldsize >= nb) { /* already big enough */
                size_t rsize = oldsize - nb;
                newp = oldp;
4256                if (rsize >= MIN_CHUNK_SIZE) {
                    mchunkptr remainder = chunk_plus_offset(newp, nb);
                    set_inuse(m, newp, nb);
                    set_inuse_and_pinuse(m, remainder, rsize);
4260                    extra = chunk2mem(remainder);
                }
            }
        }
    }
}

```



```

    }
    else if (next == m->top && oldsize + m->topsize > nb) {
4264      /* Expand into top */
        size_t newsize = oldsize + m->topsize;
        size_t newtopsize = newsize - nb;
        mchunkptr newtop = chunk_plus_offset(oldp, nb);
4268      set_inuse(m, oldp, nb);
        newtop->head = newtopsize | PINUSE_BIT;
        m->top = newtop;
        m->topsize = newtopsize;
4272      newp = oldp;
    }
}
else {
4276   USAGE_ERROR_ACTION(m, oldmem);
   POSTACTION(m);
   return 0;
}
4280 #if DEBUG
    if (newp != 0) {
        check_inuse_chunk(m, newp); /* Check requires lock */
    }
4284 #endif

    POSTACTION(m);

4288   if (newp != 0) {
        if (extra != 0) {
            internal_free(m, extra);
        }
4292   return chunk2mem(newp);
    }
    else {
        void* newmem = internal_malloc(m, bytes);
4296   if (newmem != 0) {
            size_t oc = oldsize - overhead_for(oldp);
            memcpy(newmem, oldmem, (oc < bytes)? oc : bytes);
            internal_free(m, oldmem);
4300   }
        return newmem;
    }
}
4304 return 0;
}

```

24.3 memalign support

```

static void* internal_memalign(mstate m, size_t alignment, size_t bytes) {
    if (alignment <= MALLOC_ALIGNMENT) /* Can just use malloc */
        return internal_malloc(m, bytes);
}

```

```

4312 if (alignment < MIN_CHUNK_SIZE) /* must be at least a minimum chunk size */
    alignment = MIN_CHUNK_SIZE;
if ((alignment & (alignment-SIZE_T_ONE)) != 0) { /* Ensure a power of 2 */
    size_t a = MALLOC_ALIGNMENT << 1;
4316 while (a < alignment) a <<= 1;
    alignment = a;
}

4320 if (bytes >= MAX_REQUEST - alignment) {
    if (m != 0) { /* Test isn't needed but avoids compiler warning */
        MALLOC_FAILURE_ACTION;
    }
4324 }
else {
    size_t nb = request2size(bytes);
    size_t req = nb + alignment + MIN_CHUNK_SIZE - CHUNK_OVERHEAD;
4328 char* mem = (char*)internal_malloc(m, req);
    if (mem != 0) {
        void* leader = 0;
        void* trailer = 0;
4332 mchunkptr p = mem2chunk(mem);

        if (PREACTION(m)) return 0;
        if (((size_t)(mem)) % alignment) != 0) { /* misaligned */
4336 /*
            Find an aligned spot inside chunk. Since we need to give
            back leading space in a chunk of at least MIN_CHUNK_SIZE, if
            the first calculation places us at a spot with less than
4340 MIN_CHUNK_SIZE leader, we can move to the next aligned spot.
            We've allocated enough total room so that this is always
            possible.
        */
4344 char* br = (char*)mem2chunk((size_t)(((size_t)(mem +
                                                alignment -
                                                SIZE_T_ONE)) &
                                                -alignment));

        char* pos = ((size_t)(br - (char*)(p)) >= MIN_CHUNK_SIZE)?
            br : br+alignment;
        mchunkptr newp = (mchunkptr)pos;
        size_t leadsize = pos - (char*)(p);
4352 size_t newsize = chunksize(p) - leadsize;

        if (is_mmapped(p)) { /* For mmaped chunks, just adjust offset */
            newp->prev_foot = p->prev_foot + leadsize;
4356 newp->head = newsize;
        }
        else { /* Otherwise, give back leader, use the rest */
            set_inuse(m, newp, newsize);
            set_inuse(m, p, leadsize);
4360 leader = chunk2mem(p);
        }
        p = newp;

```

```

4364     }

    /* Give back spare room at the end */
    if (!is_mmapped(p)) {
4368         size_t size = chunksize(p);
        if (size > nb + MIN_CHUNK_SIZE) {
            size_t remainder_size = size - nb;
            mchunkptr remainder = chunk_plus_offset(p, nb);
4372             set_inuse(m, p, nb);
            set_inuse(m, remainder, remainder_size);
            trailer = chunk2mem(remainder);
        }
4376     }

    assert (chunksize(p) >= nb);
    assert((((size_t)(chunk2mem(p))) % alignment) == 0);
4380     check_inuse_chunk(m, p);
    POSTACTION(m);
    if (leader != 0) {
        internal_free(m, leader);
4384     }
    if (trailer != 0) {
        internal_free(m, trailer);
    }
4388     return chunk2mem(p);
    }
}
return 0;
4392 }

```

24.4 comalloc/coalloc support

```

static void** ialloc(mstate m,
4396                     size_t n_elements,
                     size_t* sizes,
                     int opts,
                     void* chunks[]) {
4400     /*
        This provides common support for independent_X routines, handling
        all of the combinations that can result.

4404     The opts arg has:
        bit 0 set if all elements are same size (using sizes[0])
        bit 1 set if elements should be zeroed
    */

    size_t    element_size; /* chunksize of each element, if all same */
    size_t    contents_size; /* total size of elements */
    size_t    array_size;    /* request size of pointer array */
4412     void*    mem;           /* malloced aggregate space */
    mchunkptr p;             /* corresponding chunk */

```

```

size_t    remainder_size; /* remaining bytes while splitting */
void**    marray;         /* either "chunks" or malloced ptr array */
4416 mchunkptr array_chunk; /* chunk for malloced ptr array */
flag_t    was_enabled;    /* to disable mmap */
size_t    size;
size_t    i;

ensure_initialization();
/* compute array length, if needed */
if (chunks != 0) {
4424     if (n_elements == 0)
         return chunks; /* nothing to do */
     marray = chunks;
     array_size = 0;
4428 }
else {
     /* if empty req, must still return chunk representing empty array */
     if (n_elements == 0)
4432         return (void**)internal_malloc(m, 0);
     marray = 0;
     array_size = request2size(n_elements * (sizeof(void*)));
}

/* compute total element size */
if (opts & 0x1) { /* all-same-size */
     element_size = request2size(*sizes);
4440     contents_size = n_elements * element_size;
}
else { /* add up all the sizes */
     element_size = 0;
4444     contents_size = 0;
     for (i = 0; i != n_elements; ++i)
         contents_size += request2size(sizes[i]);
}

size = contents_size + array_size;

/*
4452     Allocate the aggregate chunk.  First disable direct-mmapping so
     malloc won't use it, since we would not be able to later
     free/realloc space internal to a segregated mmap region.
*/
4456 was_enabled = use_mmap(m);
disable_mmap(m);
mem = internal_malloc(m, size - CHUNK_OVERHEAD);
if (was_enabled)
4460     enable_mmap(m);
if (mem == 0)
     return 0;

4464 if (PREACTION(m)) return 0;
p = mem2chunk(mem);

```

```

    remainder_size = chunksize(p);

4468  assert(!is_mmapped(p));

    if (opts & 0x2) {          /* optionally clear the elements */
        memset((size_t*)mem, 0, remainder_size - SIZE_T_SIZE - array_size);
4472  }

    /* If not provided, allocate the pointer array as final part of chunk */
    if (marray == 0) {
4476  size_t  array_chunk_size;
        array_chunk = chunk_plus_offset(p, contents_size);
        array_chunk_size = remainder_size - contents_size;
        marray = (void**) (chunk2mem(array_chunk));
4480  set_size_and_pinuse_of_inuse_chunk(m, array_chunk, array_chunk_size);
        remainder_size = contents_size;
    }

4484  /* split out elements */
    for (i = 0; ; ++i) {
        marray[i] = chunk2mem(p);
        if (i != n_elements-1) {
4488  if (element_size != 0)
            size = element_size;
            else
                size = request2size(sizes[i]);
4492  remainder_size -= size;
            set_size_and_pinuse_of_inuse_chunk(m, p, size);
            p = chunk_plus_offset(p, size);
        }
4496  else { /* the final element absorbs any overallocation slop */
            set_size_and_pinuse_of_inuse_chunk(m, p, remainder_size);
            break;
        }
4500  }

    #if DEBUG
        if (marray != chunks) {
4504  /* final element must have exactly exhausted chunk */
            if (element_size != 0) {
                assert(remainder_size == element_size);
            }
4508  else {
                assert(remainder_size == request2size(sizes[i]));
            }
            check_inuse_chunk(m, mem2chunk(marray));
4512  }
            for (i = 0; i != n_elements; ++i)
                check_inuse_chunk(m, mem2chunk(marray[i]));

4516  #endif /* DEBUG */

```

```
        POSTACTION(m);  
        return marray;  
4520 }
```

Chapter 25

Public routines

```
#if !ONLY_MSPACES

void* dlmalloc(size_t bytes) {
4524     /*
        Basic algorithm:
        If a small request (< 256 bytes minus per-chunk overhead):
            1. If one exists, use a remainderless chunk in associated smallbin.
4528             (Remainderless means that there are too few excess bytes to
                represent as a chunk.)
            2. If it is big enough, use the dv chunk, which is normally the
                chunk adjacent to the one used for the most recent small request.
4532            3. If one exists, split the smallest available chunk in a bin,
                saving remainder in dv.
            4. If it is big enough, use the top chunk.
            5. If available, get memory from system and use it
4536        Otherwise, for a large request:
            1. Find the smallest available binned chunk that fits, and use it
                if it is better fitting than dv chunk, splitting if necessary.
            2. If better fitting than any binned chunk, use the dv chunk.
4540            3. If it is big enough, use the top chunk.
            4. If request size >= mmap threshold, try to directly mmap this chunk.
            5. If available, get memory from system and use it

4544        The ugly goto's here ensure that postaction occurs along all paths.
        */

    #if USE_LOCKS
4548        ensure_initialization(); /* initialize in sys_alloc if not using locks */
    #endif

    if (!PREACTION(gm)) {
4552        void* mem;
        size_t nb;
        if (bytes <= MAX_SMALL_REQUEST) {
            bindex_t idx;
4556            binmap_t smallbits;
            nb = (bytes < MIN_REQUEST)? MIN_CHUNK_SIZE : pad_request(bytes);
            idx = small_index(nb);
            smallbits = gm->smallmap >> idx;
```

```

if ((smallbits & 0x3U) != 0) { /* Remainderless fit to a smallbin. */
    mchunkptr b, p;
    idx += ~smallbits & 1;      /* Uses next bin if idx empty */
4564    b = smallbin_at(gm, idx);
    p = b->fd;
    assert(chunksize(p) == small_index2size(idx));
    unlink_first_small_chunk(gm, b, p, idx);
4568    set_inuse_and_pinuse(gm, p, small_index2size(idx));
    mem = chunk2mem(p);
    check_mallocated_chunk(gm, mem, nb);
    goto postaction;
4572 }

else if (nb > gm->dvsizes) {
    if (smallbits != 0) { /* Use chunk in next nonempty smallbin */
4576        mchunkptr b, p, r;
        size_t rsize;
        bindex_t i;
        binmap_t leftbits = (smallbits << idx) & left_bits(idx2bit(idx));
4580        binmap_t leastbit = least_bit(leftbits);
        compute_bit2idx(leastbit, i);
        b = smallbin_at(gm, i);
        p = b->fd;
4584        assert(chunksize(p) == small_index2size(i));
        unlink_first_small_chunk(gm, b, p, i);
        rsize = small_index2size(i) - nb;
        /* Fit here cannot be remainderless if 4byte sizes */
4588        if (SIZE_T_SIZE != 4 && rsize < MIN_CHUNK_SIZE)
            set_inuse_and_pinuse(gm, p, small_index2size(i));
        else {
            set_size_and_pinuse_of_inuse_chunk(gm, p, nb);
4592            r = chunk_plus_offset(p, nb);
            set_size_and_pinuse_of_free_chunk(r, rsize);
            replace_dv(gm, r, rsize);
        }
4596        mem = chunk2mem(p);
        check_mallocated_chunk(gm, mem, nb);
        goto postaction;
    }

    else if (gm->treemap != 0 && (mem = tmalloc_small(gm, nb)) != 0) {
        check_mallocated_chunk(gm, mem, nb);
        goto postaction;
4604    }
}

else if (bytes >= MAX_REQUEST)
4608    nb = MAX_SIZE_T; /* Too big to allocate. Force failure (in sys alloc) */
else {
    nb = pad_request(bytes);
    if (gm->treemap != 0 && (mem = tmalloc_large(gm, nb)) != 0) {

```



```

4612     check_malloced_chunk(gm, mem, nb);
        goto postaction;
    }
}

if (nb <= gm->dvsizes) {
    size_t rsize = gm->dvsizes - nb;
    mchunkptr p = gm->dv;
4620     if (rsize >= MIN_CHUNK_SIZE) { /* split dv */
        mchunkptr r = gm->dv = chunk_plus_offset(p, nb);
        gm->dvsizes = rsize;
        set_size_and_pinuse_of_free_chunk(r, rsize);
4624         set_size_and_pinuse_of_inuse_chunk(gm, p, nb);
    }
    else { /* exhaust dv */
        size_t dvs = gm->dvsizes;
4628         gm->dvsizes = 0;
        gm->dv = 0;
        set_inuse_and_pinuse(gm, p, dvs);
    }
4632     mem = chunk2mem(p);
    check_malloced_chunk(gm, mem, nb);
    goto postaction;
}

else if (nb < gm->topsize) { /* Split top */
    size_t rsize = gm->topsize -= nb;
    mchunkptr p = gm->top;
4640     mchunkptr r = gm->top = chunk_plus_offset(p, nb);
    r->head = rsize | PINUSE_BIT;
    set_size_and_pinuse_of_inuse_chunk(gm, p, nb);
    mem = chunk2mem(p);
4644     check_top_chunk(gm, gm->top);
    check_malloced_chunk(gm, mem, nb);
    goto postaction;
}

mem = sys_alloc(gm, nb);

postaction:
4652     POSTACTION(gm);
    return mem;
}

4656 return 0;
}

void dlfree(void* mem) {
4660     /*
        Consolidate freed chunks with preceeding or succeeding bordering
        free chunks, if they exist, and then place in a bin. Intermixed
        with special cases for top, dv, mmapped chunks, and usage errors.

```

```

4664  */

    if (mem != 0) {
        mchunkptr p = mem2chunk(mem);
4668 #if FOOTERS
        mstate fm = get_mstate_for(p);
        if (!ok_magic(fm)) {
            USAGE_ERROR_ACTION(fm, p);
4672         return;
        }
    #else /* FOOTERS */
    #define fm gm
4676 #endif /* FOOTERS */
        if (!PREACTION(fm)) {
            check_inuse_chunk(fm, p);
            if (RTCHECK(ok_address(fm, p) && ok_inuse(p))) {
4680         size_t psize = chunksize(p);
            mchunkptr next = chunk_plus_offset(p, psize);
            if (!pinuse(p)) {
                size_t prevsize = p->prev_foot;
4684         if (is_mmapped(p)) {
                    psize += prevsize + MMAP_FOOT_PAD;
                    if (CALL_MUNMAP((char*)p - prevsize, psize) == 0)
                        fm->footprint -= psize;
4688         goto postaction;
                }
            else {
                mchunkptr prev = chunk_minus_offset(p, prevsize);
4692         psize += prevsize;
                p = prev;
                if (RTCHECK(ok_address(fm, prev))) { /* consolidate backward */
                    if (p != fm->dv) {
4696         unlink_chunk(fm, p, prevsize);
                    }
                    else if ((next->head & INUSE_BITS) == INUSE_BITS) {
                        fm->dvsize = psize;
4700         set_free_with_pinuse(p, psize, next);
                        goto postaction;
                    }
                }
            }
4704         else
            goto erroraction;
        }
    }

    if (RTCHECK(ok_next(p, next) && ok_pinuse(next))) {
        if (!cinuse(next)) { /* consolidate forward */
            if (next == fm->top) {
4712         size_t tsize = fm->topsize += psize;
                fm->top = p;
                p->head = tsize | PINUSE_BIT;
                if (p == fm->dv) {

```

```

4716         fm->dv = 0;
            fm->dvsizesize = 0;
        }
        if (should_trim(fm, tsize))
4720         sys_trim(fm, 0);
        goto postaction;
    }
    else if (next == fm->dv) {
4724         size_t dsize = fm->dvsizesize += psize;
        fm->dv = p;
        set_size_and_pinuse_of_free_chunk(p, dsize);
        goto postaction;
4728    }
    else {
        size_t nsize = chunksize(next);
        psize += nsize;
4732        unlink_chunk(fm, next, nsize);
        set_size_and_pinuse_of_free_chunk(p, psize);
        if (p == fm->dv) {
            fm->dvsizesize = psize;
4736            goto postaction;
        }
    }
}
}
4740 else
    set_free_with_pinuse(p, psize, next);

    if (is_small(psize)) {
4744        insert_small_chunk(fm, p, psize);
        check_free_chunk(fm, p);
    }
    else {
4748        tchunkptr tp = (tchunkptr)p;
        insert_large_chunk(fm, tp, psize);
        check_free_chunk(fm, p);
        if (--fm->release_checks == 0)
4752            release_unused_segments(fm);
    }
    goto postaction;
}
}
4756 }
erroraction:
    USAGE_ERROR_ACTION(fm, p);
postaction:
4760    POSTACTION(fm);
}
}
#endif !FOOTERS
4764 #undef fm
#endif /* FOOTERS */
}

```

```

4768 void* dlcalloc(size_t n_elements, size_t elem_size) {
    void* mem;
    size_t req = 0;
    if (n_elements != 0) {
4772     req = n_elements * elem_size;
        if (((n_elements | elem_size) & ~(size_t)0xffff) &&
            (req / n_elements != elem_size))
            req = MAX_SIZE_T; /* force downstream failure on overflow */
4776     }
        mem = dlmalloc(req);
        if (mem != 0 && calloc_must_clear(mem2chunk(mem)))
            memset(mem, 0, req);
4780     return mem;
    }

    void* dlrealloc(void* oldmem, size_t bytes) {
4784     if (oldmem == 0)
        return dlmalloc(bytes);
    #ifdef REALLOC_ZERO_BYTES_FREES
        if (bytes == 0) {
4788         dlfree(oldmem);
            return 0;
        }
    #endif /* REALLOC_ZERO_BYTES_FREES */
4792     else {
    #if ! FOOTERS
        mstate m = gm;
    #else /* FOOTERS */
4796     mstate m = get_mstate_for(mem2chunk(oldmem));
        if (!ok_magic(m)) {
            USAGE_ERROR_ACTION(m, oldmem);
            return 0;
4800         }
    #endif /* FOOTERS */
        return internal_realloc(m, oldmem, bytes);
    }
4804 }

    void* dlmemalign(size_t alignment, size_t bytes) {
4808     return internal_memalign(gm, alignment, bytes);
    }

    void** dlindependent_calloc(size_t n_elements, size_t elem_size,
                                void* chunks[]) {
4812     size_t sz = elem_size; /* serves as 1-element array */
        return ialloc(gm, n_elements, &sz, 3, chunks);
    }

    void** dlindependent_comalloc(size_t n_elements, size_t sizes[],
                                void* chunks[]) {
4816     return ialloc(gm, n_elements, sizes, 0, chunks);
    }

```

```

void* dlvalloc(size_t bytes) {
    size_t pagesz;
    ensure_initialization();
4824    pagesz = mparams.page_size;
    return dlmemalign(pagesz, bytes);
}

4828 void* dlpvalloc(size_t bytes) {
    size_t pagesz;
    ensure_initialization();
    pagesz = mparams.page_size;
4832    return dlmemalign(pagesz, (bytes + pagesz - SIZE_T_ONE) & ~(pagesz - SIZE_T_ONE));
}

int dlmalloc_trim(size_t pad) {
4836    int result = 0;
    ensure_initialization();
    if (!PREACTION(gm)) {
        result = sys_trim(gm, pad);
4840    POSTACTION(gm);
    }
    return result;
}

size_t dlmalloc_footprint(void) {
    return gm->footprint;
}

size_t dlmalloc_max_footprint(void) {
    return gm->max_footprint;
}

#if !NO_MALLINFO
struct mallinfo dlmallinfo(void) {
    return internal_mallinfo(gm);
4856 }
#endif /* NO_MALLINFO */

void dlmalloc_stats() {
4860    internal_malloc_stats(gm);
}

int dlmallopt(int param_number, int value) {
4864    return change_mparam(param_number, value);
}

#endif /* !ONLY_MSPACES */

size_t dlmalloc_usable_size(void* mem) {
    if (mem != 0) {
        mchunkptr p = mem2chunk(mem);

```

```
4872     if (is_inuse(p))
        return chunksize(p) - overhead_for(p);
    }
    return 0;
4876 }
```

Chapter 26

User mspaces

```
#if MSPACES

static mstate init_user_mstate(char* tbase, size_t tsize) {
4880   size_t msize = pad_request(sizeof(struct malloc_state));
       mchunkptr mn;
       mchunkptr msp = align_as_chunk(tbase);
       mstate m = (mstate)(chunk2mem(msp));
4884   memset(m, 0, msize);
       INITIAL_LOCK(&m->mutex);
       msp->head = (msize|INUSE_BITS);
       m->seg.base = m->least_addr = tbase;
4888   m->seg.size = m->footprint = m->max_footprint = tsize;
       m->magic = mparams.magic;
       m->release_checks = MAX_RELEASE_CHECK_RATE;
       m->mflags = mparams.default_mflags;
4892   m->extp = 0;
       m->exts = 0;
       disable_contiguous(m);
       init_bins(m);
4896   mn = next_chunk(mem2chunk(m));
       init_top(m, mn, (size_t)((tbase + tsize) - (char*)mn) - TOP_FOOT_SIZE);
       check_top_chunk(m, m->top);
       return m;
4900 }

mspace create_mspace(size_t capacity, int locked) {
       mstate m = 0;
4904   size_t msize;
       ensure_initialization();
       msize = pad_request(sizeof(struct malloc_state));
       if (capacity < (size_t) -(msize + TOP_FOOT_SIZE + mparams.page_size)) {
4908   size_t rs = ((capacity == 0)? mparams.granularity :
               (capacity + TOP_FOOT_SIZE + msize));
       size_t tsize = granularity_align(rs);
       char* tbase = (char*)(CALL_MMAP(tsize));
4912   if (tbase != CMFAIL) {
       m = init_user_mstate(tbase, tsize);
       m->seg.sflags = USE_MMAP_BIT;
       set_lock(m, locked);

```

```

4916     }
    }
    return (mspace)m;
}

mspace create_mspace_with_base(void* base, size_t capacity, int locked) {
    mstate m = 0;
    size_t msize;
4924    ensure_initialization();
    msize = pad_request(sizeof(struct malloc_state));
    if (capacity > msize + TOP_FOOT_SIZE &&
        capacity < (size_t) -(msize + TOP_FOOT_SIZE + mparams.page_size)) {
4928        m = init_user_mstate((char*)base, capacity);
        m->seg.sflags = EXTERN_BIT;
        set_lock(m, locked);
    }
4932    return (mspace)m;
}

int mspace_track_large_chunks(mspace msp, int enable) {
4936    int ret = 0;
    mstate ms = (mstate)msp;
    if (!PREACTION(ms)) {
        if (!use_mmap(ms))
4940            ret = 1;
        if (!enable)
            enable_mmap(ms);
        else
4944            disable_mmap(ms);
        POSTACTION(ms);
    }
    return ret;
4948 }

size_t destroy_mspace(mspace msp) {
    size_t freed = 0;
4952    mstate ms = (mstate)msp;
    if (ok_magic(ms)) {
        msegmentptr sp = &ms->seg;
        while (sp != 0) {
4956            char* base = sp->base;
            size_t size = sp->size;
            flag_t flag = sp->sflags;
            sp = sp->next;
4960            if ((flag & USE_MMAP_BIT) && !(flag & EXTERN_BIT) &&
                CALL_MUNMAP(base, size) == 0)
                freed += size;
        }
4964    }
    else {
        USAGE_ERROR_ACTION(ms,ms);
    }
}

```



```

4968     return freed;
    }

    /*
4972     mspace versions of routines are near-clones of the global
        versions. This is not so nice but better than the alternatives.
    */

4976 void* mspace_malloc(mspace msp, size_t bytes) {
    mstate ms = (mstate)msp;
    if (!ok_magic(ms)) {
        USAGE_ERROR_ACTION(ms,ms);
4980     return 0;
    }
    if (!PREACTION(ms)) {
        void* mem;
4984     size_t nb;
        if (bytes <= MAX_SMALL_REQUEST) {
            bindex_t idx;
            binmap_t smallbits;
4988     nb = (bytes < MIN_REQUEST)? MIN_CHUNK_SIZE : pad_request(bytes);
            idx = small_index(nb);
            smallbits = ms->smallmap >> idx;

4992     if ((smallbits & 0x3U) != 0) { /* Remainderless fit to a smallbin. */
                mchunkptr b, p;
                idx += ~smallbits & 1; /* Uses next bin if idx empty */
                b = smallbin_at(ms, idx);
4996     p = b->fd;
                assert(chunksize(p) == small_index2size(idx));
                unlink_first_small_chunk(ms, b, p, idx);
                set_inuse_and_pinuse(ms, p, small_index2size(idx));
5000     mem = chunk2mem(p);
                check_mallocated_chunk(ms, mem, nb);
                goto postaction;
            }

            else if (nb > ms->dvsize) {
                if (smallbits != 0) { /* Use chunk in next nonempty smallbin */
                    mchunkptr b, p, r;
5008     size_t rsize;
                    bindex_t i;
                    binmap_t leftbits = (smallbits << idx) & left_bits(idx2bit(idx));
                    binmap_t leastbit = least_bit(leftbits);
                    compute_bit2idx(leastbit, i);
5012     b = smallbin_at(ms, i);
                    p = b->fd;
                    assert(chunksize(p) == small_index2size(i));
                    unlink_first_small_chunk(ms, b, p, i);
                    rsize = small_index2size(i) - nb;
                    /* Fit here cannot be remainderless if 4byte sizes */
                    if (SIZE_T_SIZE != 4 && rsize < MIN_CHUNK_SIZE)

```

```

5020     set_inuse_and_pinuse(ms, p, small_index2size(i));
    else {
        set_size_and_pinuse_of_inuse_chunk(ms, p, nb);
        r = chunk_plus_offset(p, nb);
5024     set_size_and_pinuse_of_free_chunk(r, rsize);
        replace_dv(ms, r, rsize);
    }
    mem = chunk2mem(p);
5028     check_malloced_chunk(ms, mem, nb);
    goto postaction;
}

5032     else if (ms->treemap != 0 && (mem = tmalloc_small(ms, nb)) != 0) {
        check_malloced_chunk(ms, mem, nb);
        goto postaction;
    }
5036 }
}
else if (bytes >= MAX_REQUEST)
    nb = MAX_SIZE_T; /* Too big to allocate. Force failure (in sys alloc) */
5040 else {
    nb = pad_request(bytes);
    if (ms->treemap != 0 && (mem = tmalloc_large(ms, nb)) != 0) {
        check_malloced_chunk(ms, mem, nb);
5044     goto postaction;
    }
}

5048 if (nb <= ms->dvsize) {
    size_t rsize = ms->dvsize - nb;
    mchunkptr p = ms->dv;
    if (rsize >= MIN_CHUNK_SIZE) { /* split dv */
5052     mchunkptr r = ms->dv = chunk_plus_offset(p, nb);
        ms->dvsize = rsize;
        set_size_and_pinuse_of_free_chunk(r, rsize);
        set_size_and_pinuse_of_inuse_chunk(ms, p, nb);
5056     }
    else { /* exhaust dv */
        size_t dvs = ms->dvsize;
        ms->dvsize = 0;
5060     ms->dv = 0;
        set_inuse_and_pinuse(ms, p, dvs);
    }
    mem = chunk2mem(p);
5064     check_malloced_chunk(ms, mem, nb);
    goto postaction;
}

5068 else if (nb < ms->topsize) { /* Split top */
    size_t rsize = ms->topsize - nb;
    mchunkptr p = ms->top;
    mchunkptr r = ms->top = chunk_plus_offset(p, nb);

```

```

5072     r->head = rsize | PINUSE_BIT;
        set_size_and_pinuse_of_inuse_chunk(ms, p, nb);
        mem = chunk2mem(p);
        check_top_chunk(ms, ms->top);
5076     check_mallocated_chunk(ms, mem, nb);
        goto postaction;
    }

5080     mem = sys_alloc(ms, nb);

    postaction:
        POSTACTION(ms);
5084     return mem;
    }

    return 0;
5088 }

void mspace_free(mspace msp, void* mem) {
    if (mem != 0) {
5092         mchunkptr p = mem2chunk(mem);
        #if FOOTERS
            mstate fm = get_mstate_for(p);
            msp = msp; /* placate people compiling -Wunused */
5096 #else /* FOOTERS */
            mstate fm = (mstate)msp;
        #endif /* FOOTERS */
        if (!ok_magic(fm)) {
5100             USAGE_ERROR_ACTION(fm, p);
            return;
        }
        if (!PREACTION(fm)) {
5104             check_inuse_chunk(fm, p);
            if (RTCHECK(ok_address(fm, p) && ok_inuse(p))) {
                size_t psize = chunksize(p);
                mchunkptr next = chunk_plus_offset(p, psize);
5108                 if (!pinuse(p)) {
                    size_t prevsize = p->prev_foot;
                    if (is_mmapped(p)) {
                        psize += prevsize + MMAP_FOOT_PAD;
5112                         if (CALL_MUNMAP((char*)p - prevsize, psize) == 0)
                            fm->footprint -= psize;
                        goto postaction;
                    }
                }
                else {
5116                     mchunkptr prev = chunk_minus_offset(p, prevsize);
                    psize += prevsize;
                    p = prev;
5120                     if (RTCHECK(ok_address(fm, prev))) { /* consolidate backward */
                        if (p != fm->dv) {
                            unlink_chunk(fm, p, prevsize);
                        }
                    }
                }
            }
        }
    }
}

```

```

5124         else if ((next->head & INUSE_BITS) == INUSE_BITS) {
            fm->dvsizesize = psize;
            set_free_with_pinuse(p, psize, next);
            goto postaction;
5128         }
        }
        else
            goto erroraction;
5132     }
}

if (RTCHECK(ok_next(p, next) && ok_pinuse(next))) {
5136     if (!cinuse(next)) { /* consolidate forward */
        if (next == fm->top) {
            size_t tsize = fm->topsize += psize;
            fm->top = p;
5140            p->head = tsize | PINUSE_BIT;
            if (p == fm->dv) {
                fm->dv = 0;
                fm->dvsizesize = 0;
5144            }
            if (should_trim(fm, tsize))
                sys_trim(fm, 0);
            goto postaction;
5148        }
        else if (next == fm->dv) {
            size_t dsize = fm->dvsizesize += psize;
            fm->dv = p;
5152            set_size_and_pinuse_of_free_chunk(p, dsize);
            goto postaction;
        }
        else {
5156            size_t nsize = chunksize(next);
            psize += nsize;
            unlink_chunk(fm, next, nsize);
            set_size_and_pinuse_of_free_chunk(p, psize);
5160            if (p == fm->dv) {
                fm->dvsizesize = psize;
                goto postaction;
            }
5164        }
    }
}
else
    set_free_with_pinuse(p, psize, next);

if (is_small(psize)) {
    insert_small_chunk(fm, p, psize);
    check_free_chunk(fm, p);
5172 }
else {
    tchunkptr tp = (tchunkptr)p;
    insert_large_chunk(fm, tp, psize);

```

```

5176         check_free_chunk(fm, p);
            if (--fm->release_checks == 0)
                release_unused_segments(fm);
        }
5180     goto postaction;
    }
    }
    erroraction:
5184     USAGE_ERROR_ACTION(fm, p);
    postaction:
        POSTACTION(fm);
    }
5188 }
}

void* mspace_calloc(mspace msp, size_t n_elements, size_t elem_size) {
5192     void* mem;
    size_t req = 0;
    mstate ms = (mstate)msp;
    if (!ok_magic(ms)) {
5196         USAGE_ERROR_ACTION(ms,ms);
        return 0;
    }
    if (n_elements != 0) {
5200         req = n_elements * elem_size;
        if (((n_elements | elem_size) & ~(size_t)0xffff) &&
            (req / n_elements != elem_size))
            req = MAX_SIZE_T; /* force downstream failure on overflow */
5204     }
    mem = internal_malloc(ms, req);
    if (mem != 0 && calloc_must_clear(mem2chunk(mem)))
        memset(mem, 0, req);
5208     return mem;
}

void* mspace_realloc(mspace msp, void* oldmem, size_t bytes) {
5212     if (oldmem == 0)
        return mspace_malloc(msp, bytes);
#ifdef REALLOC_ZERO_BYTES_FREES
    if (bytes == 0) {
5216         mspace_free(msp, oldmem);
        return 0;
    }
#endif /* REALLOC_ZERO_BYTES_FREES */
    else {
5220         #if FOOTERS
            mchunkptr p = mem2chunk(oldmem);
            mstate ms = get_mstate_for(p);
5224         #else /* FOOTERS */
            mstate ms = (mstate)msp;
        #endif /* FOOTERS */
        if (!ok_magic(ms)) {

```

```

5228     USAGE_ERROR_ACTION(ms,ms);
        return 0;
    }
    return internal_realloc(ms, oldmem, bytes);
5232 }
}

void* mspace_memalign(mspace msp, size_t alignment, size_t bytes) {
5236     mstate ms = (mstate)msp;
    if (!ok_magic(ms)) {
        USAGE_ERROR_ACTION(ms,ms);
        return 0;
5240     }
    return internal_memalign(ms, alignment, bytes);
}

5244 void** mspace_independent_calloc(mspace msp, size_t n_elements,
                                   size_t elem_size, void* chunks[]) {
    size_t sz = elem_size; /* serves as 1-element array */
    mstate ms = (mstate)msp;
5248     if (!ok_magic(ms)) {
        USAGE_ERROR_ACTION(ms,ms);
        return 0;
    }
5252     return ialloc(ms, n_elements, &sz, 3, chunks);
}

void** mspace_independent_comalloc(mspace msp, size_t n_elements,
5256                                   size_t sizes[], void* chunks[]) {
    mstate ms = (mstate)msp;
    if (!ok_magic(ms)) {
        USAGE_ERROR_ACTION(ms,ms);
5260     return 0;
    }
    return ialloc(ms, n_elements, sizes, 0, chunks);
}

int mspace_trim(mspace msp, size_t pad) {
    int result = 0;
    mstate ms = (mstate)msp;
5268     if (ok_magic(ms)) {
        if (!PREACTION(ms)) {
            result = sys_trim(ms, pad);
            POSTACTION(ms);
5272     }
    }
    else {
        USAGE_ERROR_ACTION(ms,ms);
5276     }
    return result;
}

```

```

5280 void mspace_malloc_stats(mspace msp) {
    mstate ms = (mstate)msp;
    if (ok_magic(ms)) {
        internal_malloc_stats(ms);
5284    }
    else {
        USAGE_ERROR_ACTION(ms,ms);
    }
5288 }

size_t mspace_footprint(mspace msp) {
    size_t result = 0;
5292    mstate ms = (mstate)msp;
    if (ok_magic(ms)) {
        result = ms->footprint;
    }
5296    else {
        USAGE_ERROR_ACTION(ms,ms);
    }
    return result;
5300 }

size_t mspace_max_footprint(mspace msp) {
5304    size_t result = 0;
    mstate ms = (mstate)msp;
    if (ok_magic(ms)) {
        result = ms->max_footprint;
5308    }
    else {
        USAGE_ERROR_ACTION(ms,ms);
    }
5312    return result;
}

5316 #if !NO_MALLINFO
struct mallinfo mspace_mallinfo(mspace msp) {
    mstate ms = (mstate)msp;
    if (!ok_magic(ms)) {
5320        USAGE_ERROR_ACTION(ms,ms);
    }
    return internal_mallinfo(ms);
}
5324 #endif /* NO_MALLINFO */

size_t mspace_usable_size(void* mem) {
    if (mem != 0) {
5328        mchunkptr p = mem2chunk(mem);
        if (is_inuse(p))
            return chunksize(p) - overhead_for(p);
    }
}

```

```
5332     return 0;
    }

    int mspace_mallopt(int param_number, int value) {
5336     return change_mparam(param_number, value);
    }

    #endif /* MSPACES */
```


Chapter 27

Postscript

5340

27.1 Alternative MORECORE functions

/*

Guidelines for creating a custom version of MORECORE:

5344

* For best performance, MORECORE should allocate in multiples of pagesize.

* MORECORE may allocate more memory than requested. (Or even less, but this will usually result in a malloc failure.)

5348

* MORECORE must not allocate memory when given argument zero, but instead return one past the end address of memory from previous nonzero call.

5352

* For best performance, consecutive calls to MORECORE with positive arguments should return increasing addresses, indicating that space has been contiguously extended.

* Even though consecutive calls to MORECORE need not return contiguous addresses, it must be OK for malloc'ed chunks to span multiple regions in those cases where they do happen to be contiguous.

5356

* MORECORE need not handle negative arguments -- it may instead just return MFAIL when given negative arguments.

5360

Negative arguments are always multiples of pagesize. MORECORE must not misinterpret negative args as large positive unsigned args. You can suppress all such calls from even occurring by defining MORECORE_CANNOT_TRIM,

As an example alternative MORECORE, here is a custom allocator kindly contributed for pre-OSX macOS. It uses virtually but not necessarily physically contiguous non-paged memory (locked in, present and won't get swapped out). You can use it by uncommenting this section, adding some #includes, and setting up the appropriate defines above:

5368

```
#define MORECORE osMoreCore
```

5372

There is also a shutdown routine that should somehow be called for cleanup upon program exit.

```
#define MAX_POOL_ENTRIES 100
```

5376

```
#define MINIMUM_MORECORE_SIZE (64 * 1024U)
```

```

static int next_os_pool;
void *our_os_pools[MAX_POOL_ENTRIES];

5380 void *osMoreCore(int size)
{
    void *ptr = 0;
    static void *sbrk_top = 0;

    if (size > 0)
    {
        if (size < MINIMUM_MORECORE_SIZE)
5388     size = MINIMUM_MORECORE_SIZE;
        if (CurrentExecutionLevel() == kTaskLevel)
            ptr = PoolAllocateResident(size + RM_PAGE_SIZE, 0);
        if (ptr == 0)
5392     {
            return (void *) MFAIL;
        }
        // save ptrs so they can be freed during cleanup
5396     our_os_pools[next_os_pool] = ptr;
        next_os_pool++;
        ptr = (void *) (((size_t) ptr) + RM_PAGE_MASK) & ~RM_PAGE_MASK;
        sbrk_top = (char *) ptr + size;
5400     return ptr;
    }
    else if (size < 0)
    {
5404     // we don't currently support shrink behavior
        return (void *) MFAIL;
    }
    else
5408     {
        return sbrk_top;
    }
}

// cleanup any allocated memory pools
// called as last thing before shutting down driver

5416 void osCleanupMem(void)
{
    void **ptr;

5420     for (ptr = our_os_pools; ptr < &our_os_pools[MAX_POOL_ENTRIES]; ptr++)
        if (*ptr)
        {
            PoolDeallocate(*ptr);
5424     *ptr = 0;
        }
}

5428 */

```

27.2 History

```

/*
V2.8.4 Wed May 27 09:56:23 2009  Doug Lea  (dl at gee)
5432  * Use zeros instead of prev foot for is_mmapped
      * Add mspace_track_large_chunks; thanks to Jean Brouwers
      * Fix set_inuse in internal_realloc; thanks to Jean Brouwers
      * Fix insufficient sys_alloc padding when using 16byte alignment
5436  * Fix bad error check in mspace_footprint
      * Adaptations for ptmalloc; thanks to Wolfram Gloger.
      * Reentrant spin locks; thanks to Earl Chew and others
      * Win32 improvements; thanks to Niall Douglas and Earl Chew
5440  * Add NO_SEGMENT_TRAVERSAL and MAX_RELEASE_CHECK_RATE options
      * Extension hook in malloc_state
      * Various small adjustments to reduce warnings on some compilers
      * Various configuration extensions/changes for more platforms. Thanks
5444  to all who contributed these.

V2.8.3 Thu Sep 22 11:16:32 2005  Doug Lea  (dl at gee)
      * Add max_footprint functions
5448  * Ensure all appropriate literals are size_t
      * Fix conditional compilation problem for some #define settings
      * Avoid concatenating segments with the one provided
        in create_mspace_with_base
5452  * Rename some variables to avoid compiler shadowing warnings
      * Use explicit lock initialization.
      * Better handling of sbrk interference.
      * Simplify and fix segment insertion, trimming and mspace_destroy
5456  * Reinstate REALLOC_ZERO_BYTES_FREES option from 2.7.x
      * Thanks especially to Dennis Flanagan for help on these.

V2.8.2 Sun Jun 12 16:01:10 2005  Doug Lea  (dl at gee)
5460  * Fix memalign brace error.

V2.8.1 Wed Jun  8 16:11:46 2005  Doug Lea  (dl at gee)
      * Fix improper #endif nesting in C++
5464  * Add explicit casts needed for C++

V2.8.0 Mon May 30 14:09:02 2005  Doug Lea  (dl at gee)
      * Use trees for large bins
5468  * Support mspace
      * Use segments to unify sbrk-based and mmap-based system allocation,
        removing need for emulation on most platforms without sbrk.
      * Default safety checks
5472  * Optional footer checks. Thanks to William Robertson for the idea.
      * Internal code refactoring
      * Incorporate suggestions and platform-specific changes.
        Thanks to Dennis Flanagan, Colin Plumb, Niall Douglas,
5476  Aaron Bachmann, Emery Berger, and others.
      * Speed up non-fastbin processing enough to remove fastbins.
      * Remove useless cfree() to avoid conflicts with other apps.

```

5480 * Remove internal memcpy, memset. Compilers handle builtins better.
 * Remove some options that no one ever used and rename others.

V2.7.2 Sat Aug 17 09:07:30 2002 Doug Lea (dl at gee)

 * Fix malloc_state bitmap array misdeclaration

V2.7.1 Thu Jul 25 10:58:03 2002 Doug Lea (dl at gee)

5488 * Allow tuning of FIRST_SORTED_BIN_SIZE
 * Use PTR_UINT as type for all ptr->int casts. Thanks to John Belmonte.
 * Better detection and support for non-contiguousness of MORECORE.
 Thanks to Andreas Mueller, Conal Walsh, and Wolfram Gloger
 * Bypass most of malloc if no frees. Thanks To Emery Berger.
 * Fix freeing of old top non-contiguous chunk in sysmalloc.
 5492 * Raised default trim and map thresholds to 256K.
 * Fix mmap-related #defines. Thanks to Lubos Lunak.
 * Fix copy macros; added LACKS_FCNTL_H. Thanks to Neal Walfield.
 * Branch-free bin calculation
 5496 * Default trim and mmap thresholds now 256K.

V2.7.0 Sun Mar 11 14:14:06 2001 Doug Lea (dl at gee)

5500 * Introduce independent_comalloc and independent_calloc.
 Thanks to Michael Pachos for motivation and help.
 * Make optional .h file available
 * Allow > 2GB requests on 32bit systems.
 * new WIN32 sbrk, mmap, munmap, lock code from <Walter@GeNeSys-e.de>.
 5504 Thanks also to Andreas Mueller <a.mueller at paradatac.de>,
 and Anonymous.
 * Allow override of MALLOC_ALIGNMENT (Thanks to Ruud Waij for
 helping test this.)
 5508 * memalign: check alignment arg
 * realloc: don't try to shift chunks backwards, since this
 leads to more fragmentation in some programs and doesn't
 seem to help in any others.
 5512 * Collect all cases in malloc requiring system memory into sysmalloc
 * Use mmap as backup to sbrk
 * Place all internal state in malloc_state
 * Introduce fastbins (although similar to 2.5.1)
 5516 * Many minor tunings and cosmetic improvements
 * Introduce USE_PUBLIC_MALLOC_WRAPPERS, USE_MALLOC_LOCK
 * Introduce MALLOC_FAILURE_ACTION, MORECORE_CONTIGUOUS
 Thanks to Tony E. Bennett <tbennett@nvidia.com> and others.
 5520 * Include errno.h to support default failure action.

V2.6.6 Sun Dec 5 07:42:19 1999 Doug Lea (dl at gee)

5524 * return null for negative arguments
 * Added Several WIN32 cleanups from Martin C. Fong <mcfong at yahoo.com>
 * Add 'LACKS_SYS_PARAM_H' for those systems without 'sys/param.h'
 (e.g. WIN32 platforms)
 * Cleanup header file inclusion for WIN32 platforms
 5528 * Cleanup code to avoid Microsoft Visual C++ compiler complaints
 * Add 'USE_DL_PREFIX' to quickly allow co-existence with existing
 memory allocation routines

```

5532      * Set 'malloc_getpagesize' for WIN32 platforms (needs more work)
      * Use 'assert' rather than 'ASSERT' in WIN32 code to conform to
        usage of 'assert' in non-WIN32 code
      * Improve WIN32 'sbrk()' emulation's 'findRegion()' routine to
        avoid infinite loop
5536      * Always call 'fREe()' rather than 'free()'

```

V2.6.5 Wed Jun 17 15:57:31 1998 Doug Lea (dl at gee)

```

      * Fixed ordering problem with boundary-stamping

```

V2.6.3 Sun May 19 08:17:58 1996 Doug Lea (dl at gee)

```

      * Added pvalloc, as recommended by H.J. Liu
      * Added 64bit pointer support mainly from Wolfram Gloger
5544      * Added anonymously donated WIN32 sbrk emulation
      * Malloc, calloc, getpagesize: add optimizations from Raymond Nijssen
      * malloc_extend_top: fix mask error that caused wastage after
        foreign sbrks
5548      * Add linux mmap support code from HJ Liu

```

V2.6.2 Tue Dec 5 06:52:55 1995 Doug Lea (dl at gee)

```

      * Integrated most documentation with the code.
5552      * Add support for mmap, with help from
        Wolfram Gloger (Gloger@lrz.uni-muenchen.de).
      * Use last_remainder in more cases.
      * Pack bins using idea from colin@nyx10.cs.du.edu
5556      * Use ordered bins instead of best-fit threshold
      * Eliminate block-local decls to simplify tracing and debugging.
      * Support another case of realloc via move into top
      * Fix error occuring when initial sbrk_base not word-aligned.
5560      * Rely on page size for units instead of SBRK_UNIT to
        avoid surprises about sbrk alignment conventions.
      * Add mallinfo, mallopt. Thanks to Raymond Nijssen
        (raymond@es.ele.tue.nl) for the suggestion.
5564      * Add 'pad' argument to malloc_trim and top_pad mallopt parameter.
      * More precautions for cases where other routines call sbrk,
        courtesy of Wolfram Gloger (Gloger@lrz.uni-muenchen.de).
      * Added macros etc., allowing use in linux libc from
5568      H.J. Lu (hjl@gnu.ai.mit.edu)
      * Inverted this history list

```

V2.6.1 Sat Dec 2 14:10:57 1995 Doug Lea (dl at gee)

```

5572      * Re-tuned and fixed to behave more nicely with V2.6.0 changes.
      * Removed all preallocation code since under current scheme
        the work required to undo bad preallocations exceeds
        the work saved in good cases for most test programs.
5576      * No longer use return list or unconsolidated bins since
        no scheme using them consistently outperforms those that don't
        given above changes.
      * Use best fit for very large chunks to prevent some worst-cases.
5580      * Added some support for debugging

```

V2.6.0 Sat Nov 4 07:05:23 1995 Doug Lea (dl at gee)

```

    * Removed footers when chunks are in use. Thanks to
5584     Paul Wilson (wilson@cs.texas.edu) for the suggestion.

V2.5.4 Wed Nov  1 07:54:51 1995  Doug Lea  (dl at gee)
    * Added malloc_trim, with help from Wolfram Gloger
5588     (wmglo@Dent.MED.Uni-Muenchen.DE).

V2.5.3 Tue Apr 26 10:16:01 1994  Doug Lea  (dl at g)

5592 V2.5.2 Tue Apr  5 16:20:40 1994  Doug Lea  (dl at g)
    * realloc: try to expand in both directions
    * malloc: swap order of clean-bin strategy;
    * realloc: only conditionally expand backwards
5596    * Try not to scavenge used bins
    * Use bin counts as a guide to preallocation
    * Occasionally bin return list chunks in first scan
    * Add a few optimizations from colin@nyx10.cs.du.edu

V2.5.1 Sat Aug 14 15:40:43 1993  Doug Lea  (dl at g)
    * faster bin computation & slightly different binning
    * merged all consolidations to one part of malloc proper
5604     (eliminating old malloc_find_space & malloc_clean_bin)
    * Scan 2 returns chunks (not just 1)
    * Propagate failure in realloc if malloc returns 0
    * Add stuff to allow compilation on non-ANSI compilers
5608     from kpv@research.att.com

V2.5 Sat Aug  7 07:41:59 1993  Doug Lea  (dl at g.oswego.edu)
    * removed potential for odd address access in prev_chunk
5612    * removed dependency on getpagesize.h
    * misc cosmetics and a bit more internal documentation
    * anticosmetics: mangled names in macros to evade debugger strangeness
    * tested on sparc, hp-700, dec-mips, rs6000
5616     with gcc & native cc (hp, dec only) allowing
    Detlefs & Zorn comparison study (in SIGPLAN Notices.)

Trial version Fri Aug 28 13:14:29 1992  Doug Lea  (dl at g.oswego.edu)
5620    * Based loosely on libg++-1.2X malloc. (It retains some of the overall
    structure of old version, but most details differ.)

*/

```